



Clustering 101

Dorothy Quincy

Dorothy.quincy@ibm.com

Agenda


- Clustering, what is it and why is it needed?
- How Clustering works and basic setup
- Benefits of Clustering
 - Workload Balancing
 - Service Availability
- Uniform clusters
 - What are they?
 - How to configure them

Why is clustering needed?

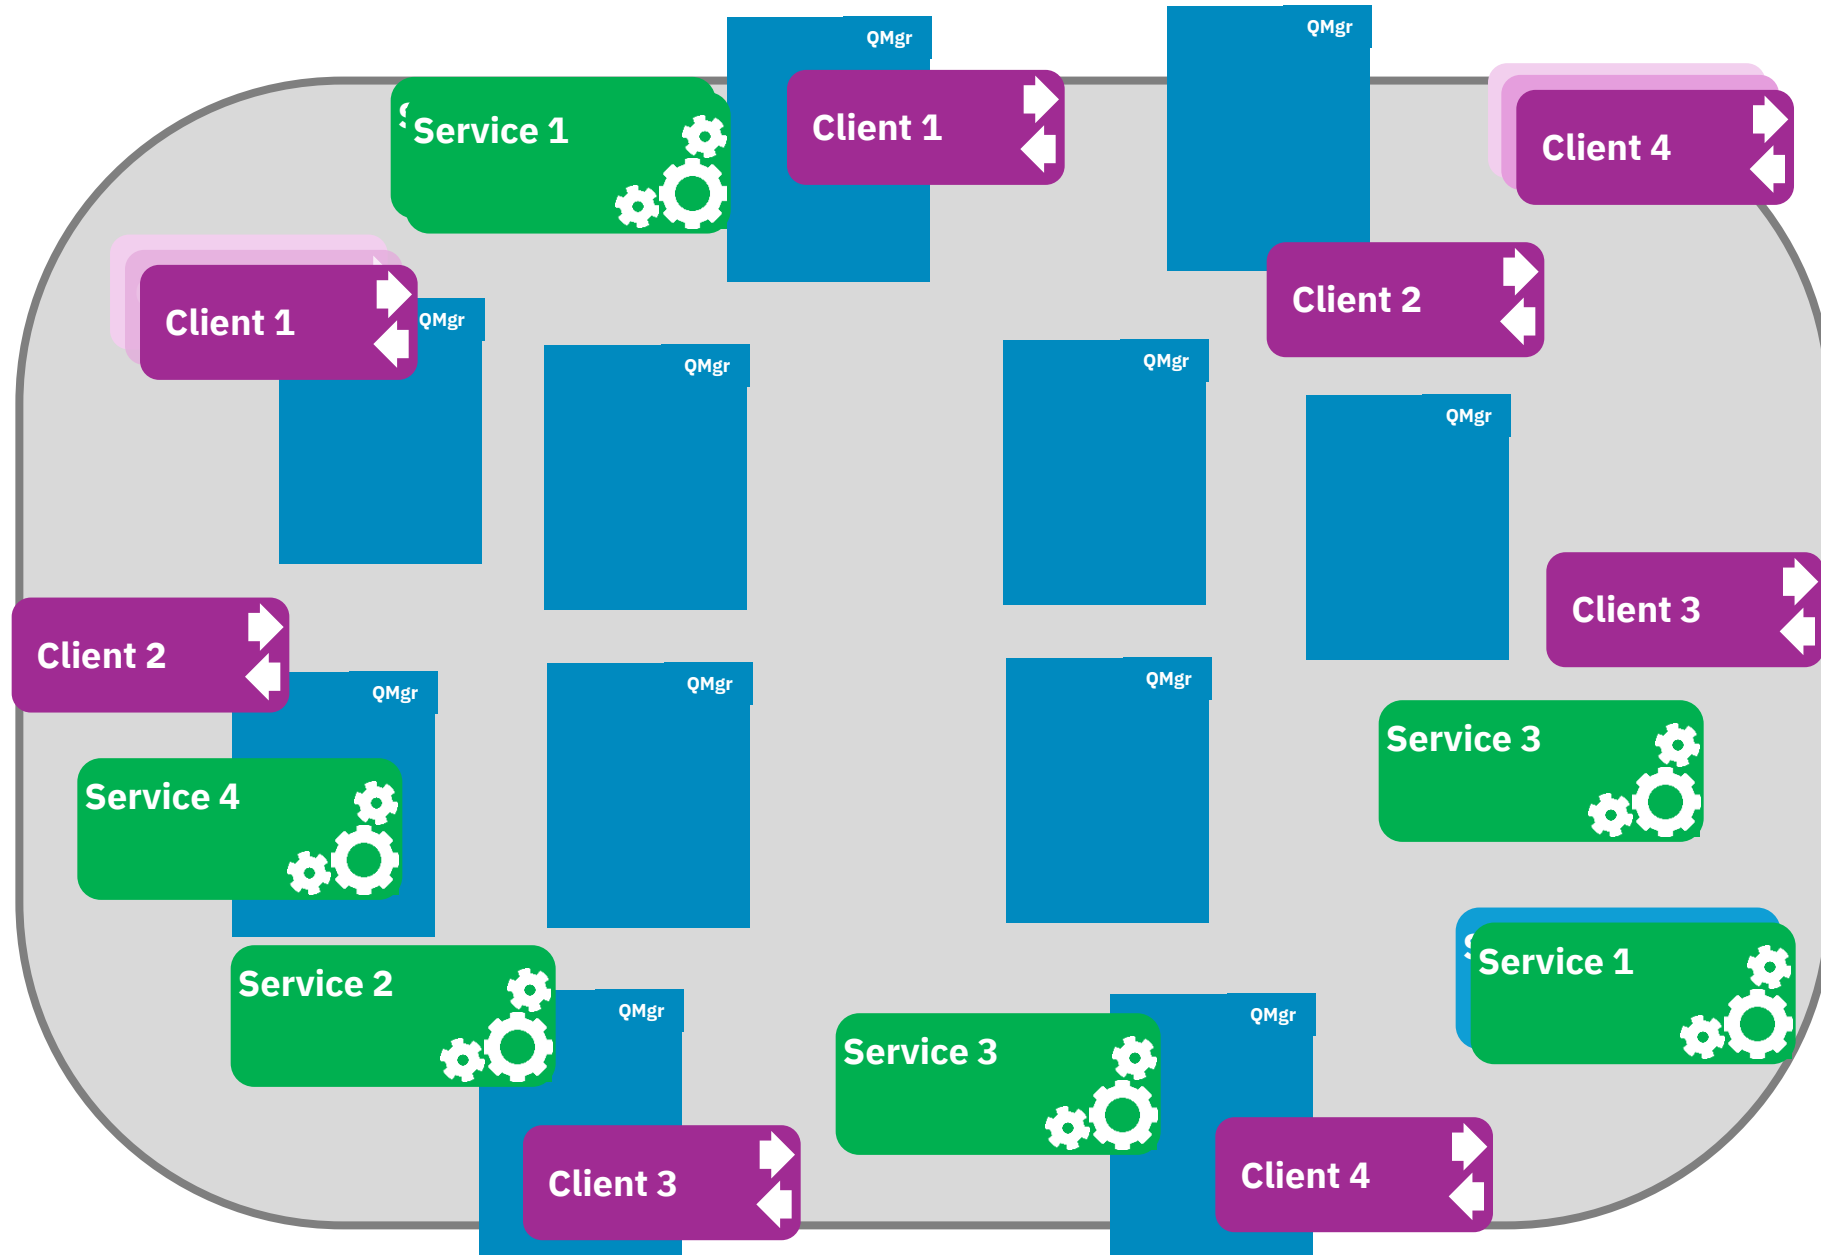
Where it all begins...



Over time...

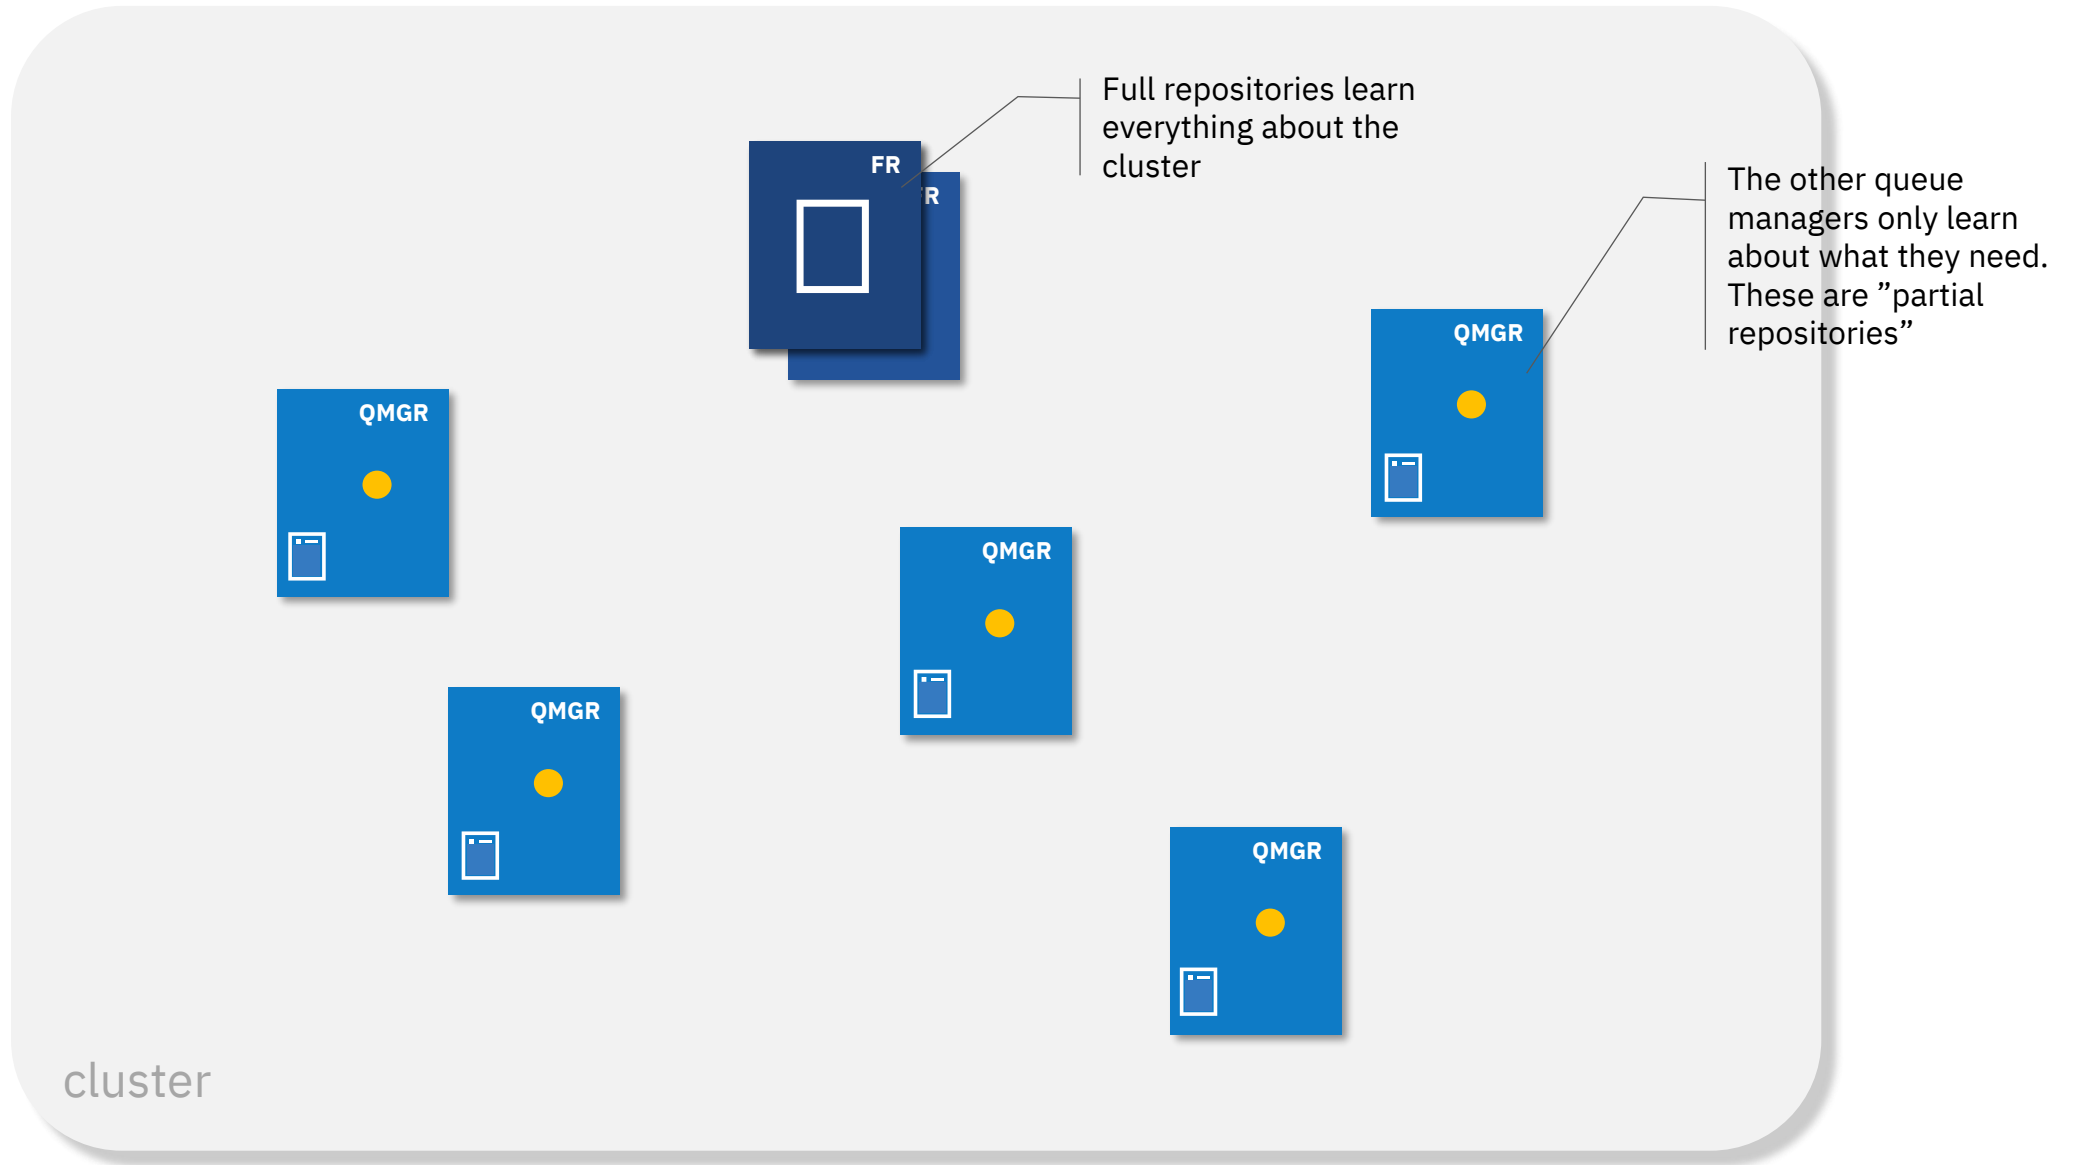


...Now you need a cluster



What is a Cluster?

Cluster overview

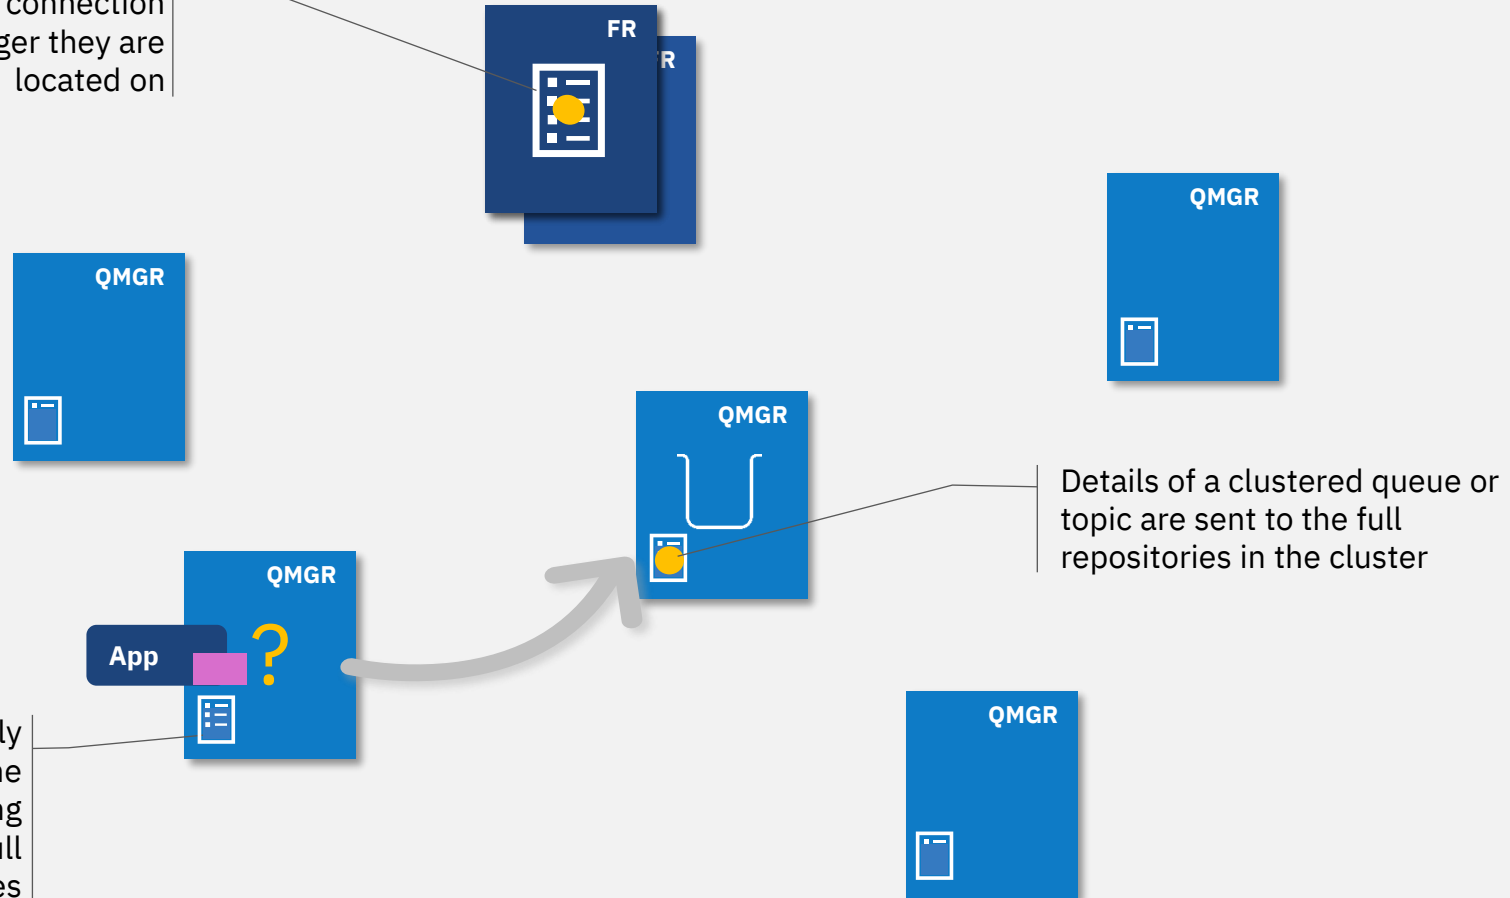


Cluster management


Full repositories will pass on the details of cluster queues and the connection details of the queue manager they are located on

Queue managers persistently cache their knowledge of the cluster resources, limiting interactions with the full repositories

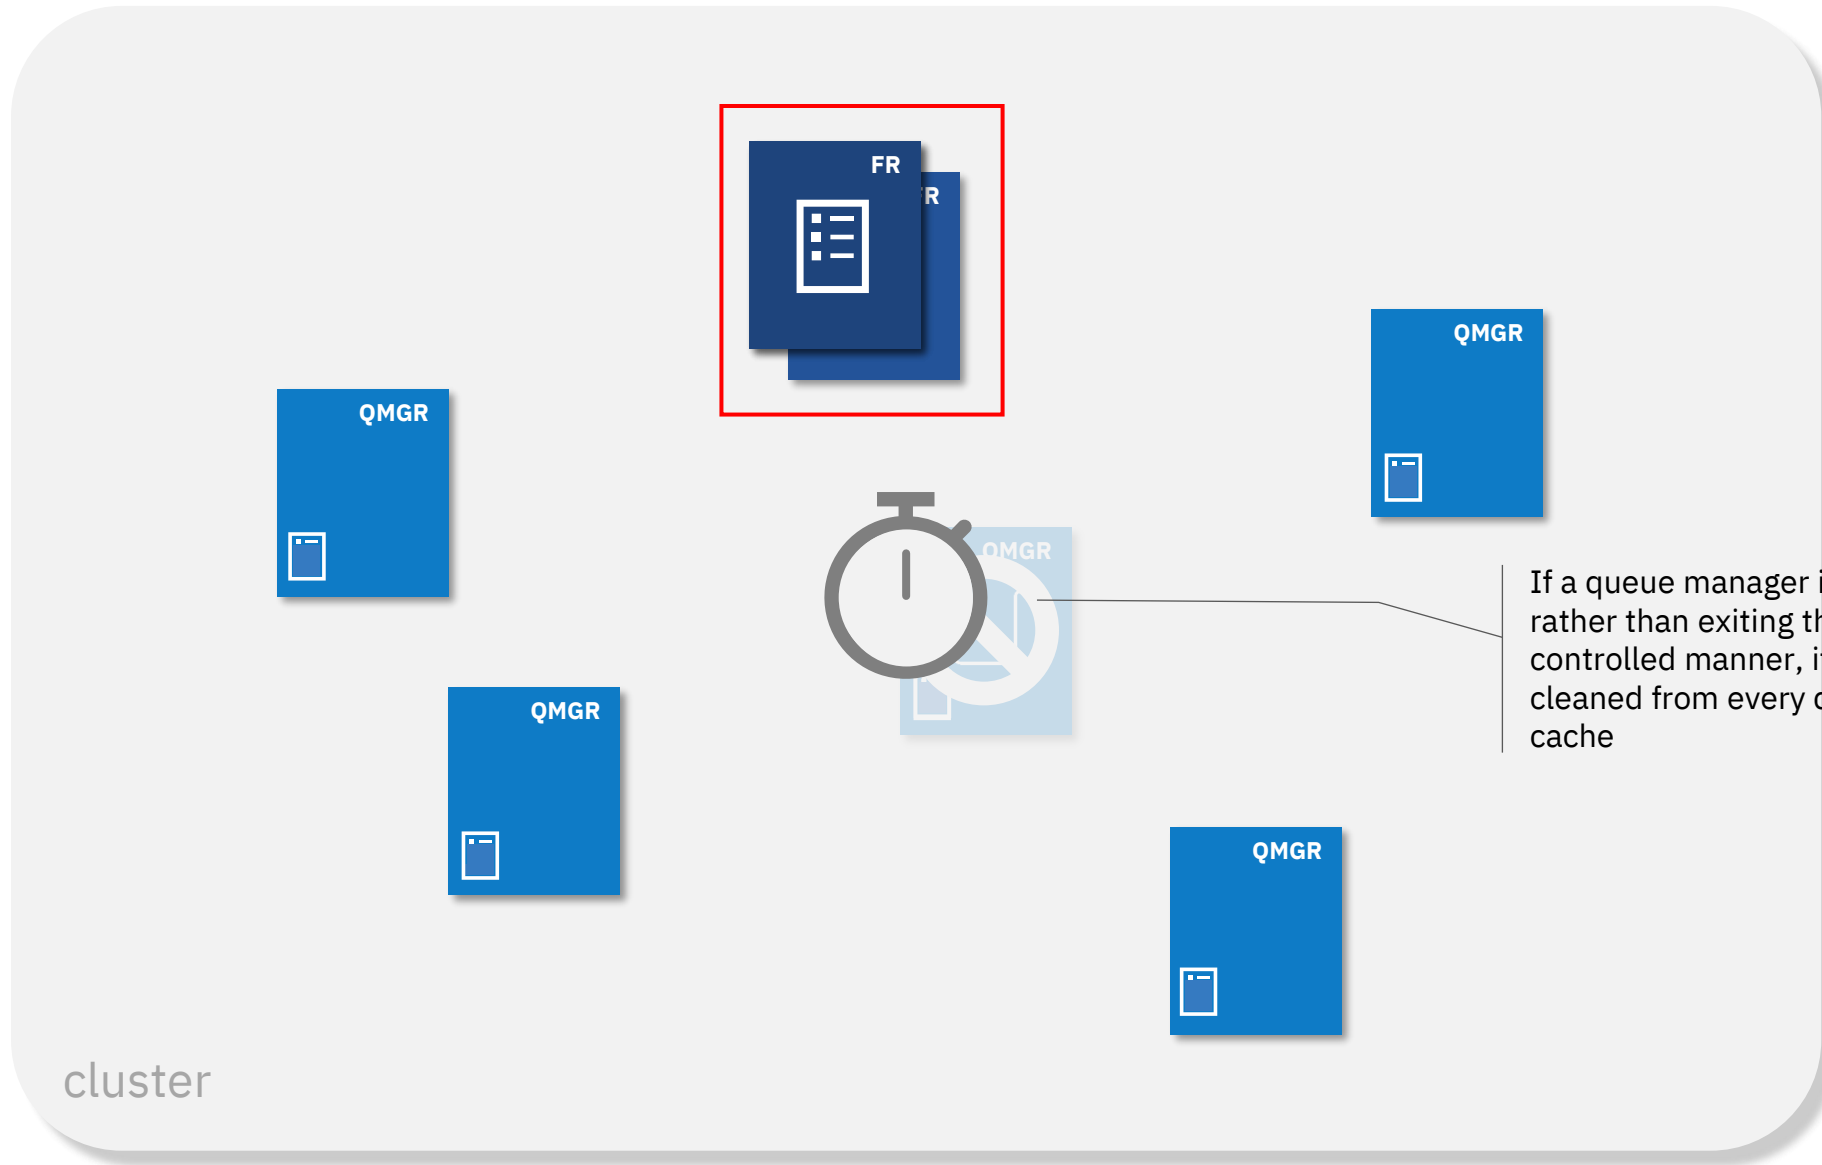
cluster



Cluster overview




Cluster overview



How to create a cluster

Step 1: Create your two full repositories



```
+ZQS1 ALTER QMGR REPOS('ZQCLUSTER')
```

```
+ZQS1 DEFINE CHANNEL('TO.ZQS1') CHLTYPE(CLUSRCVR)  
CLUSTER('ZQCLUSTER') CONNAME(FR1 location) TRPTYPE(TCP)
```

```
+ZQS1 DEFINE CHANNEL('TO.ZQS2') CHLTYPE(CLUSSDR)  
CLUSTER('ZQCLUSTER') CONNAME(FR2 location) TRPTYPE(TCP)
```

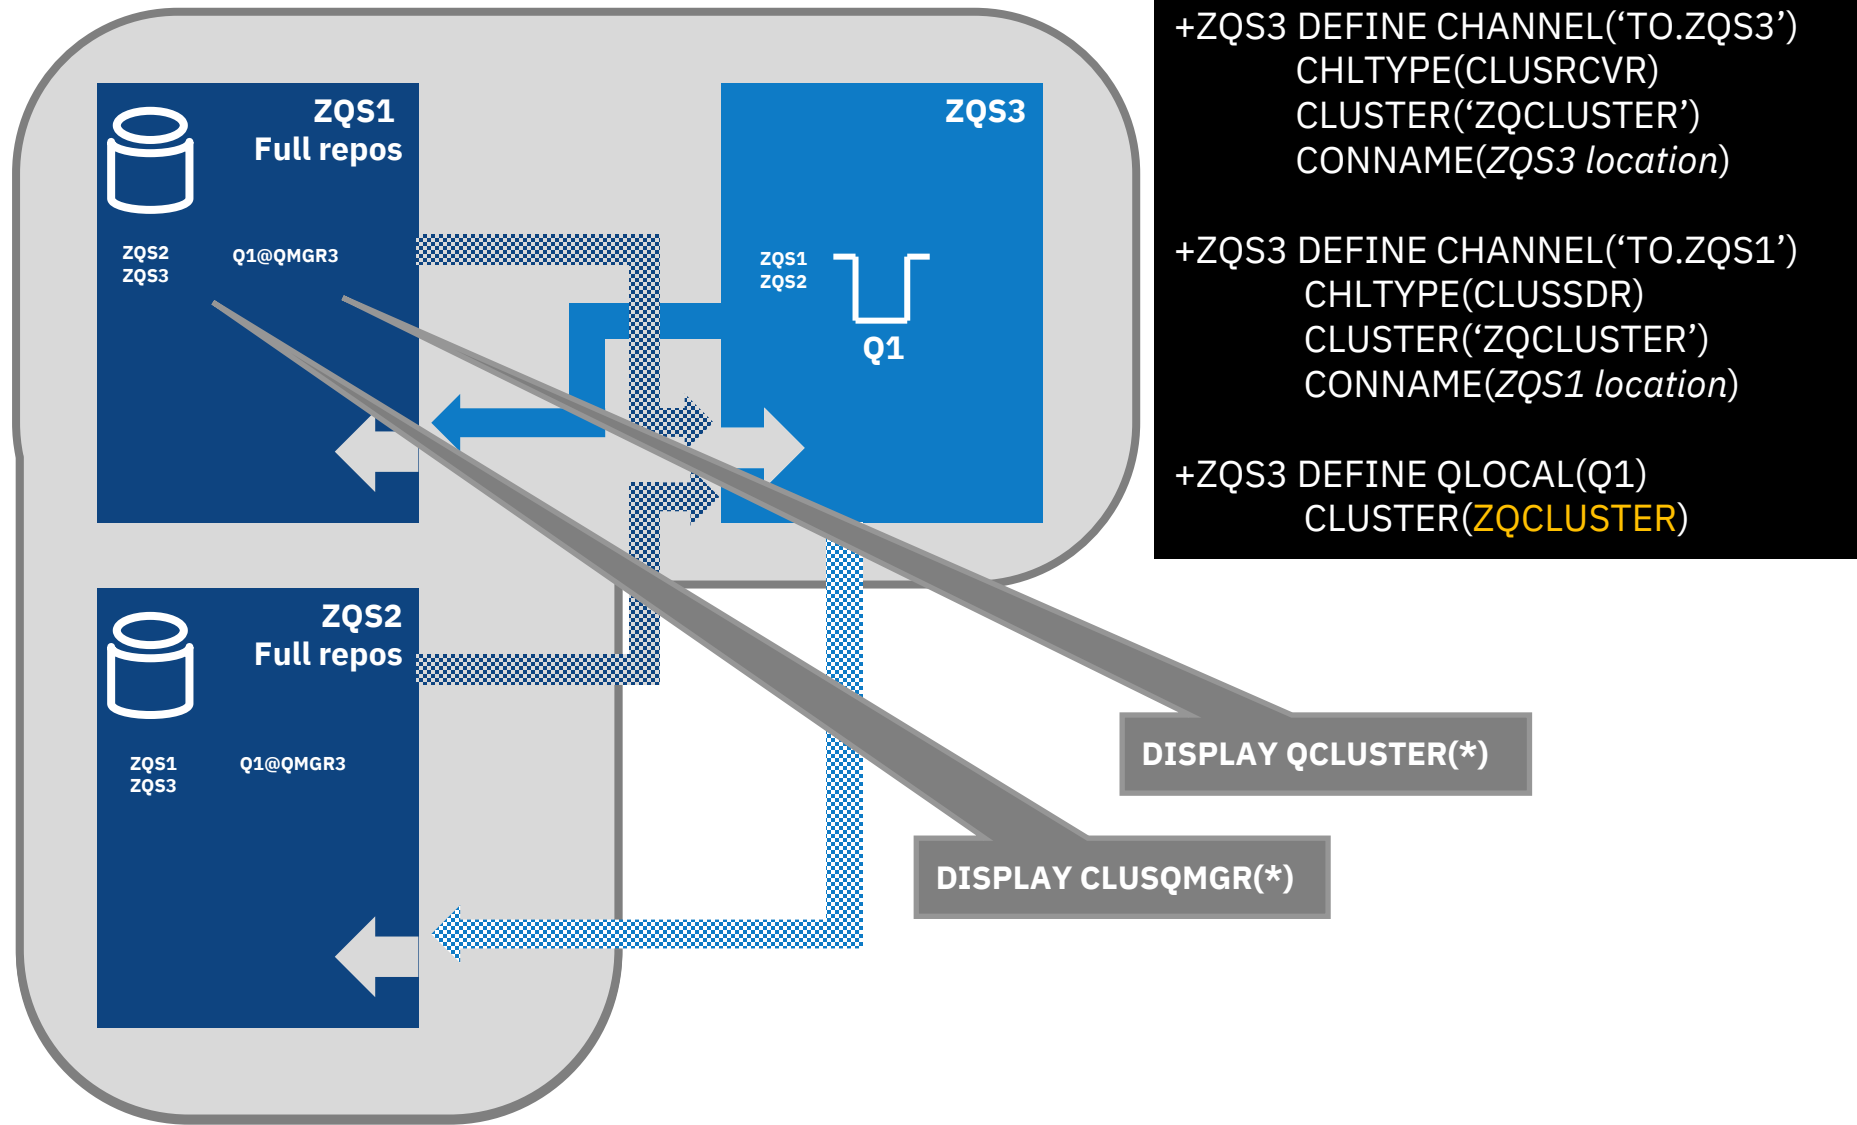
```
+ZQS2 ALTER QMGR REPOS('ZQCLUSTER')
```

```
+ZQS2 DEFINE CHANNEL('TO.ZQS2') CHLTYPE(CLUSRCVR)  
CLUSTER('ZQCLUSTER') CONNAME(FR2 location) TRPTYPE(TCP)
```

```
+ZQS2 DEFINE CHANNEL('TO.ZQS1') CHLTYPE(CLUSSDR)  
CLUSTER('ZQCLUSTER') CONNAME(FR1 location) TRPTYPE(TCP)
```

*This object creation is handled by MQ Explorer tooling

Step 2: Add in more queue managers



```

ZQS1 DISPLAY CLUSQMGR(*)
CSQM293I ZQS1 CSQMDRTC 2 CLUSQMGR FOUND MATCHING REQUEST CRITERIA
CSQM201I ZQS1 CSQMDRTC DISPLAY CLUSQMGR DETAILS 769
CLUSQMGR(ZQS1)
  CLUSTER(ZQCLUSTER)
  CHANNEL(TO.ZQS1)
END CLUSQMGR DETAILS
CSQM201I ZQS1 CSQMDRTC DISPLAY CLUSQMGR DETAILS 770
CLUSQMGR(ZQS2)
  CLUSTER(ZQCLUSTER)
  CHANNEL(TO.ZQS2)
END CLUSQMGR DETAILS

```

information about remote queue managers **and** the channel definitions by which they can be reached


```

ZQS1 DISPLAY QCLUSTER(*)
CSQM293I ZQS1 CSQMDRTC 1 QCLUSTER FOUND MATCHING REQUEST CRITERIA
CSQM201I ZQS1 CSQMDRTC DISPLAY QCLUSTER DETAILS 811
QUEUE(Q1)
  TYPE(QCLUSTER)
  QSGDISP(QMGR)
END QCLUSTER DETAILS
CSQ9022I ZQS1 CSQMDRTC ' DISPLAY QCLUSTER' NORMAL COMPLETION

```

information about the queue and topic objects that we can see via a cluster

Step 2: Add in more queue managers



```
+ZQS3 DEFINE CHANNEL('ZQCLUSTER.ZQS3')  
  CHLTYPE(CLUSRCVR)  
  CLUSTER('ZQCLUSTER')  
  CONNAME(QMGR1 location)
```

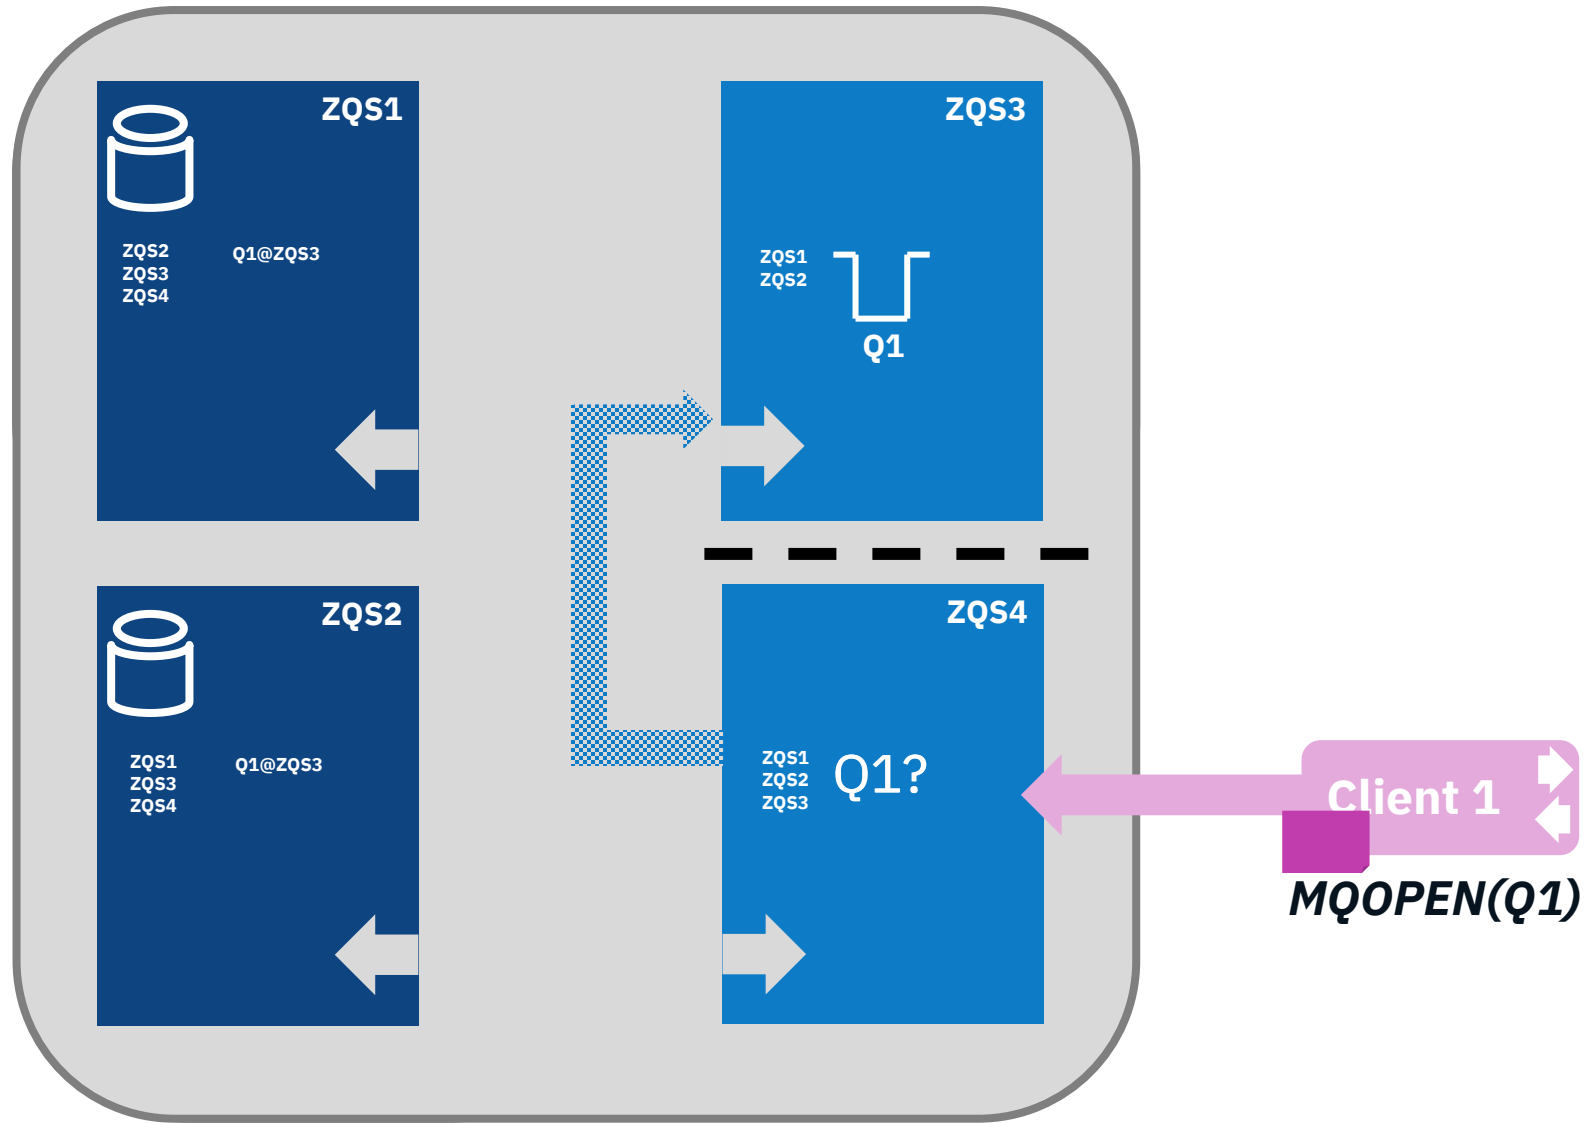
```
+ZQS3 DEFINE CHANNEL('ZQCLUSTER.ZQS1')  
  CHLTYPE(CLUSSDR)  
  CLUSTER('ZQCLUSTER')  
  CONNAME(FR1 location)
```

```
+ZQS3 DEFINE QLOCAL(Q1)  
  CLUSTER(ZQCLUSTER)
```


```
+ZQS4 DEFINE CHANNEL('ZQCLUSTER.ZQS4')  
  CHLTYPE(CLUSRCVR)  
  CLUSTER('ZQCLUSTER')  
  CONNAME(ZQS4 location)
```

```
+ZQS4 DEFINE CHANNEL('ZQCLUSTER.ZQS1')  
  CHLTYPE(CLUSSDR)  
  CLUSTER('ZQCLUSTER')  
  CONNAME(ZQS1 location)
```


Step 3: Start sending messages



So all you needed...




- Two full repository queue managers
- A cluster receiver channel each
- A single cluster sender each
- No need to manage pairs of channels between each queue manager combination or their transmission queues
- No need for remote queue definitions

Shared Cluster Queues


Channel workload balancing

- Cluster workload balancing applies when there are multiple cluster queues of the same name
- Cluster workload balancing will be applied in one of three ways:
 - When the putting application opens the queue
 - bind on open
 - When a message group is started
 - bind on group
 - When a message is put to the queue
 - bind not fixed
- When workload balancing is applied:
 - The source queue manager builds a list of all potential targets based on the queue name
 - Eliminates the impossible options
 - Prioritises the remainder
 - If more than one come out equal, workload balancing ensues ...
- Balancing is based on:
 - The **channel** – not the target queue
 - Channel traffic to all queues is taken into account
 - Weightings can be applied to the channel
- ... this is used to send the messages to the chosen target




Channel workload balancing

- Cluster workload balancing applies when there are **multiple cluster queues of the same name**
- Cluster workload balancing will be applied in **one of three ways**:
 - When the putting application opens the queue - **bind on open**
 - When a message group is started - **bind on group**
 - When a message is put to the queue - **bind not fixed**
- When workload balancing is applied:
 - The source queue manager builds a list of all potential targets based on the queue name
 - **Eliminates** the impossible options
 - **Prioritises** the remainder
 - If more than one come out equal, **workload balancing** ensues ...
- Balancing is based on:
 - The **channel** – not the target queue
 - **Channel traffic to all queues** is taken into account
 - **Weightings** can be applied to the channel
- ... this is used to send the messages to the chosen target



Channel workload balancing


- Cluster workload balancing applies when there are **multiple cluster queues of the same name**
- Cluster workload balancing will be applied in **one of three ways**:
 - When the putting application opens the queue - **bind on open**
 - When a message group is started - **bind on group**
 - When a message is put to the queue - **bind not fixed**
- When workload balancing is applied:
 - The source queue manager builds a list of all potential targets based on the queue name
 - **Eliminates** the impossible options
 - **Prioritises** the remainder
 - If more than one come out equal, **workload balancing** ensues ...
- Balancing is based on:
 - The **channel** – not the target queue
 - **Channel traffic to all queues** is taken into account
 - **Weightings** can be applied to the channel
- ... this is used to send the messages to the chosen target



Availability routing

Clustering for availability

Is MQ Clustering a high availability solution?



- Having multiple potential targets for any message can improve the availability of the solution, always providing an option to process new messages.
- A queue manager in a cluster has the ability to route new and old messages based on the **availability of the channels**, routing messages to running queue managers.
- Clustering can be used to route messages to active consuming applications.


Channel availability routing

- When performing workload balancing, the availability of the channel to reach the target is a factor
- All things being equal, messages will be routed to those targets with a working channel

Routing of messages based on availability doesn't just happen when they're first put, it also occurs for queued transmission messages every time the channel is retried. So blocked messages can be re-routed, if they're not prevented...

Things that can prevent routing

- Applications targeting messages at a specific queue manager
- Using "cluster workload rank"
- Binding messages to a target



Pros and cons of binding

Bind on open

Bind on group

Bind context:

Duration of an *open*

Duration of *logical group*

- All messages put within the bind context will go to same target*
- Message order can be preserved**
- Workload balancing logic is only driven at the start of the context
- Once a target has been chosen it **cannot change**
 - Whether it's available or not
 - Even if all the messages could be redirected

Bind not fixed

Bind context:

None

- Greater *availability*, a message will be redirected to an available target***
- Overhead of workload balancing logic for every message
- Message order may be affected

Bind on open is the default

It could be set on the cluster queue (don't forget aliases) or in the app

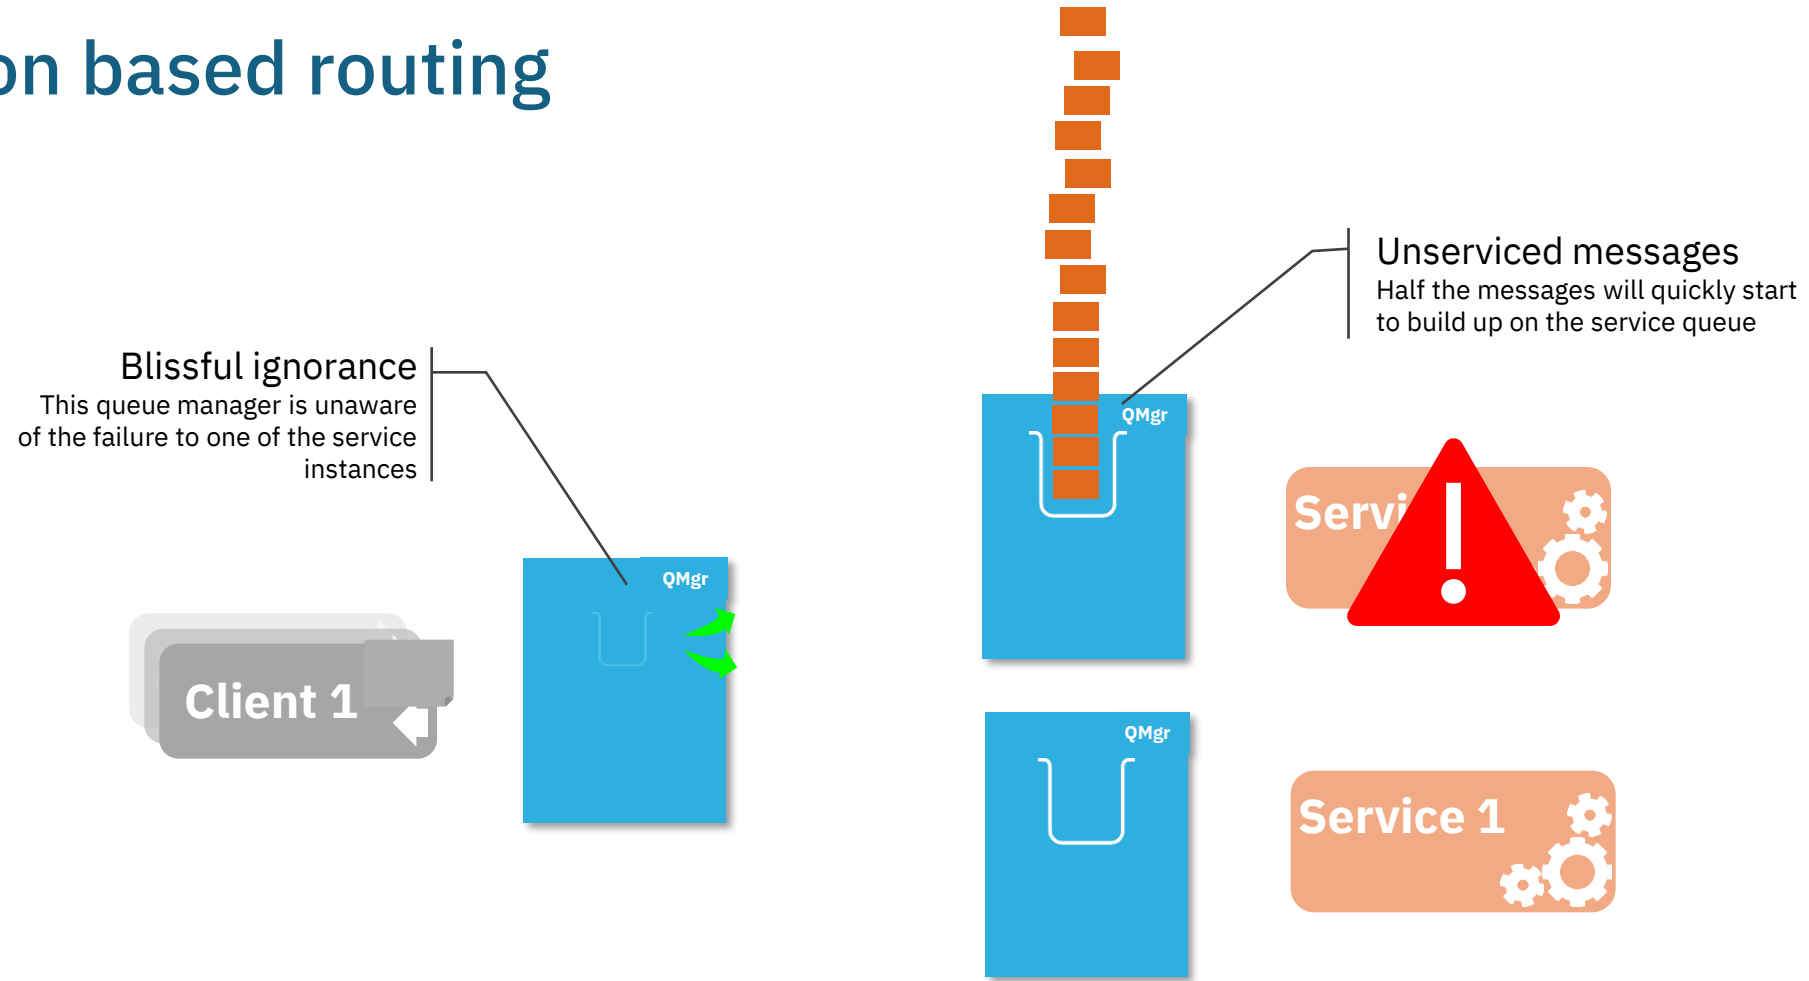
* While a route is known by the source queue manager, it won't be rebalanced, but it could be DLQd

** Other aspects may affect ordering (e.g. deadletter queueing)

*** Unless it's fixed for another reason (e.g. specifying a target queue manager)


Application availability routing

Application based routing

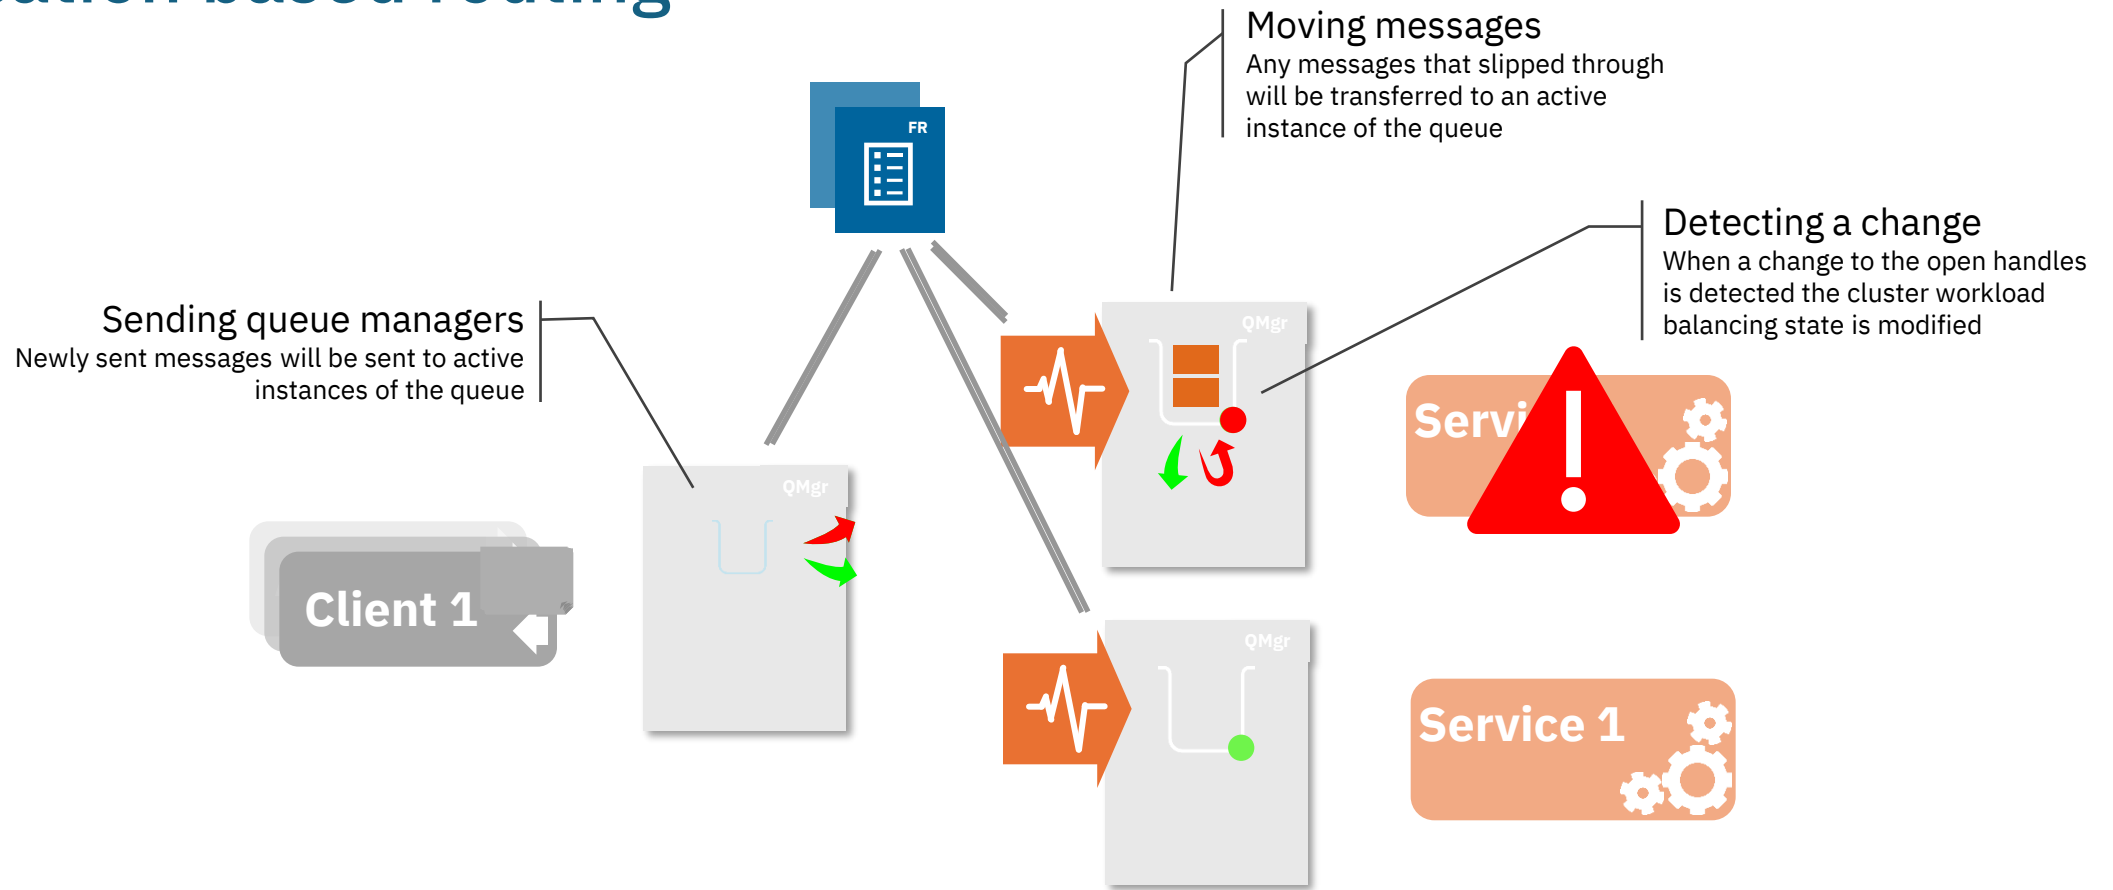


- Cluster workload balancing does not take into account the availability of receiving applications
- Or a build up of messages on a queue

Application based routing



Application based routing



- MQ provides a sample monitoring service tool, **amqsclm**
- It regularly checks for attached consuming applications (IPPROCS)
- And automatically adjusts the cluster queue definitions to route messages intelligently (CLWLPRTY)
- That information is automatically distributed around the cluster

Thank You !

Any Questions Please ?



Your feedback is important!

Submit a session evaluation for each session you attend:

www.share.org/evaluation

