

IBM MQ for z/OS SMF Skills Workshop



Dorothy Quincy
Washington Systems Center

IBM MQ for z/OS Client
Technical Specialist



Today's agenda

Introduction to Queue Manager Internals (45 min)

Walk through the internal functionality of a queue manager to understand the internal resource threads and storage facilities underpinning each queue manager.

Introduction to Interpreting SMF data for MQ on z/OS (1 hour)

Recognize key metrics and the associated impact in SMF 115 and 116 data to identify where tuning can occur to improve MQ's performance.

Demonstration of SMF processing (45 minutes)

Demonstration of the process IBM uses evaluate customer environments for IBM MQ health checks. IBM will demonstrate how the data gets from z/OS to readable and interpretable spreadsheets.

My goals for today

Make it clear how IBM processes your SMF data to make recommendations for your MQ environment during our health checks

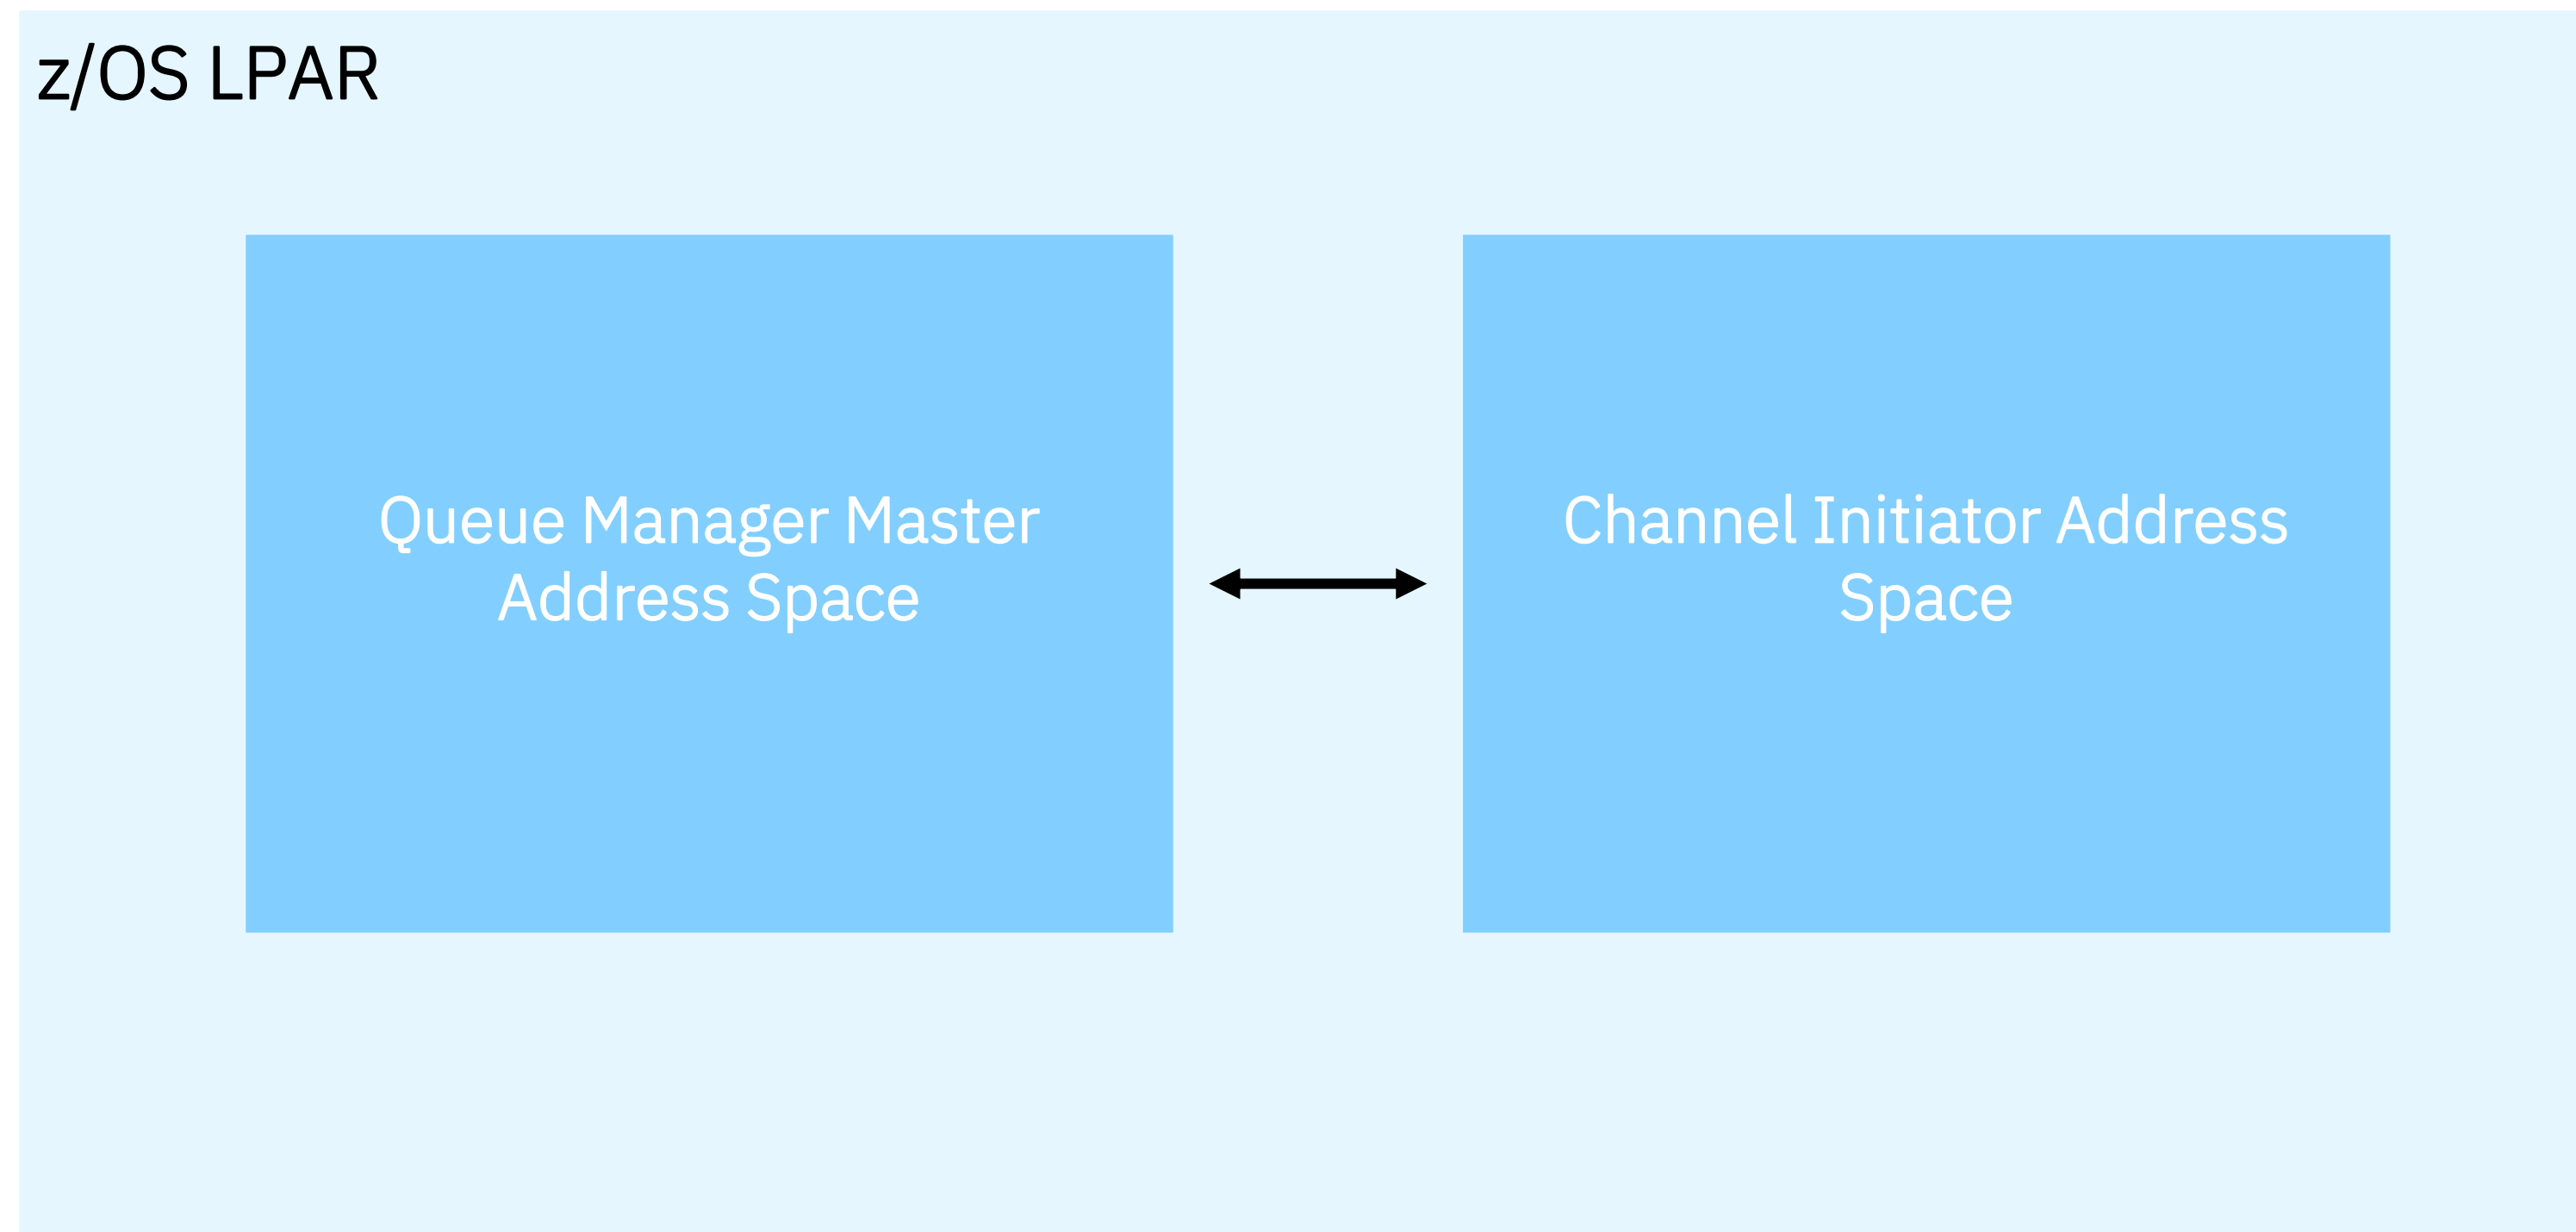
Give you a sense of the KPI's to pay attention to when evaluating SMF data for your environment

Show you various options for how to view SMF data


Connect the SMF data and its KPI's back to what is actually going on inside a queue manager running on IBM MQ for z/OS

Introduction to Queue Manager Internals


At a glance



Diving deeper: Private queues



Diving deeper: Private queues



Digging into local queuing

Message A

```
Sent message:

JMSMessage class: jms_text
JMSType:          null
JMSDeliveryMode:  2
JMSDeliveryDelay: 0
JMSDeliveryTime:  1585562399950
JMSExpiration:    0
JMSPriority:       4
JMSMessageID:     ID:414d5120514d3120202020202020200ac2815e024ce120
JMSTimestamp:     1585562399950
JMSCorrelationID: null
JMSDestination:   queue:///DEV.QUEUE.1
JMSReplyTo:       null
JMSRedelivered:   false
  JMSXAppID: JmsPutGet (JMS)
  JMSXDeliveryCount: 0
  JMSXUserID: app
  JMS_IBM_PutApplType: 28
  JMS_IBM_PutDate: 20200330
  JMS_IBM_PutTime: 09595997
Your lucky number today is 926
```

Message details ×

Messages properties

Message ID	ID:414d5120514d3120202020202020200ac2815e024ce120
Timestamp	2022-2-18 16:37:23
Character set	UTF-8
Delivery mode	Persistent
Application ID	JmsPutGet (JMS)
Format	MQSTR
Expiration	0
Priority	4
Encoding	273
User ID	app

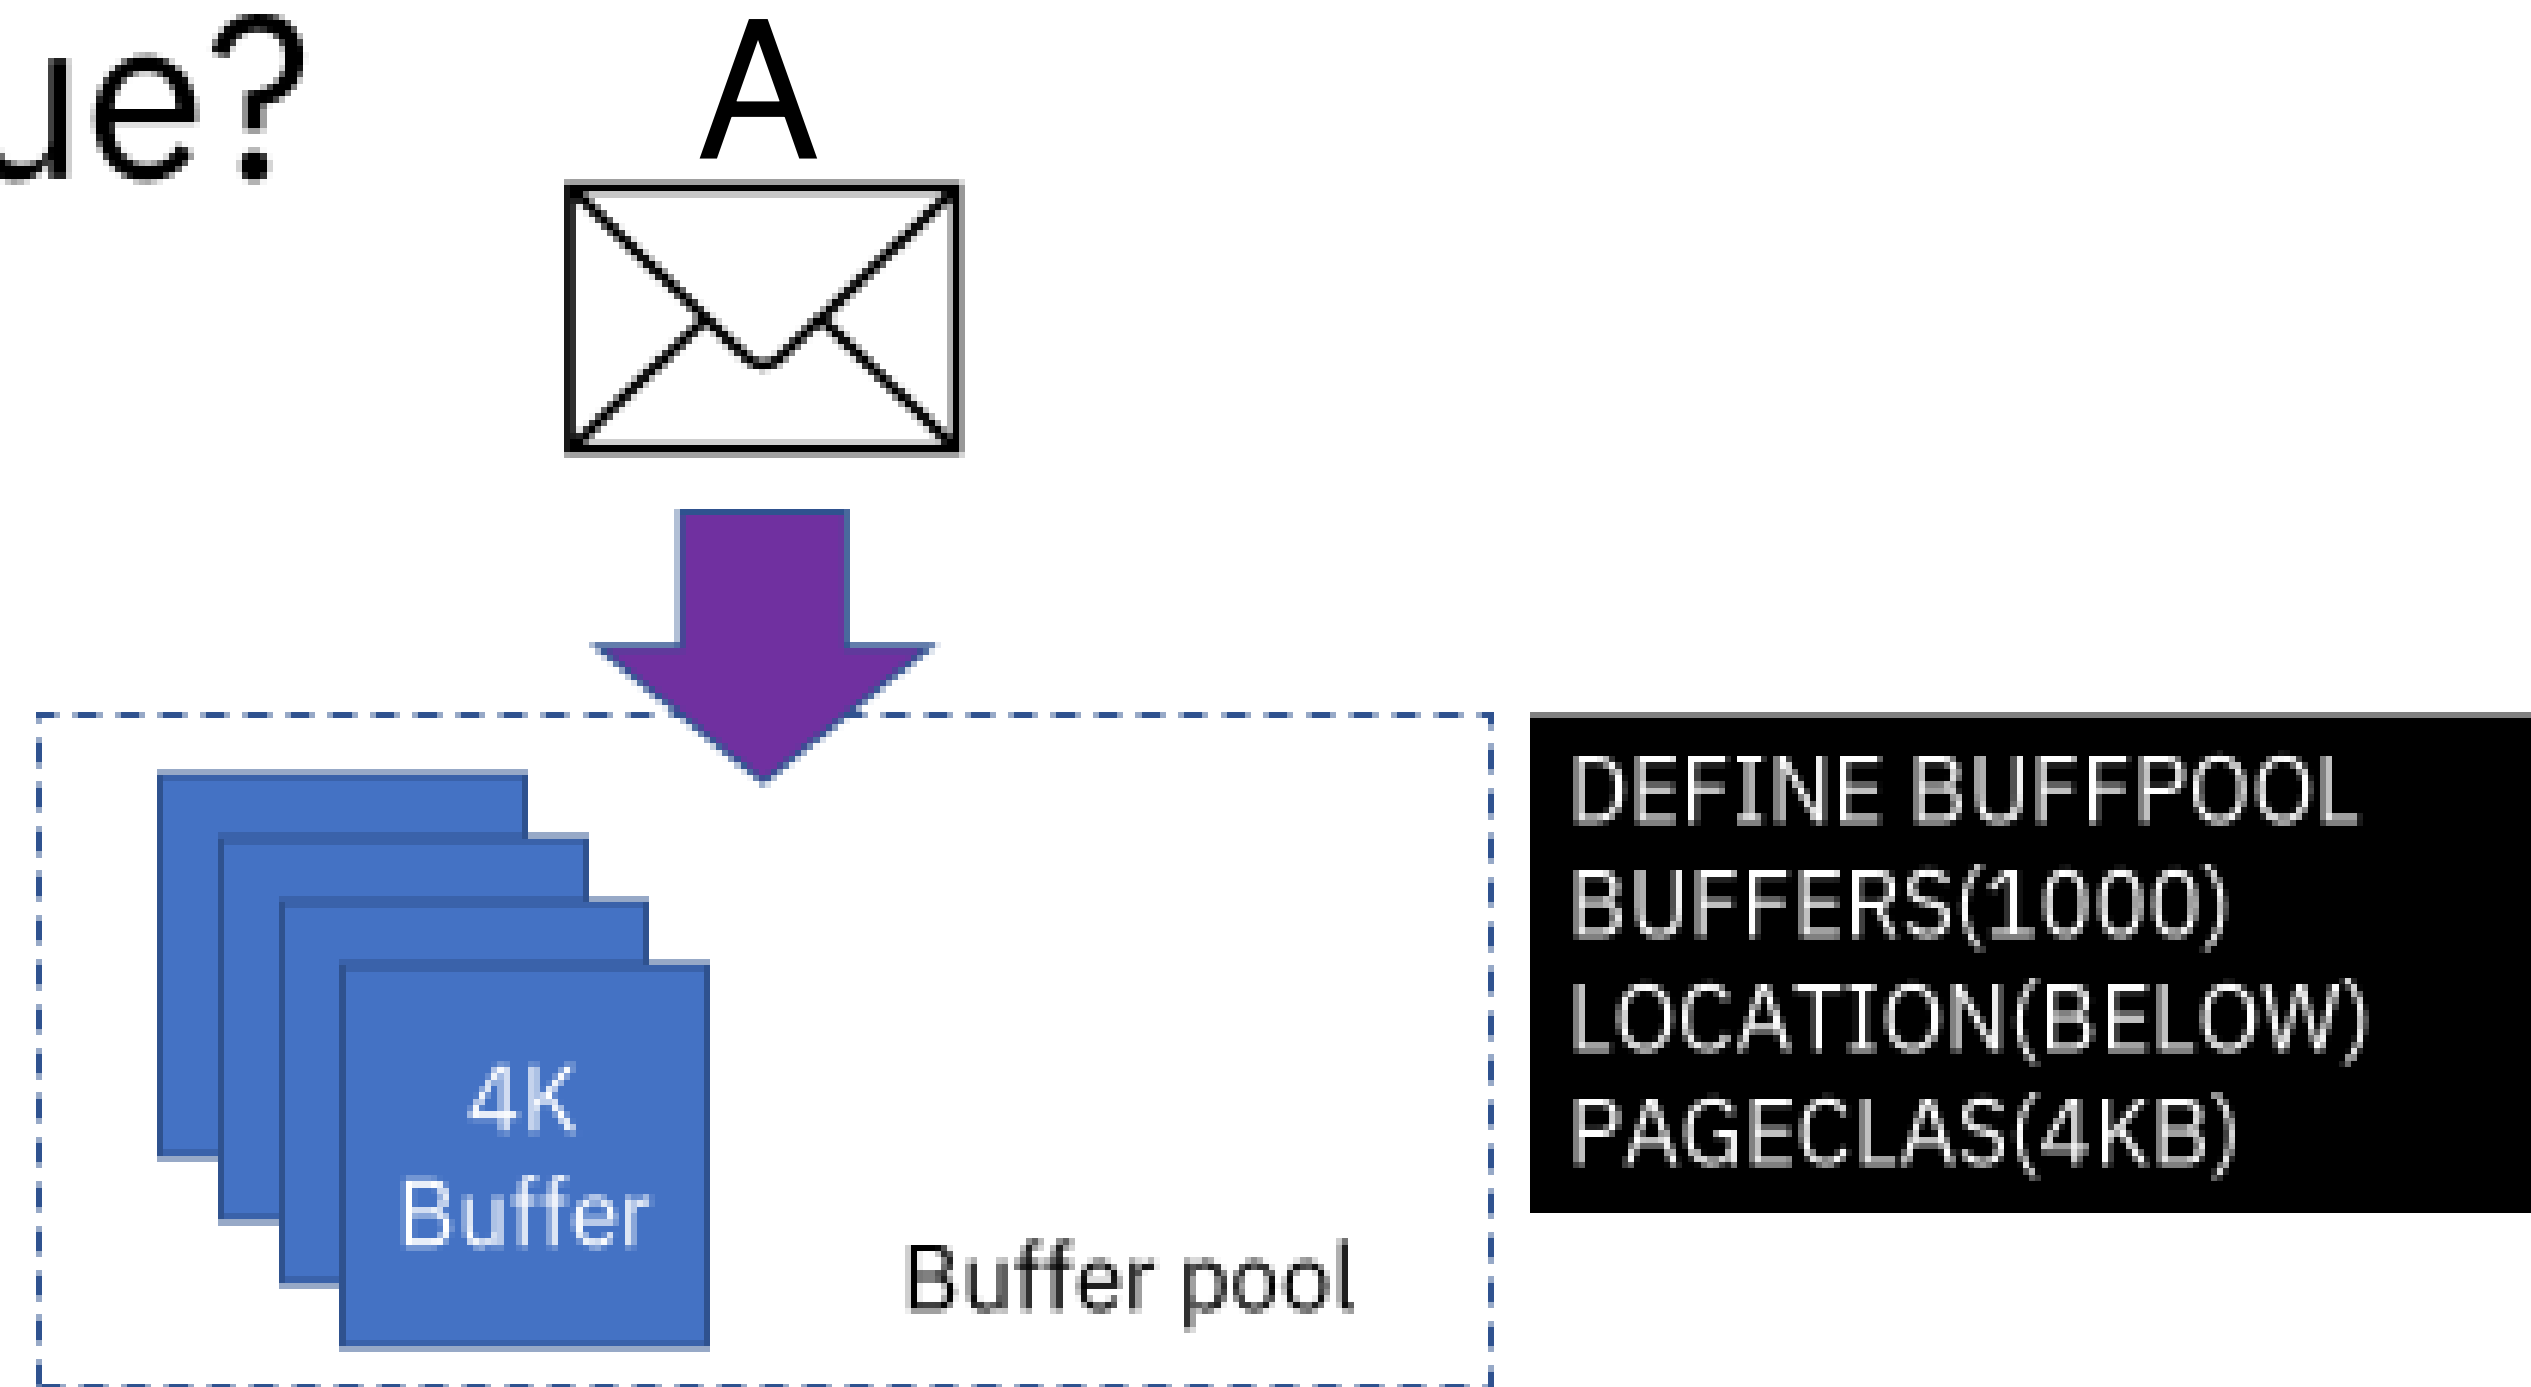
Application data

Your lucky number today is 369

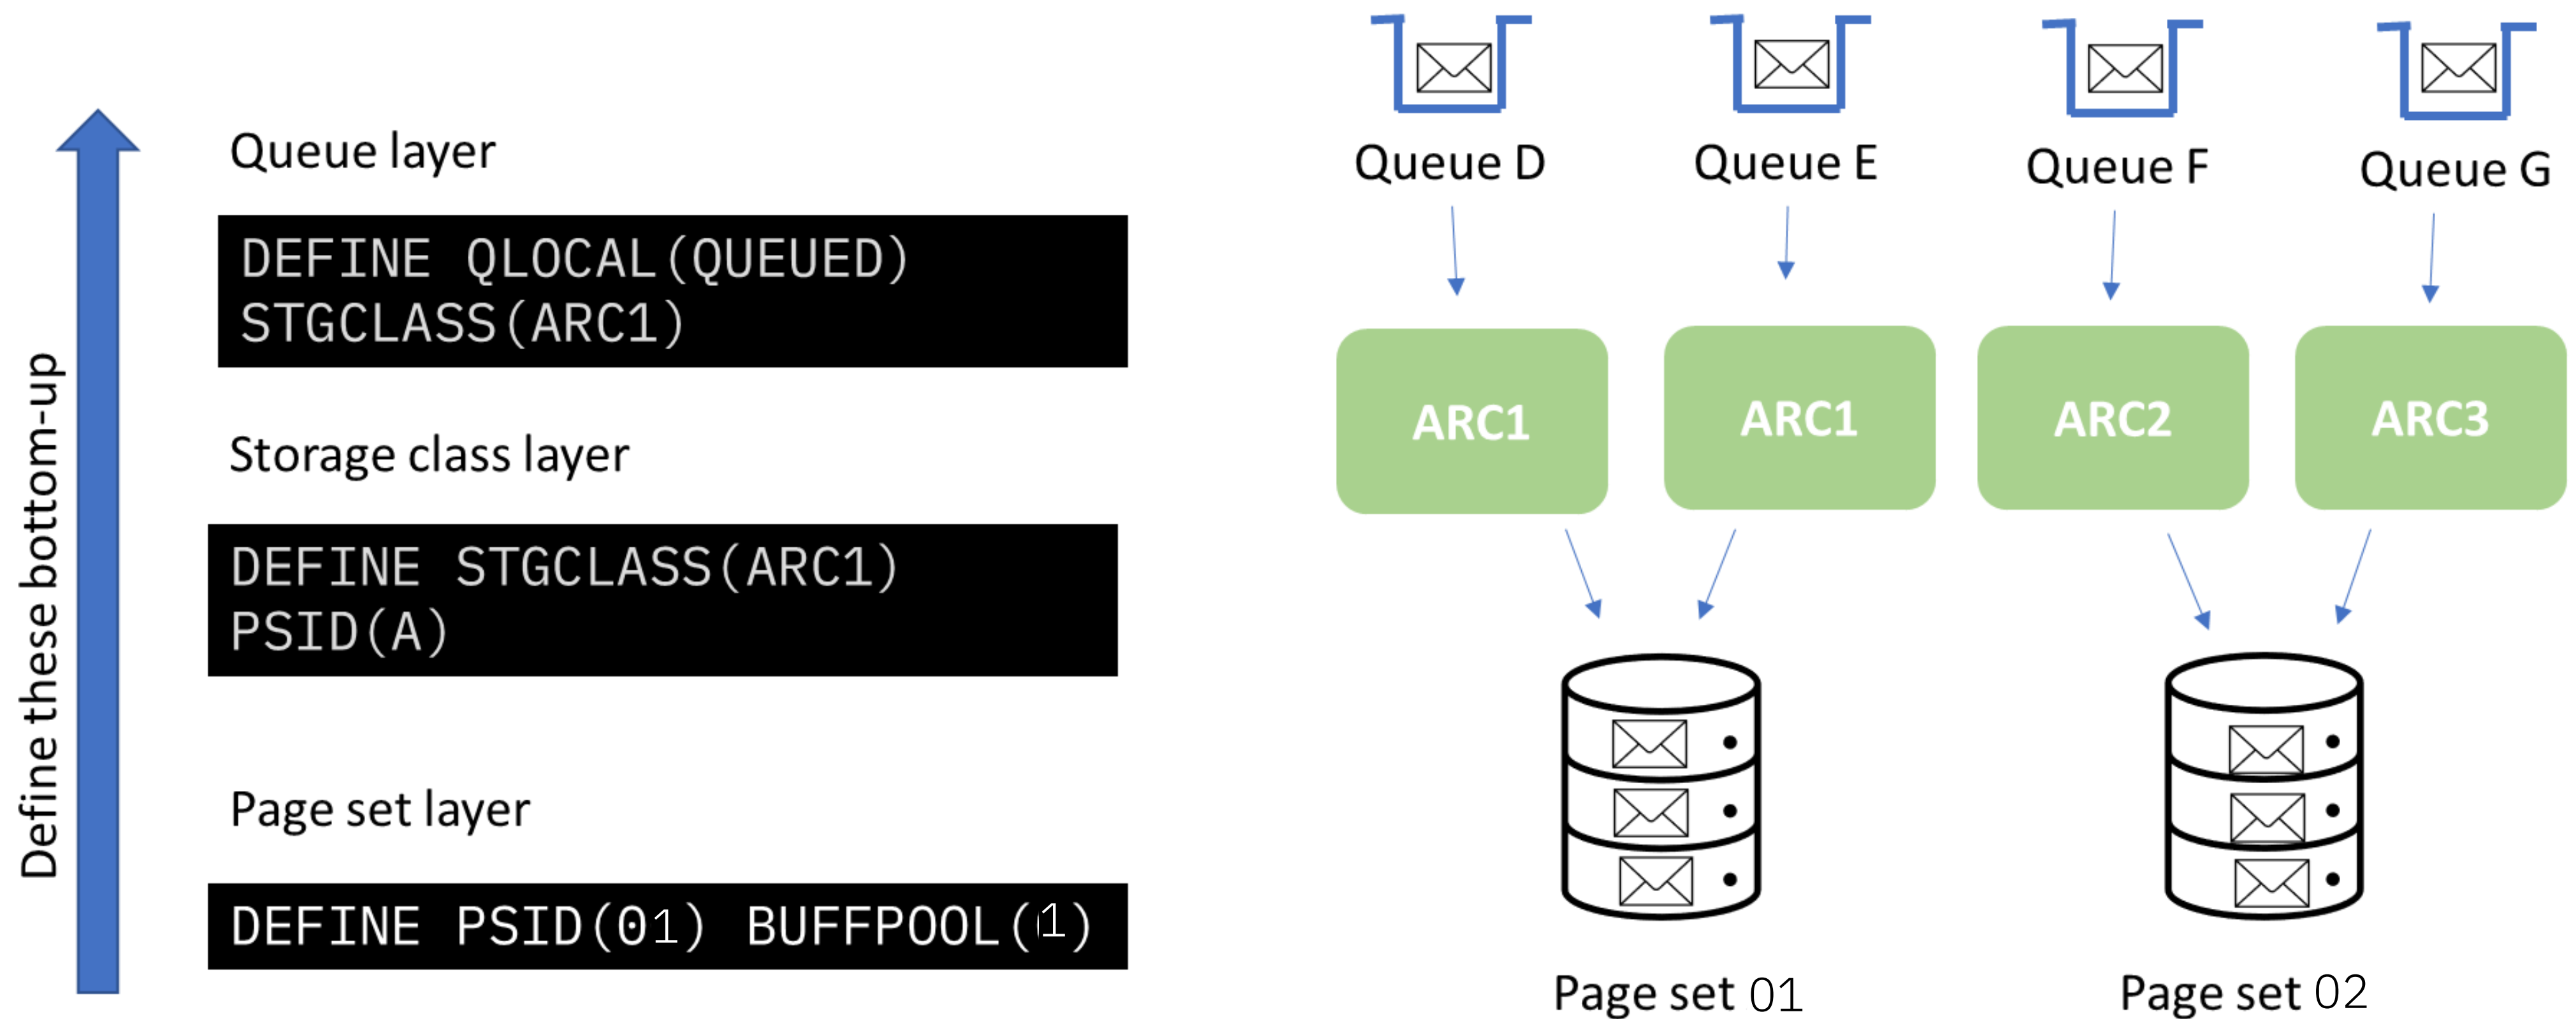
How does physical storage work on a private queue?

Messages are written to buffer pools when any of these conditions are met:


1. When messages have been in the buffer pool for 2 log checkpoints
2. When buffer pool usage exceeds the deferred write threshold
3. When buffer pool usage exceeds the buffer pool threshold



How are private queues associated with physical storage?

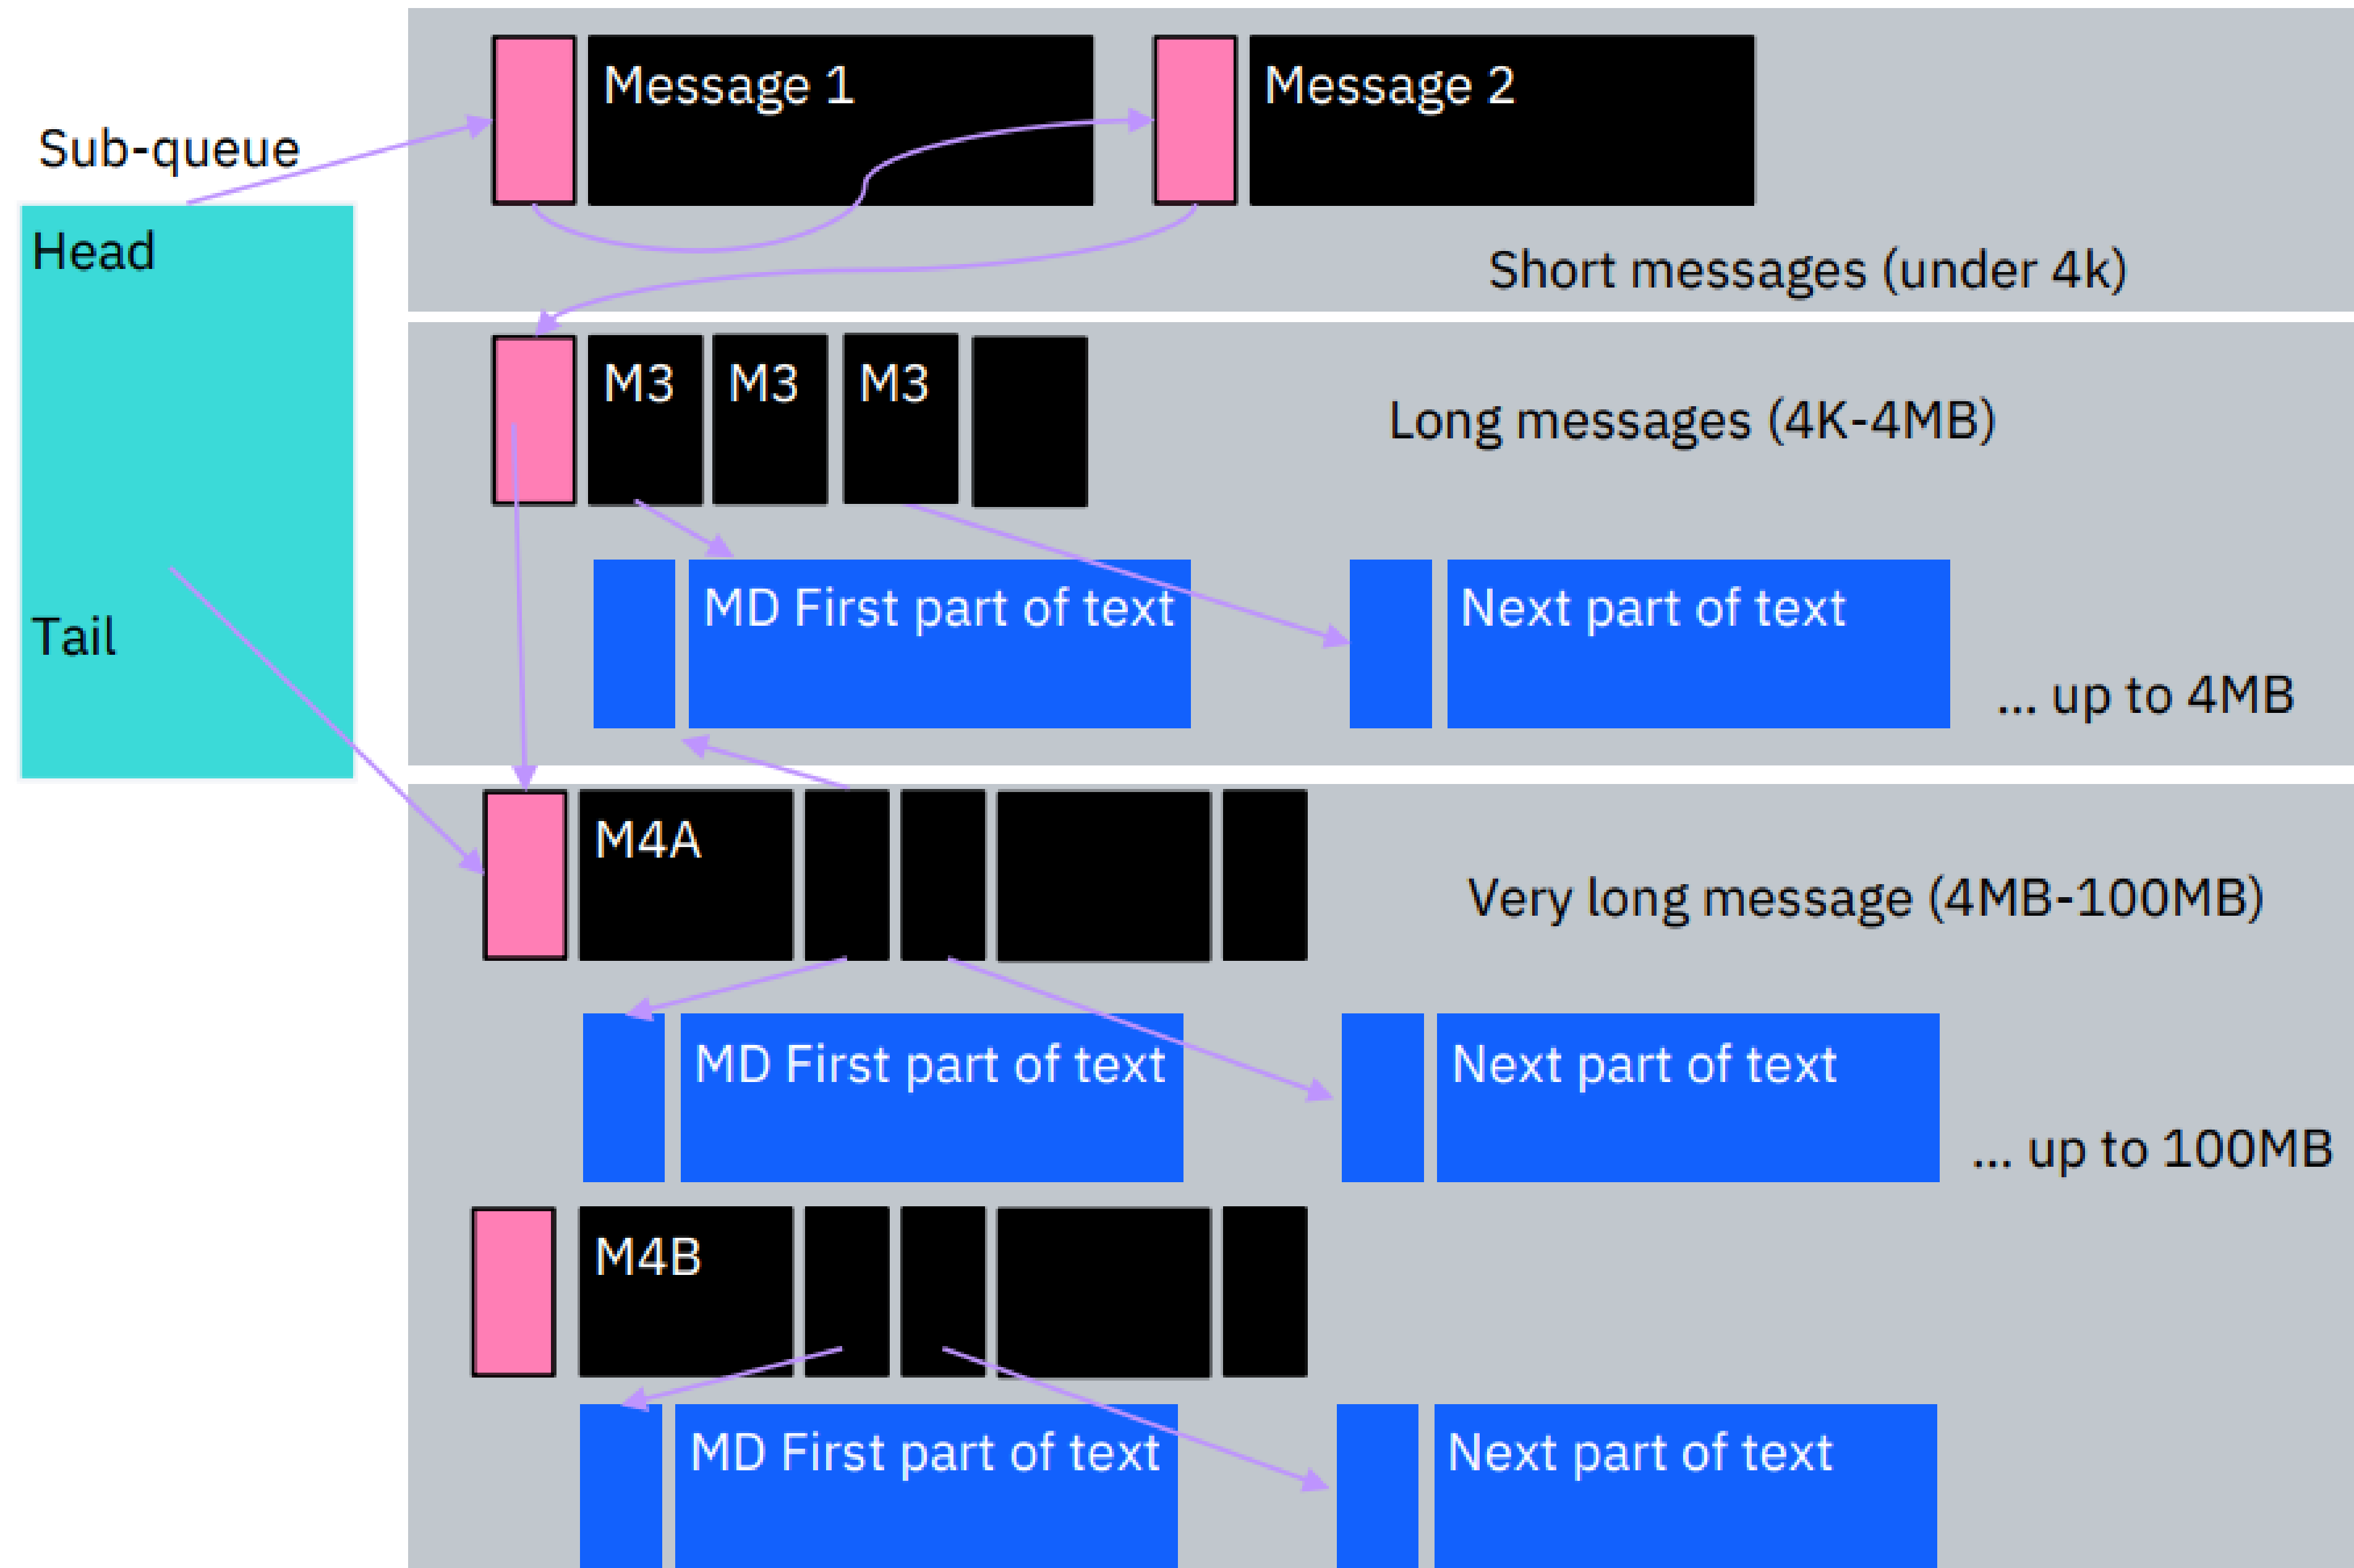


Internal Representation of a Private Queue

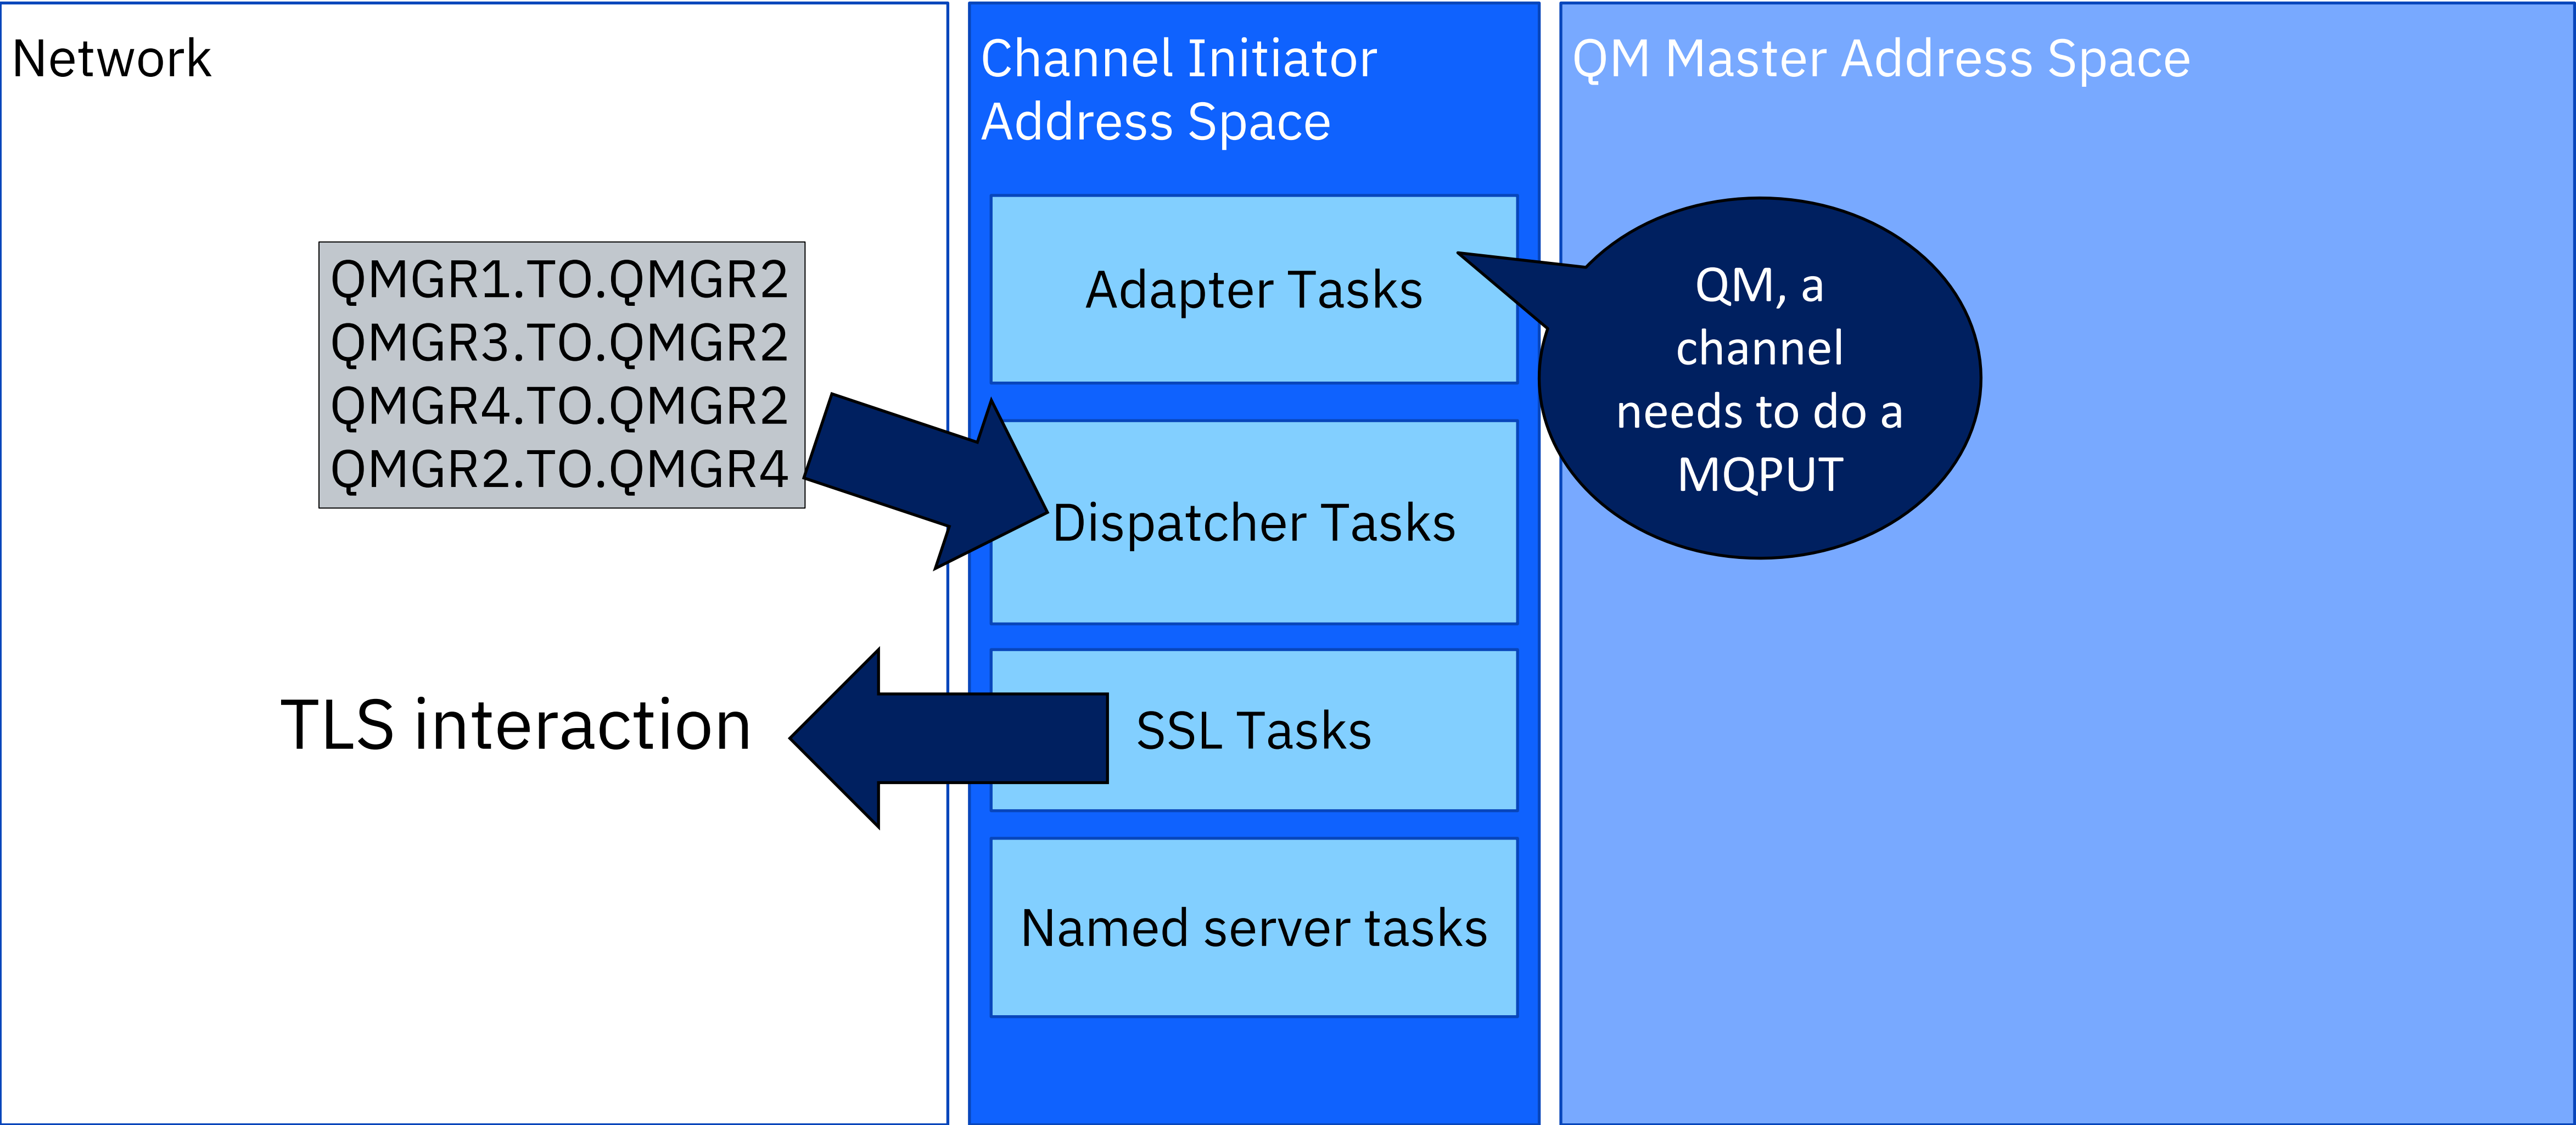


Sub-queue Internal View

in

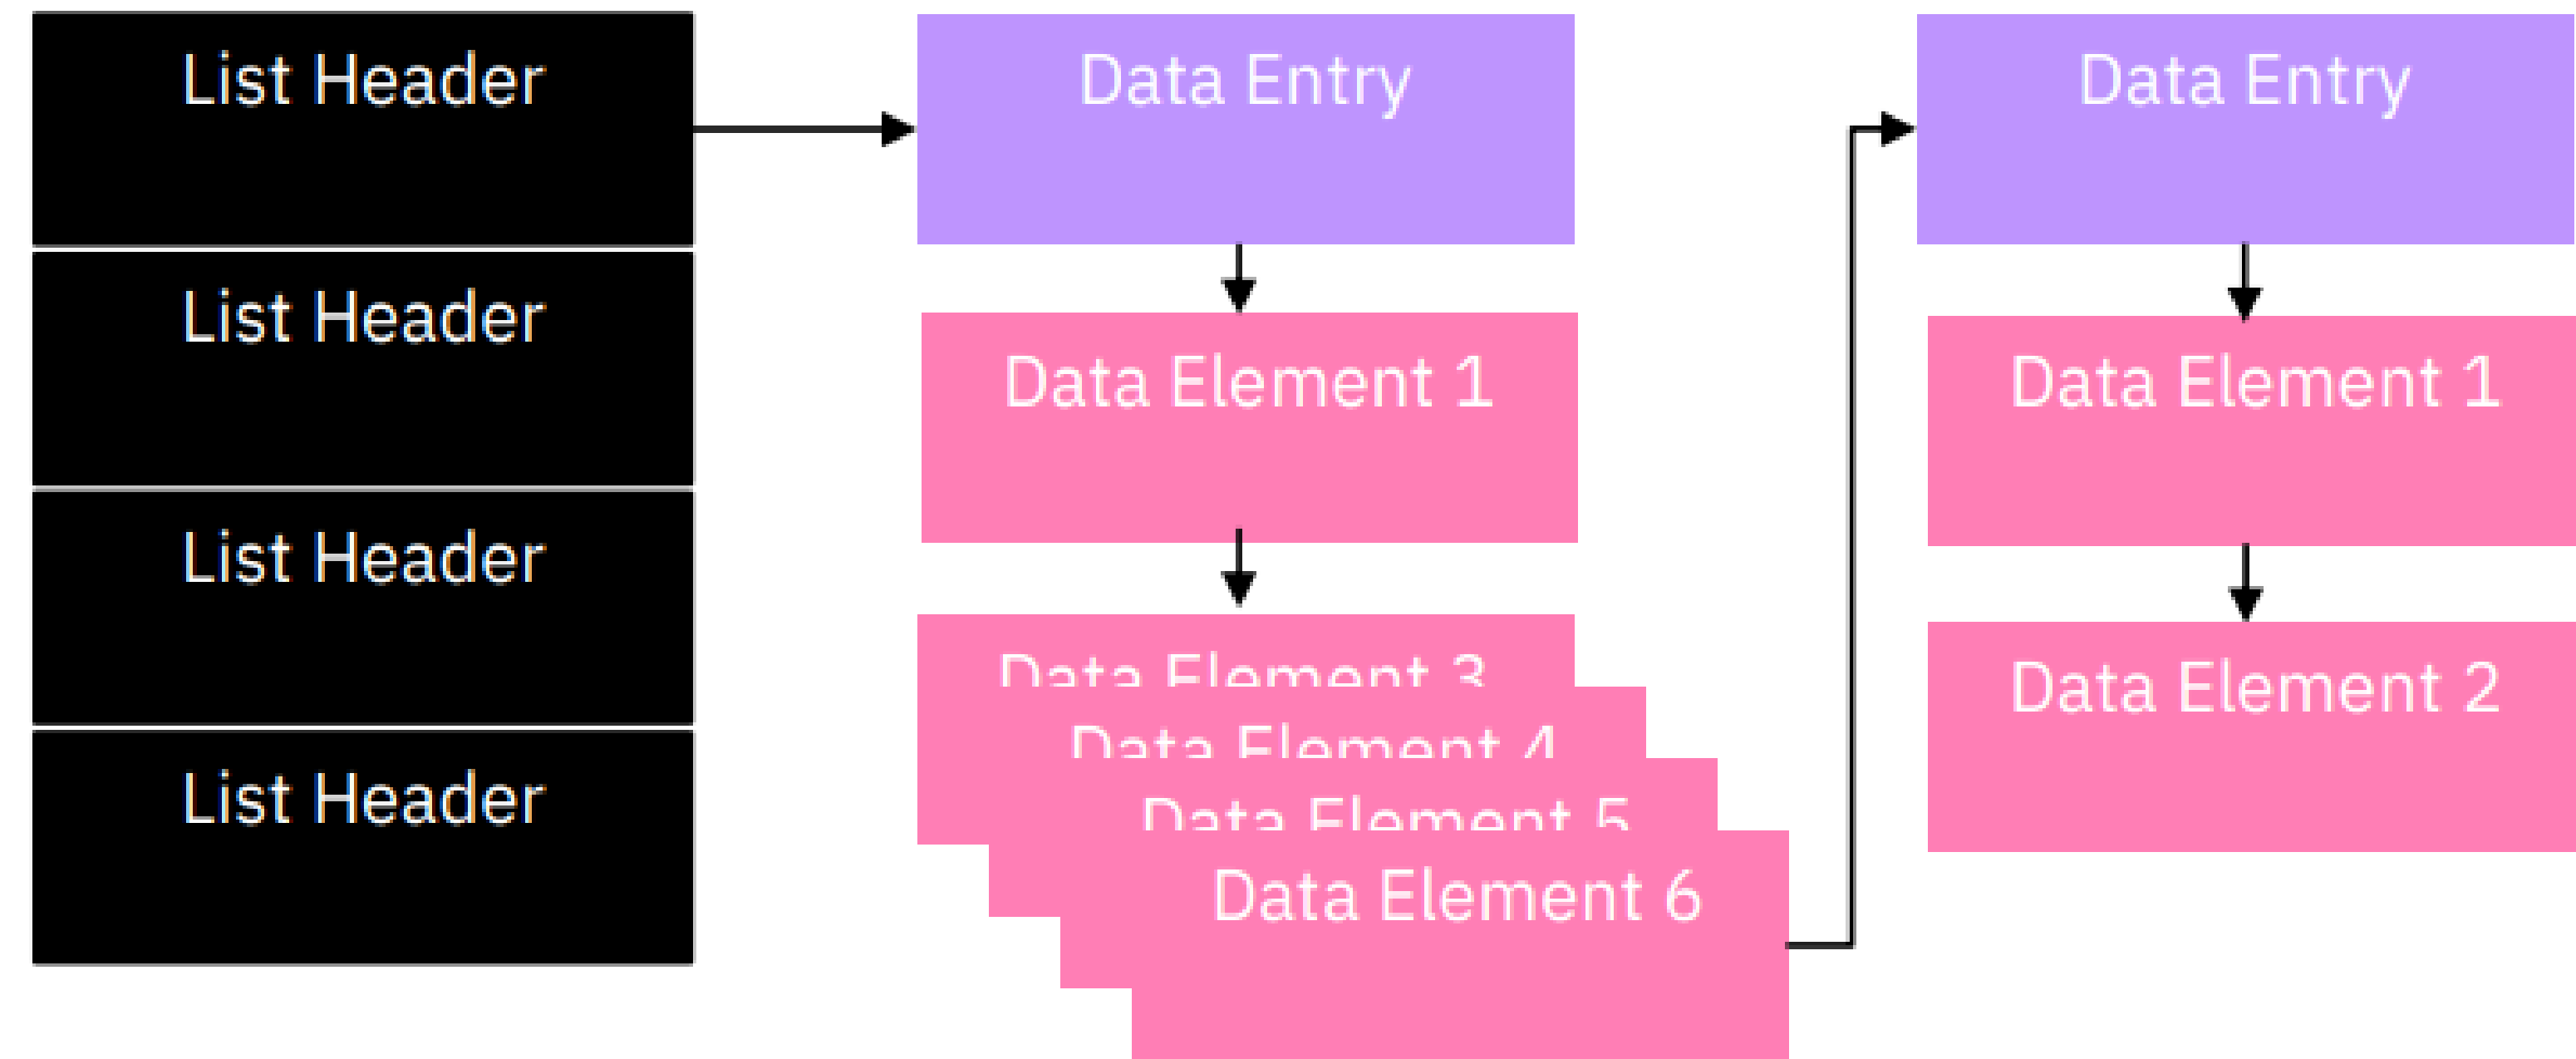


CHINIT Address Space Structure

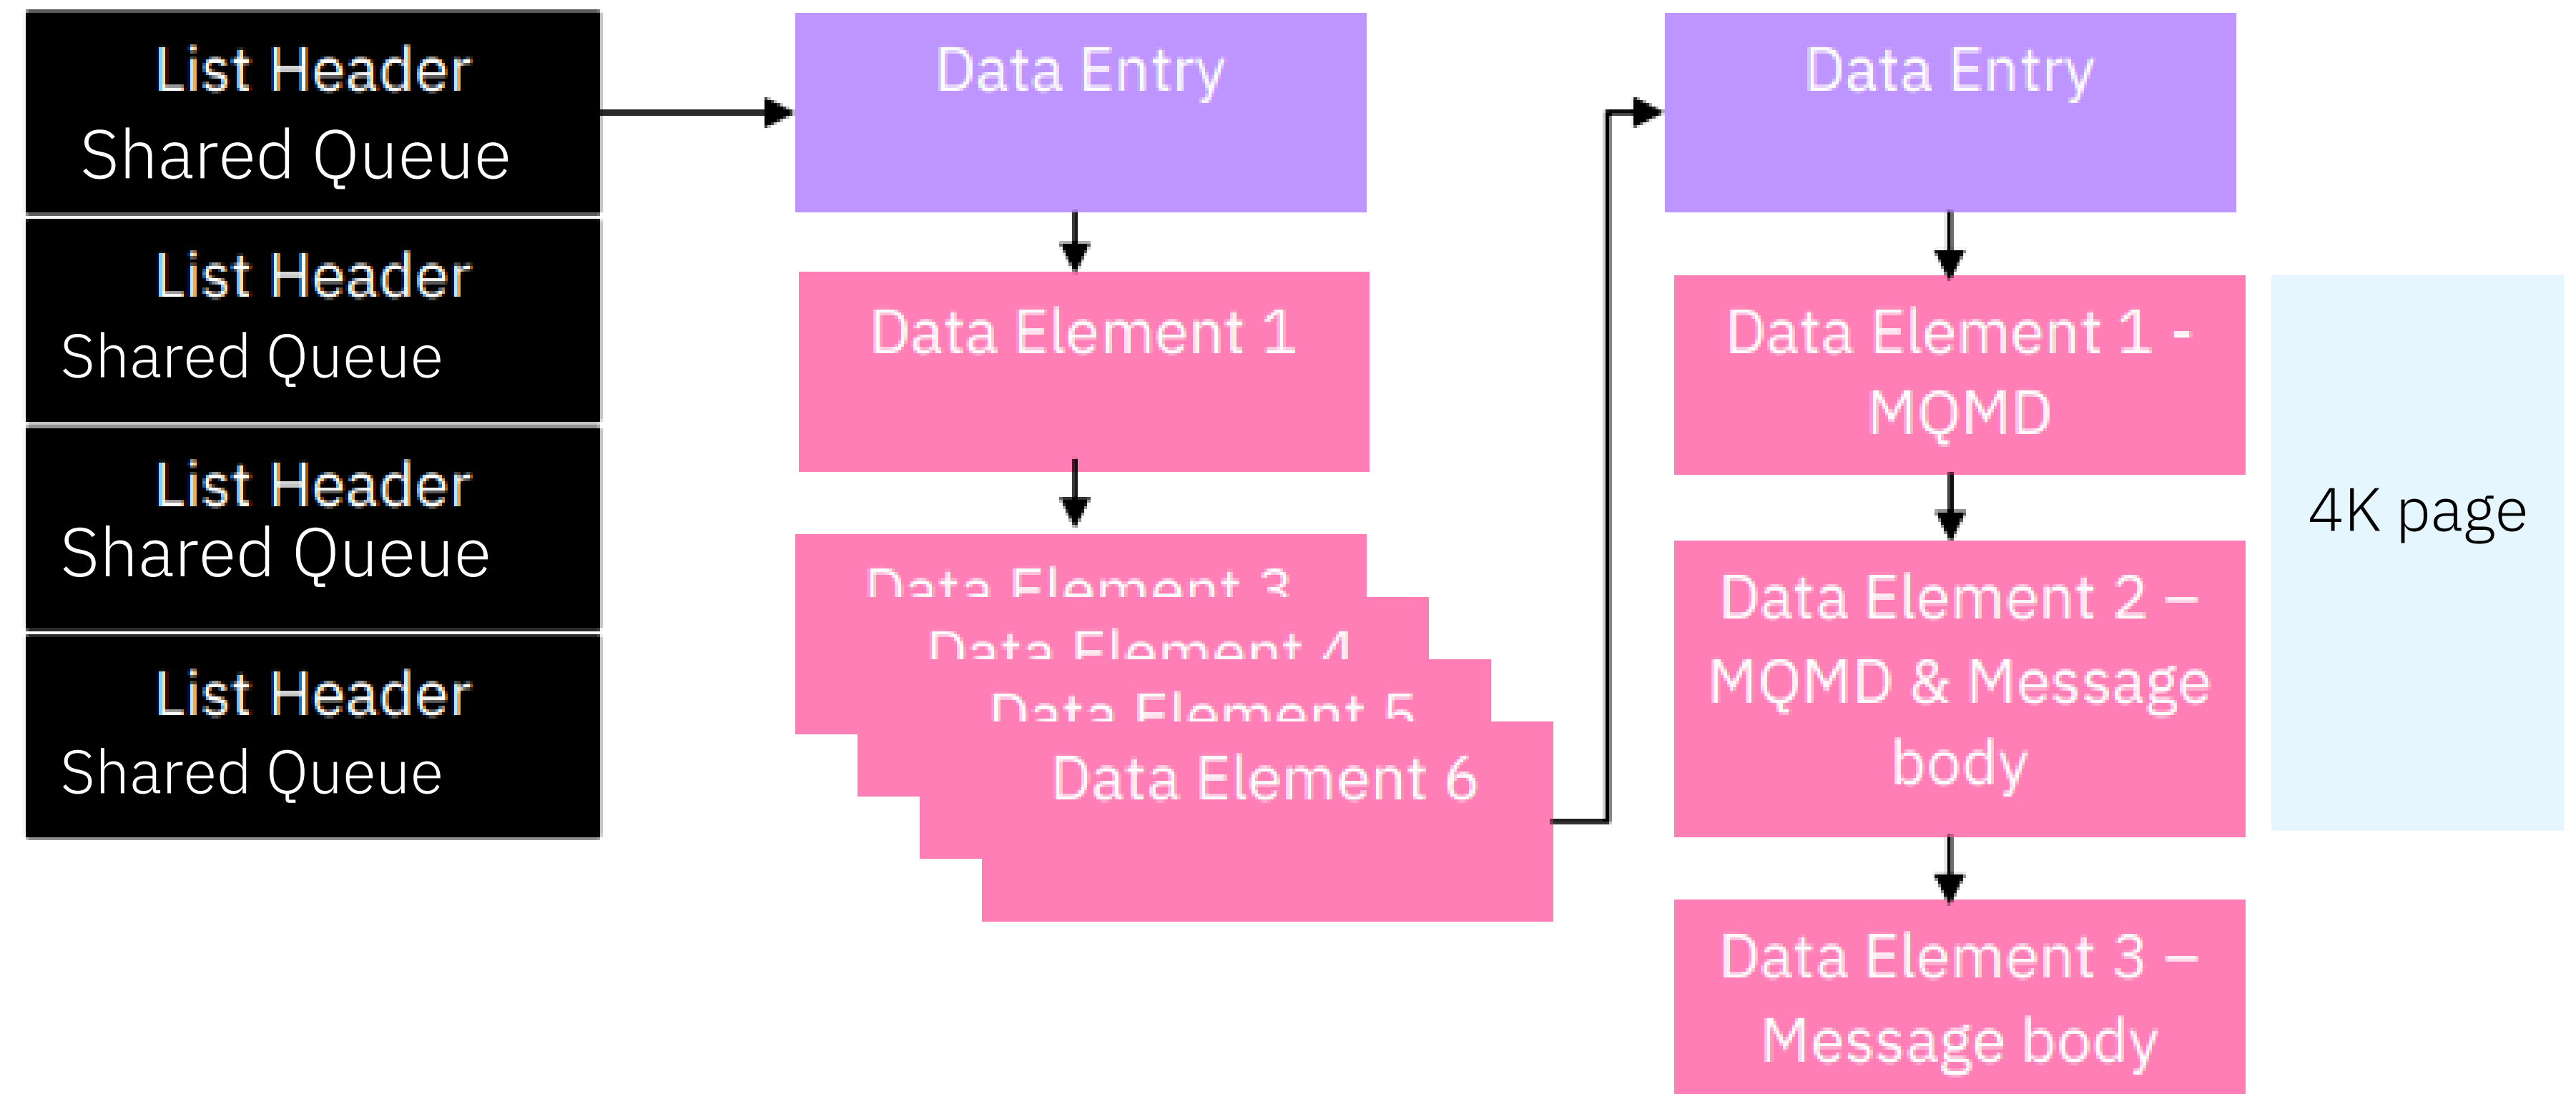


Internals of a shared queue

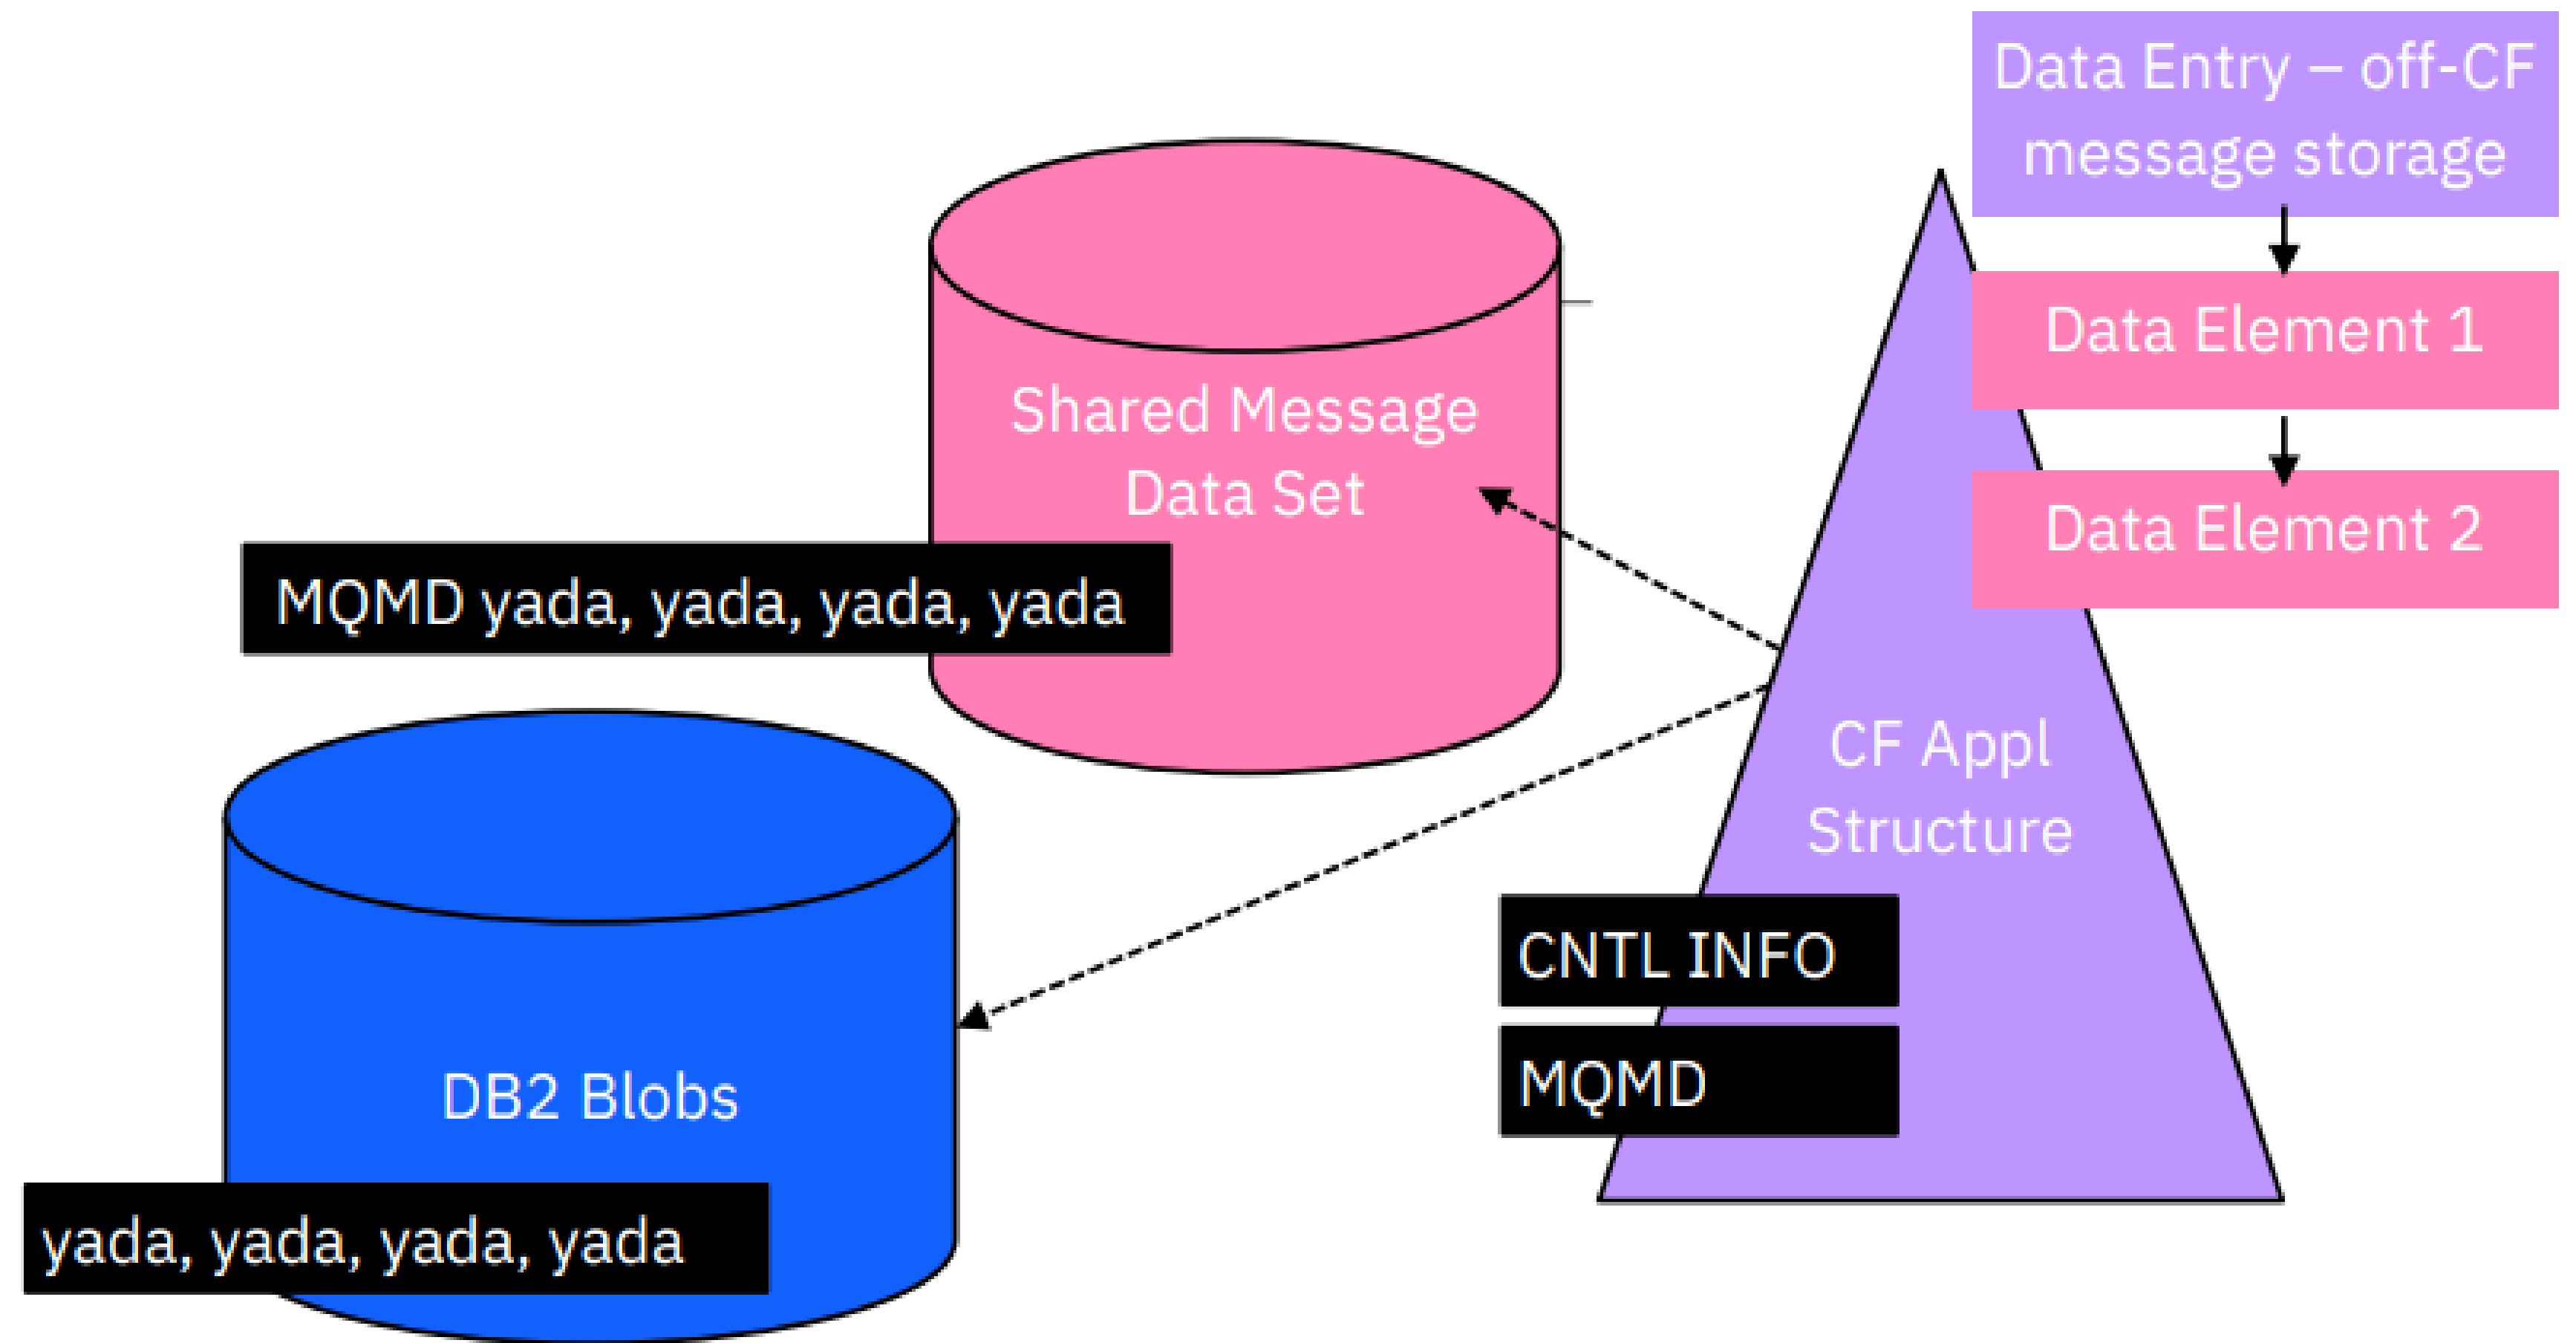
Internal Representation of a Coupling facility list structure



Internal Representation of a Shared Queue




Shared Queue Message Storage




Where does logging come in?

1) Persistent



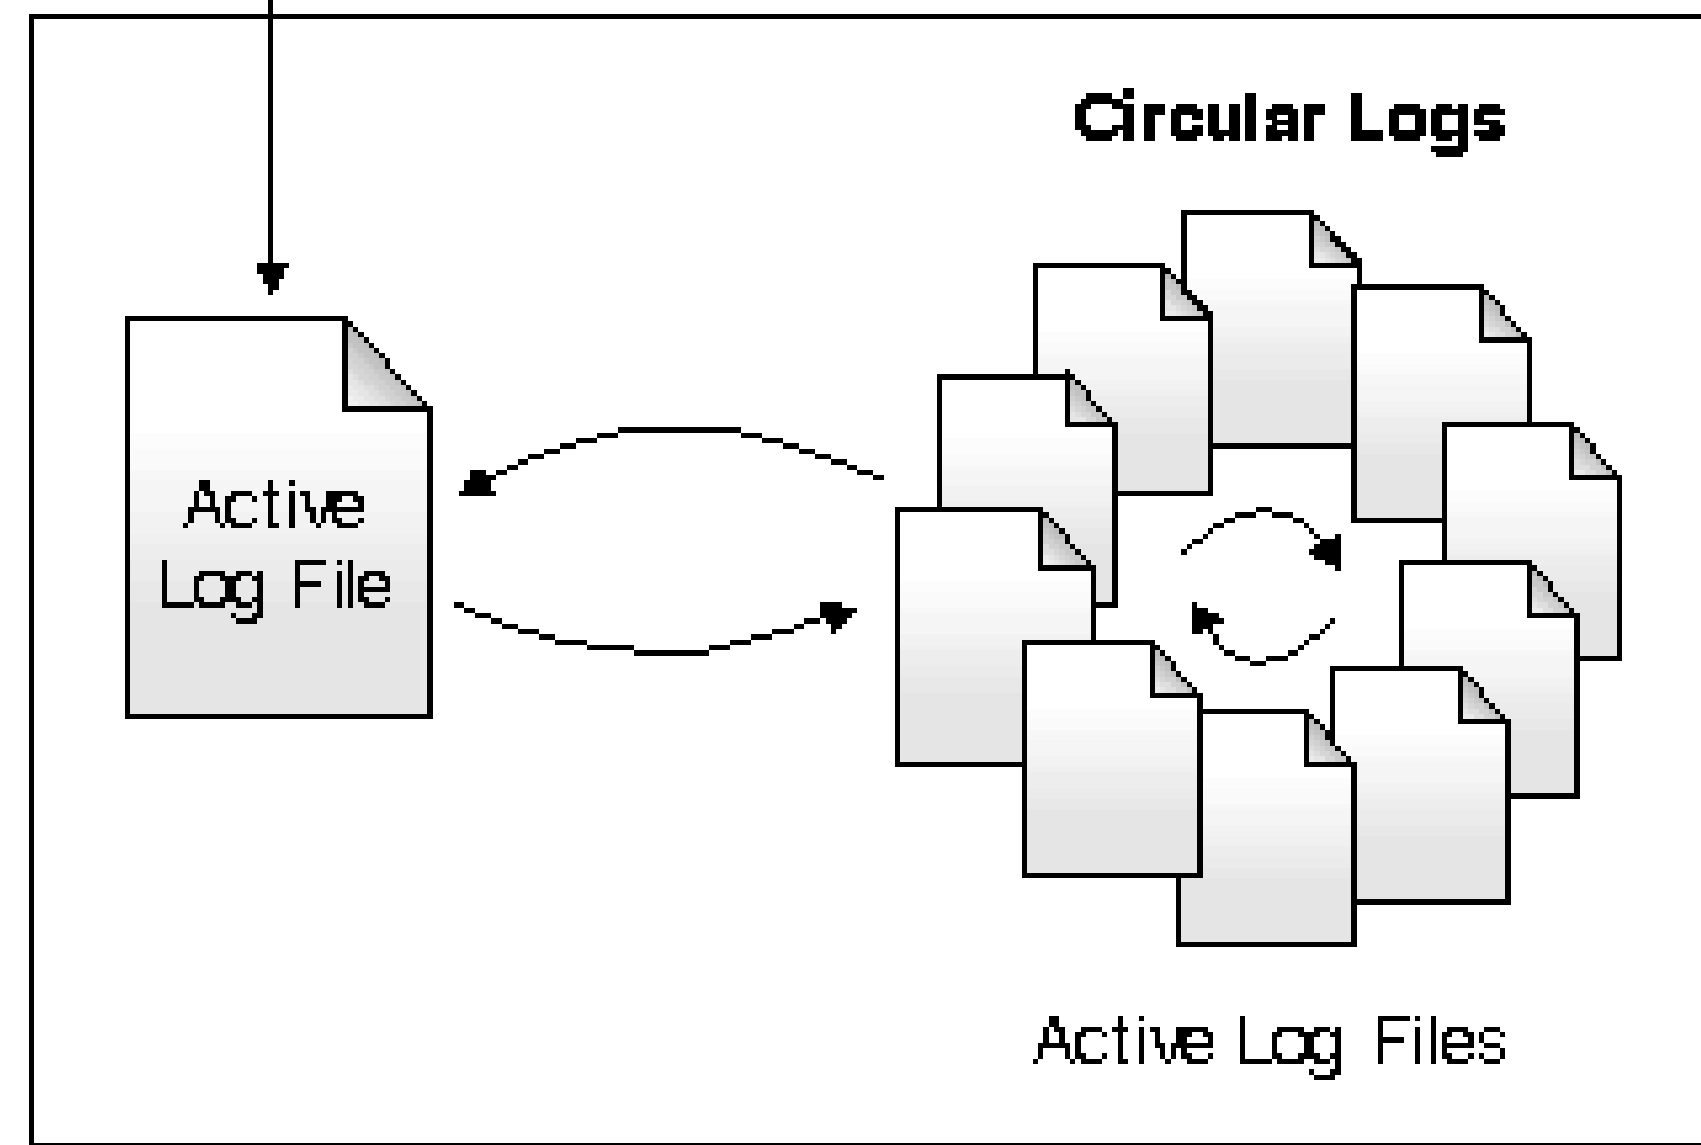
2) MQ Object

3) Queue Manager



An inventory of all active and archived log data sets known to IBM MQ.

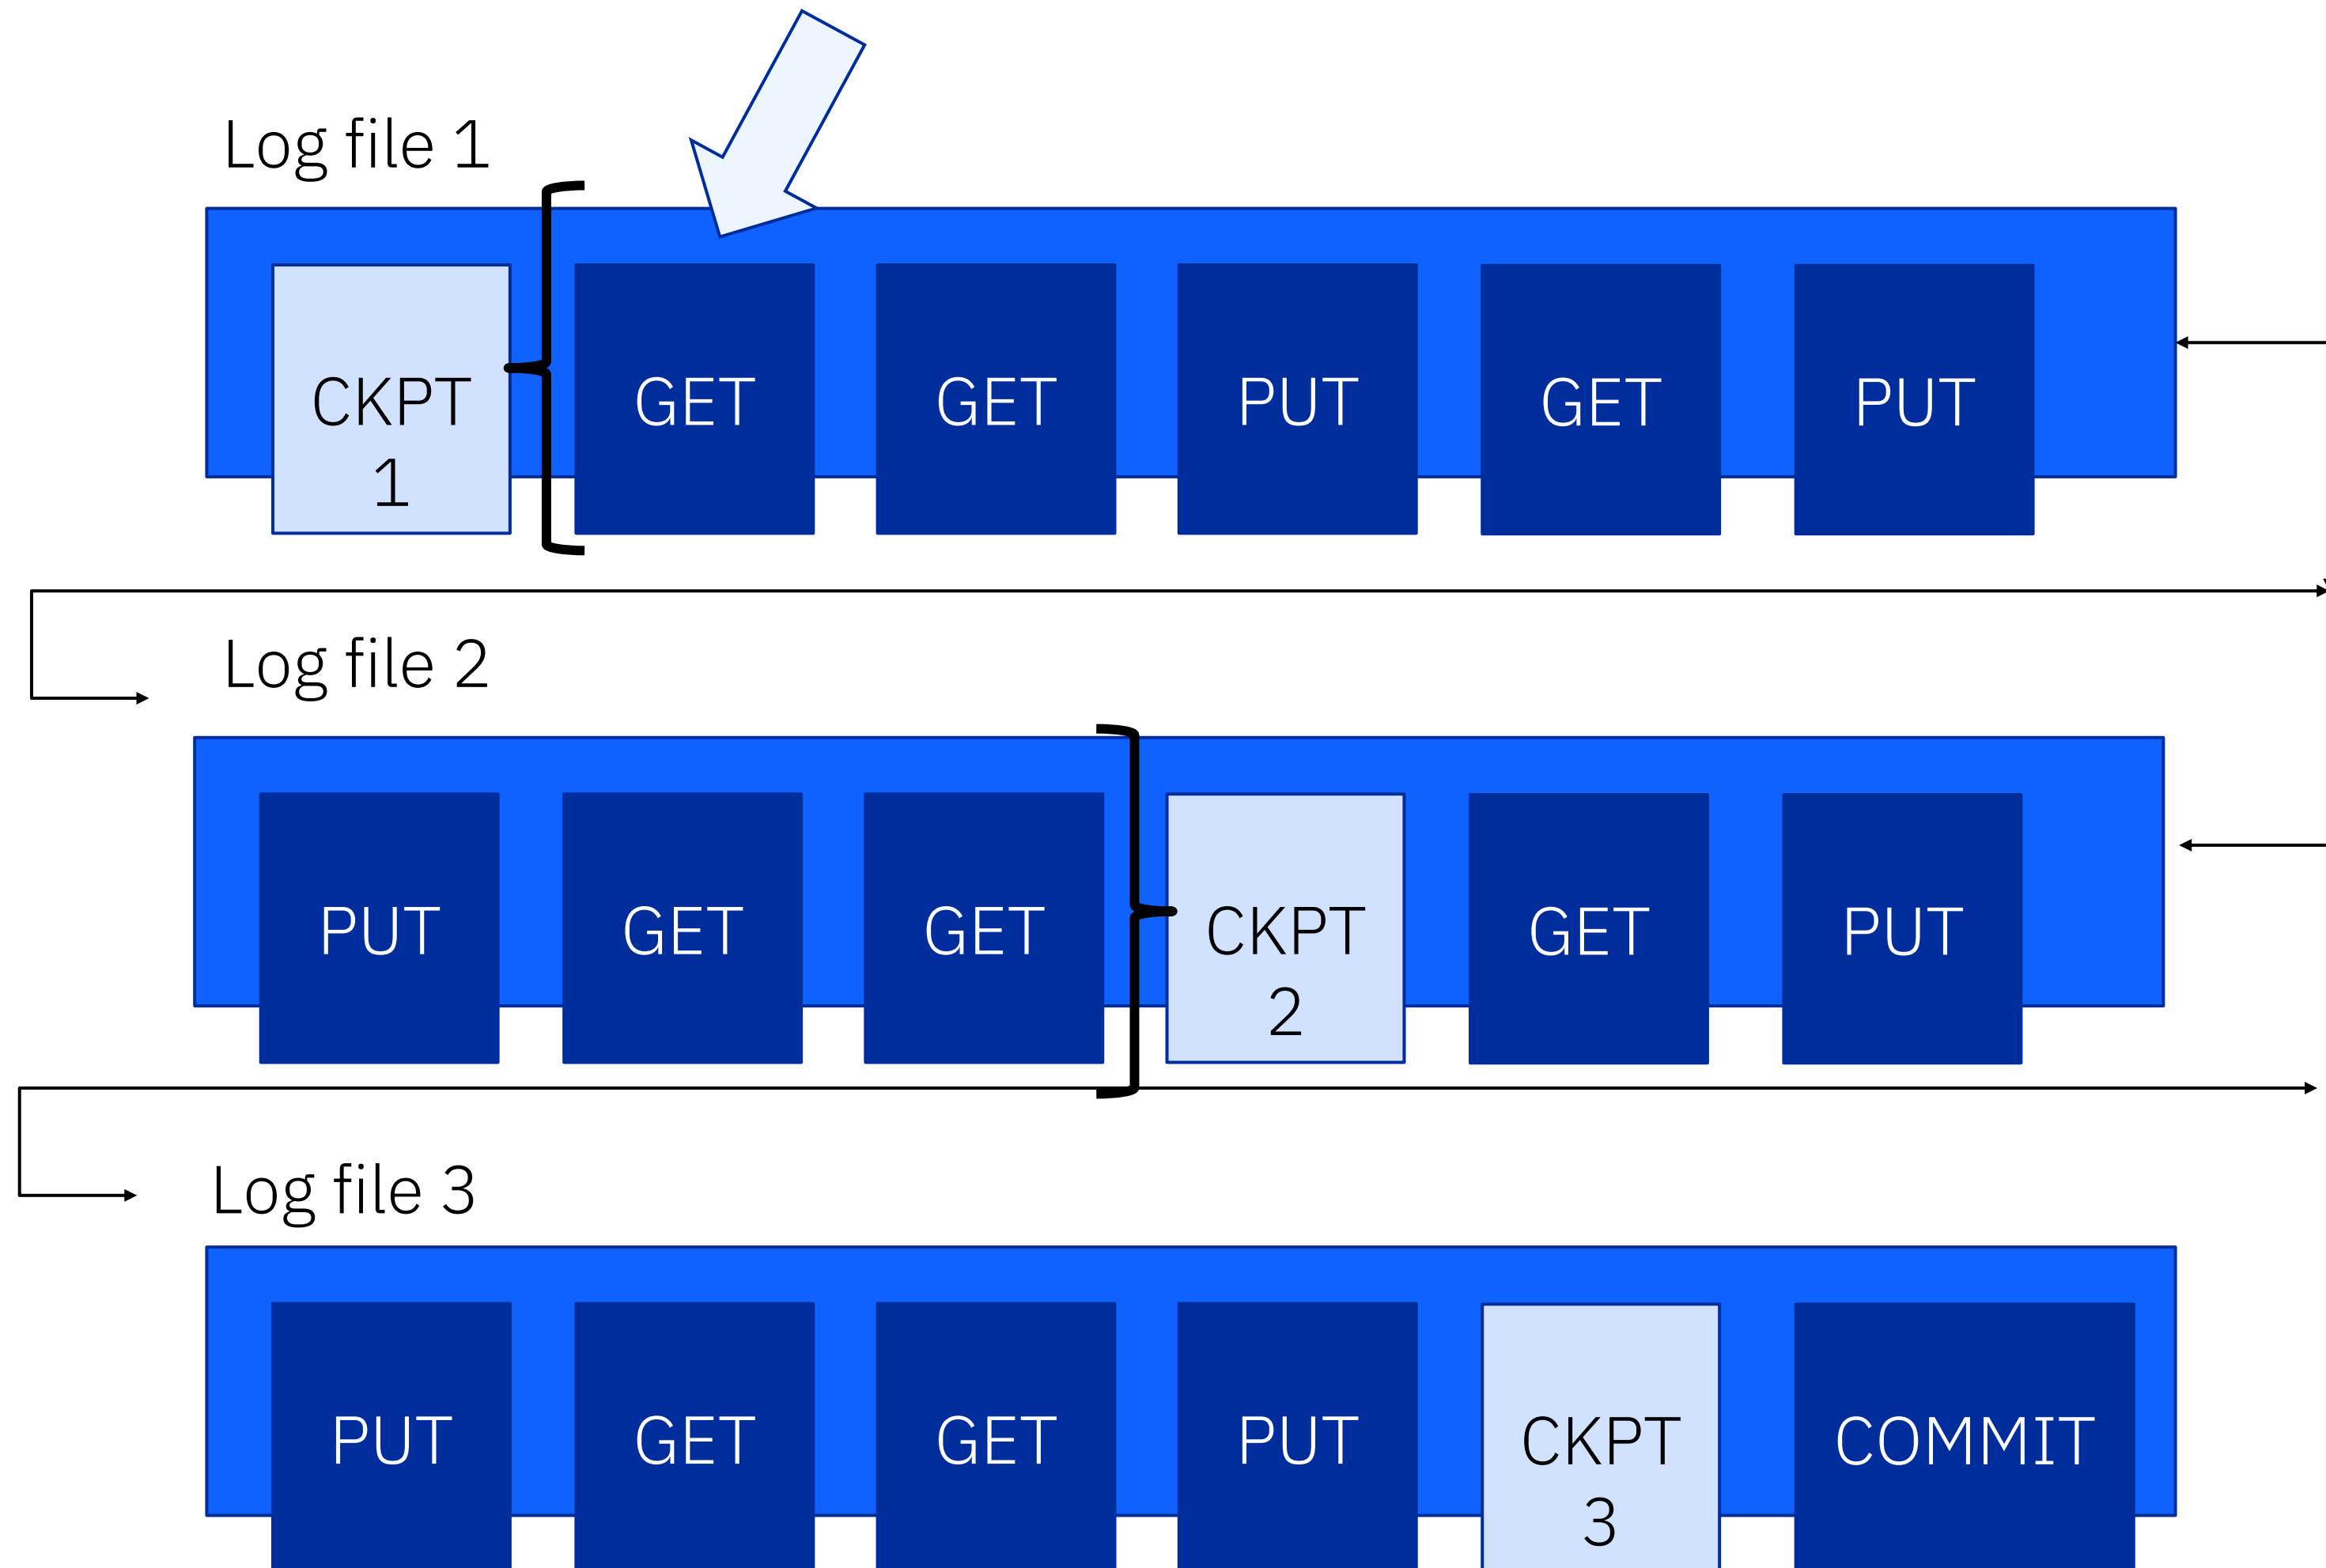
1. Unit of recovery log records
2. Checkpoint records
3. Page set control records
4. CF structure backup records



Offloading



What does a log file look like?



Define these bottom-up

Queue layer

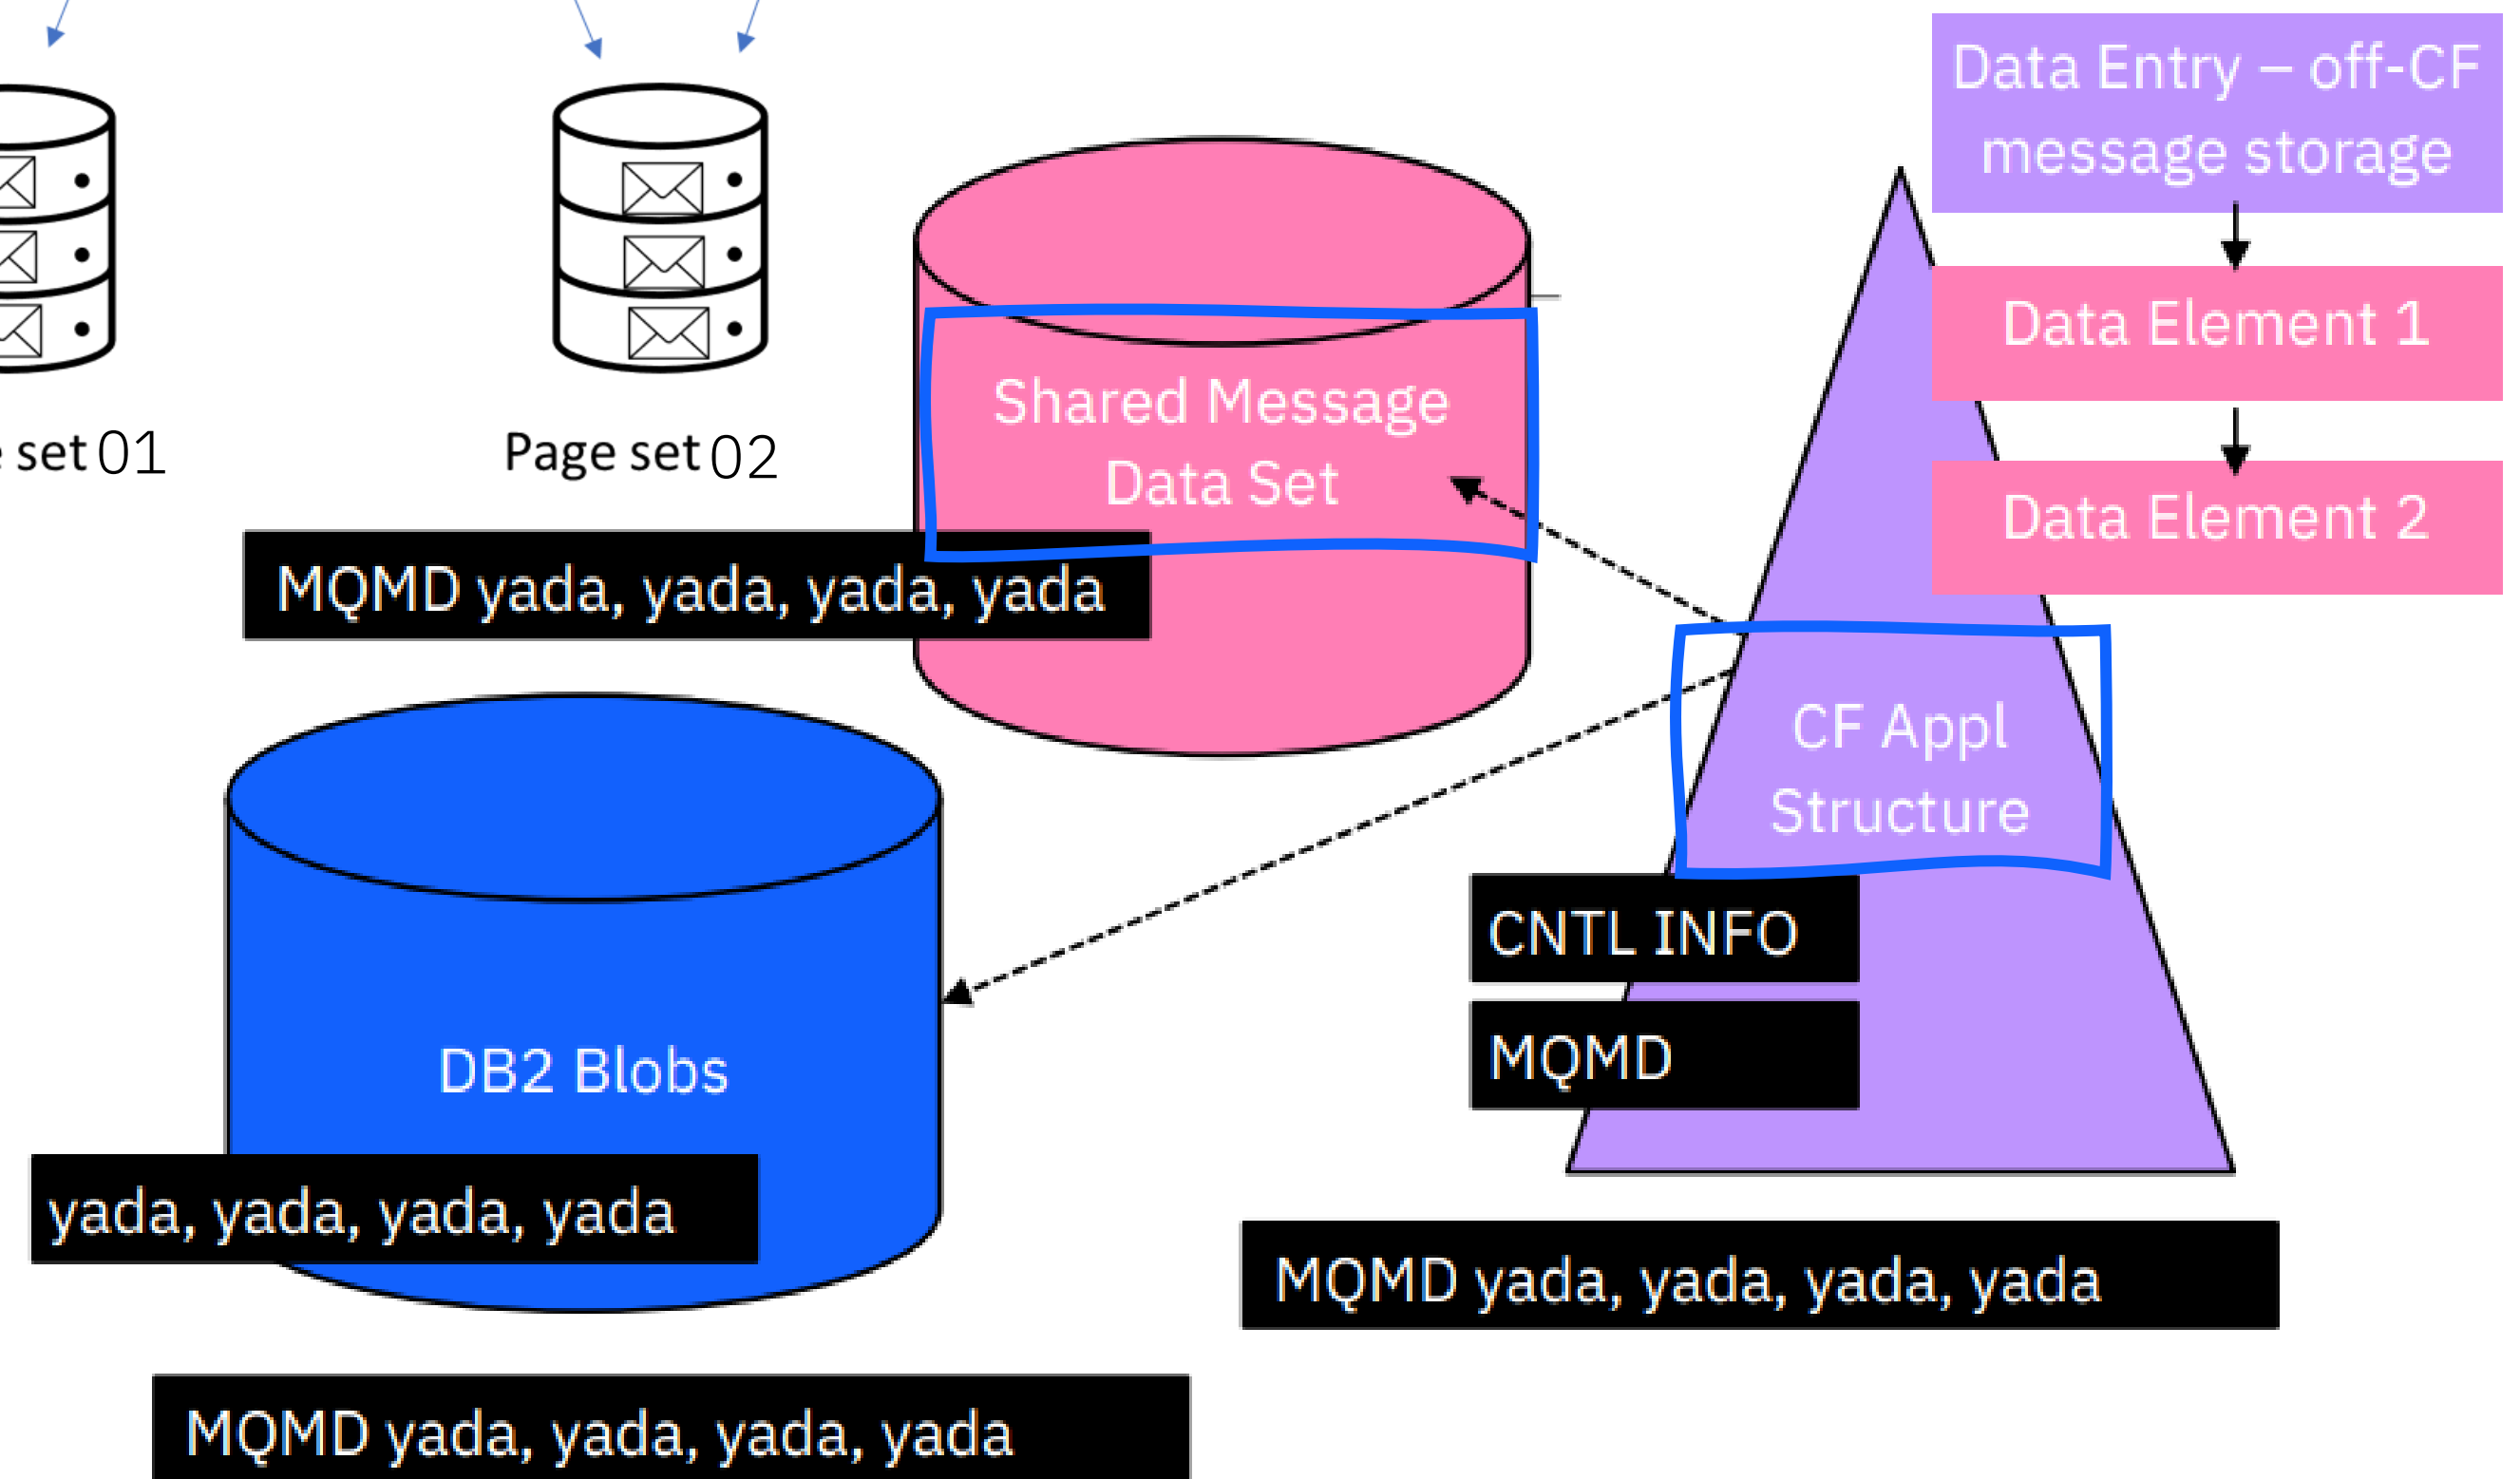
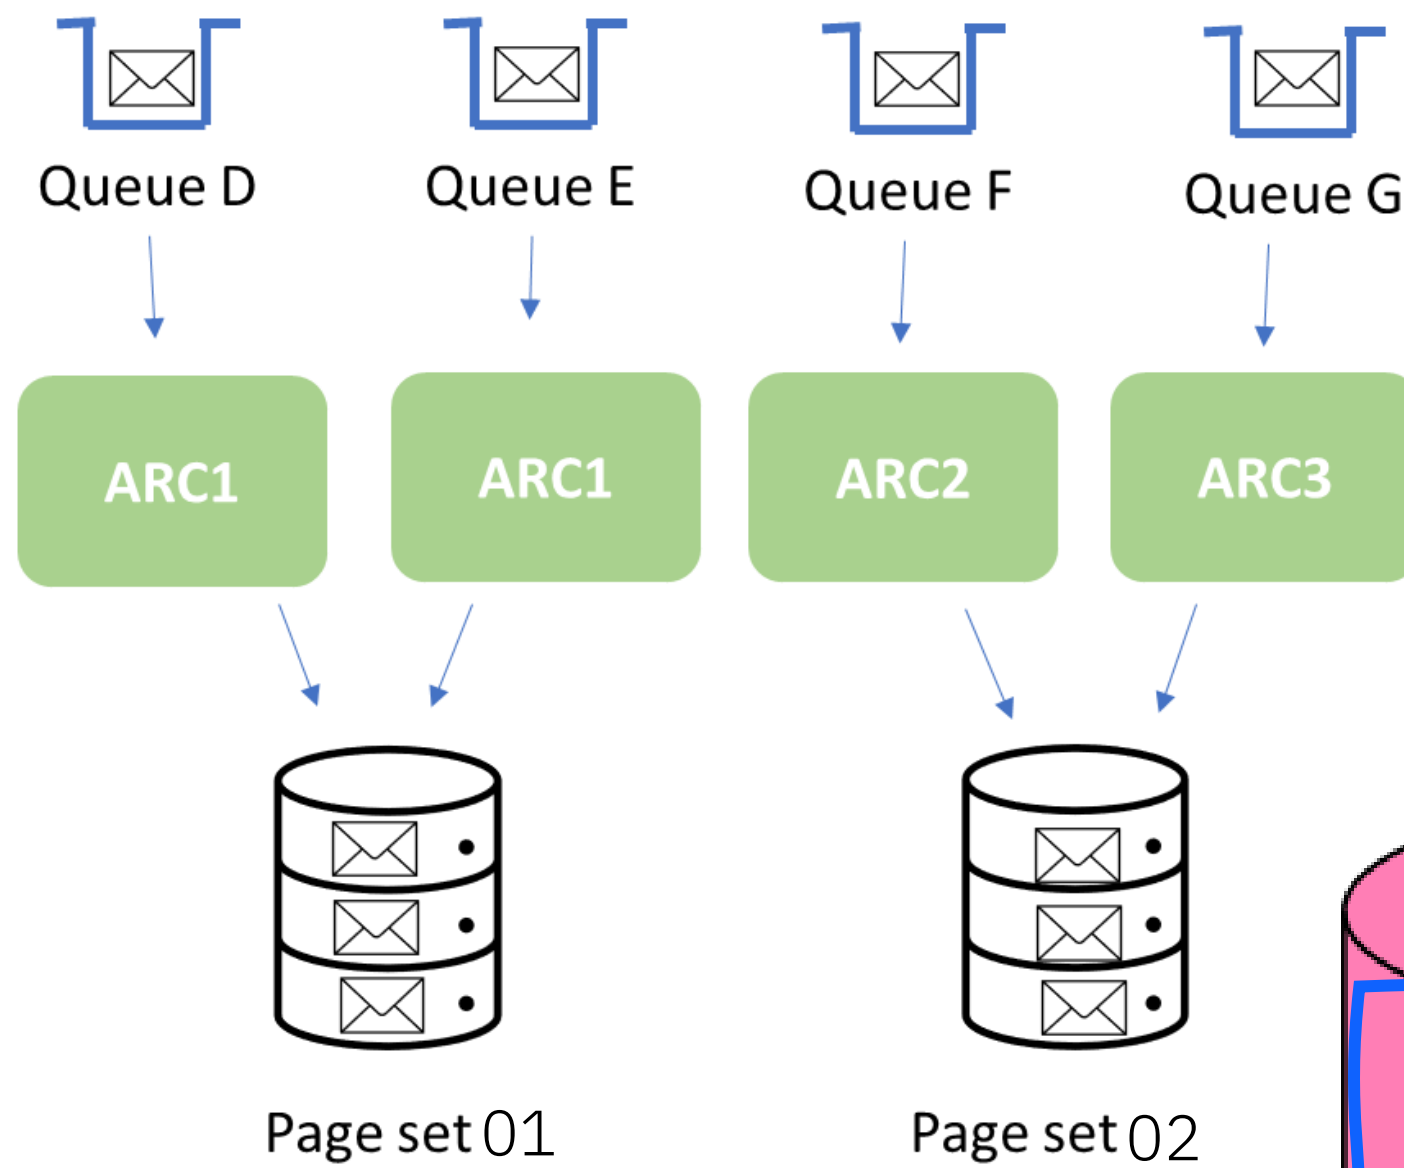
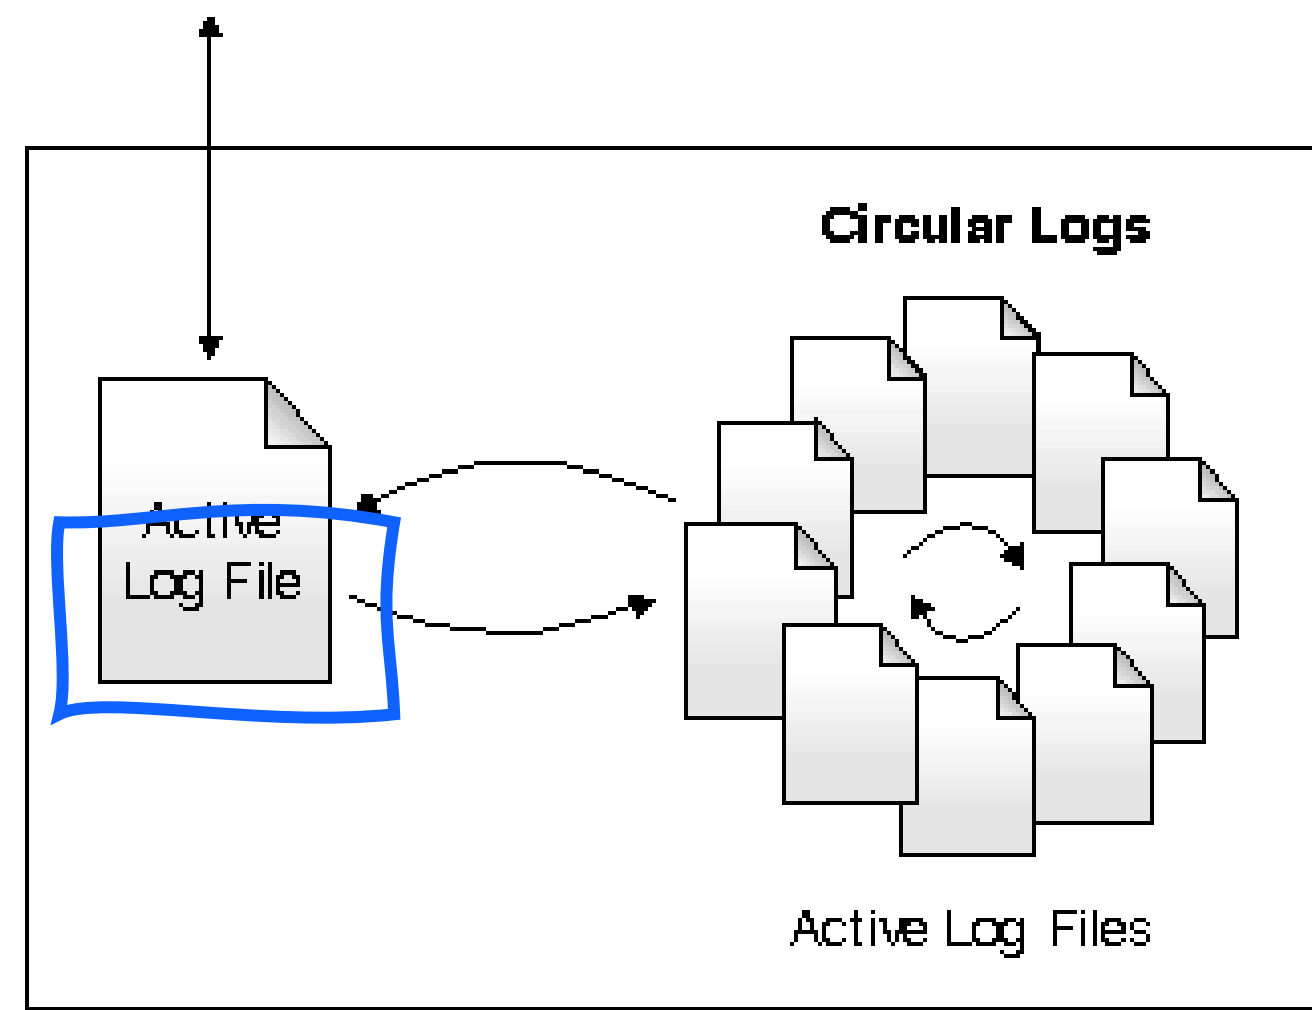
```
DEFINE QLOCAL (QUEUED)  
STGCLASS (ARC1)
```

Storage class layer

```
DEFINE STGCLASS (ARC1)  
PSID (A)
```

Page set layer

```
DEFINE PSID (01) BUFFPOOL (1)
```



Concept check

When I want to offload messages from my list structure, I should use...

(a) DB2 Blobs

(b) Shared Message Data sets

(c) Page sets

Why might a short message be classified as a 4k or less?

Which address space is the Pub/Sub engine associated with?

(a) QMGR master address space

(b) CHIN address space

What is the size of an element in a list structure?

To recap...

Private queues use buffer pools, storage classes, and page sets to underpin queuing

Shared queues use CF list structures, shared message data sets, and BLOBs to underpin queuing

Both private and shared queues use **logging** for recovery

Understanding storage and logging sets the stage for our next presentation...

How to analyze your IBM MQ for z/OS SMF data

There are two types of SMF records that are relevant to MQ for z/OS:

SMF 115:

Statistics data produced by an IBM MQ queue manager

SMF 116:

Accounting data produced by an IBM MQ queue manager

* You can look at this data in two ways – on z/OS and through exporting to CSV files

SMF 115

MQ
Storage

[SMF-QIS1.csv – Page Set Statistics](#)

[SMF-QPST.csv – Buffer Manager](#)

[SMF-QJST.csv – Log Manager](#)

SMF-QSGM.csv - Storage

SMF-QSPH.csv - Storage

SMF-QSRS.csv - Storage

SMF-QSST.csv – Storage

[SMF-QESD.csv – Shared Message Data Set](#)

[SMF-QEST.csv – Coupling Facility Statistics](#)

[SMF-Q5ST.csv – BLOB Statistics](#)

MQ
Requests

[SMF-QLST.csv – Lock Manager](#)

[SMF-QMST.csv – Message Manager](#)

[SMF-QIST.csv – Data Manager Statistics](#)

[SMF-QCCT.csv – Channel Statistics](#)

[SMF-QCTADP.csv – Adapter Task Statistics](#)

[SMF-QCTDSP.csv – Dispatcher Task Statistics](#)

SMF-QCTSSL.csv – SSL Statistics

[SMF-QTST.csv – Publications Statistics](#)

SMF 116

MQ
Tasks

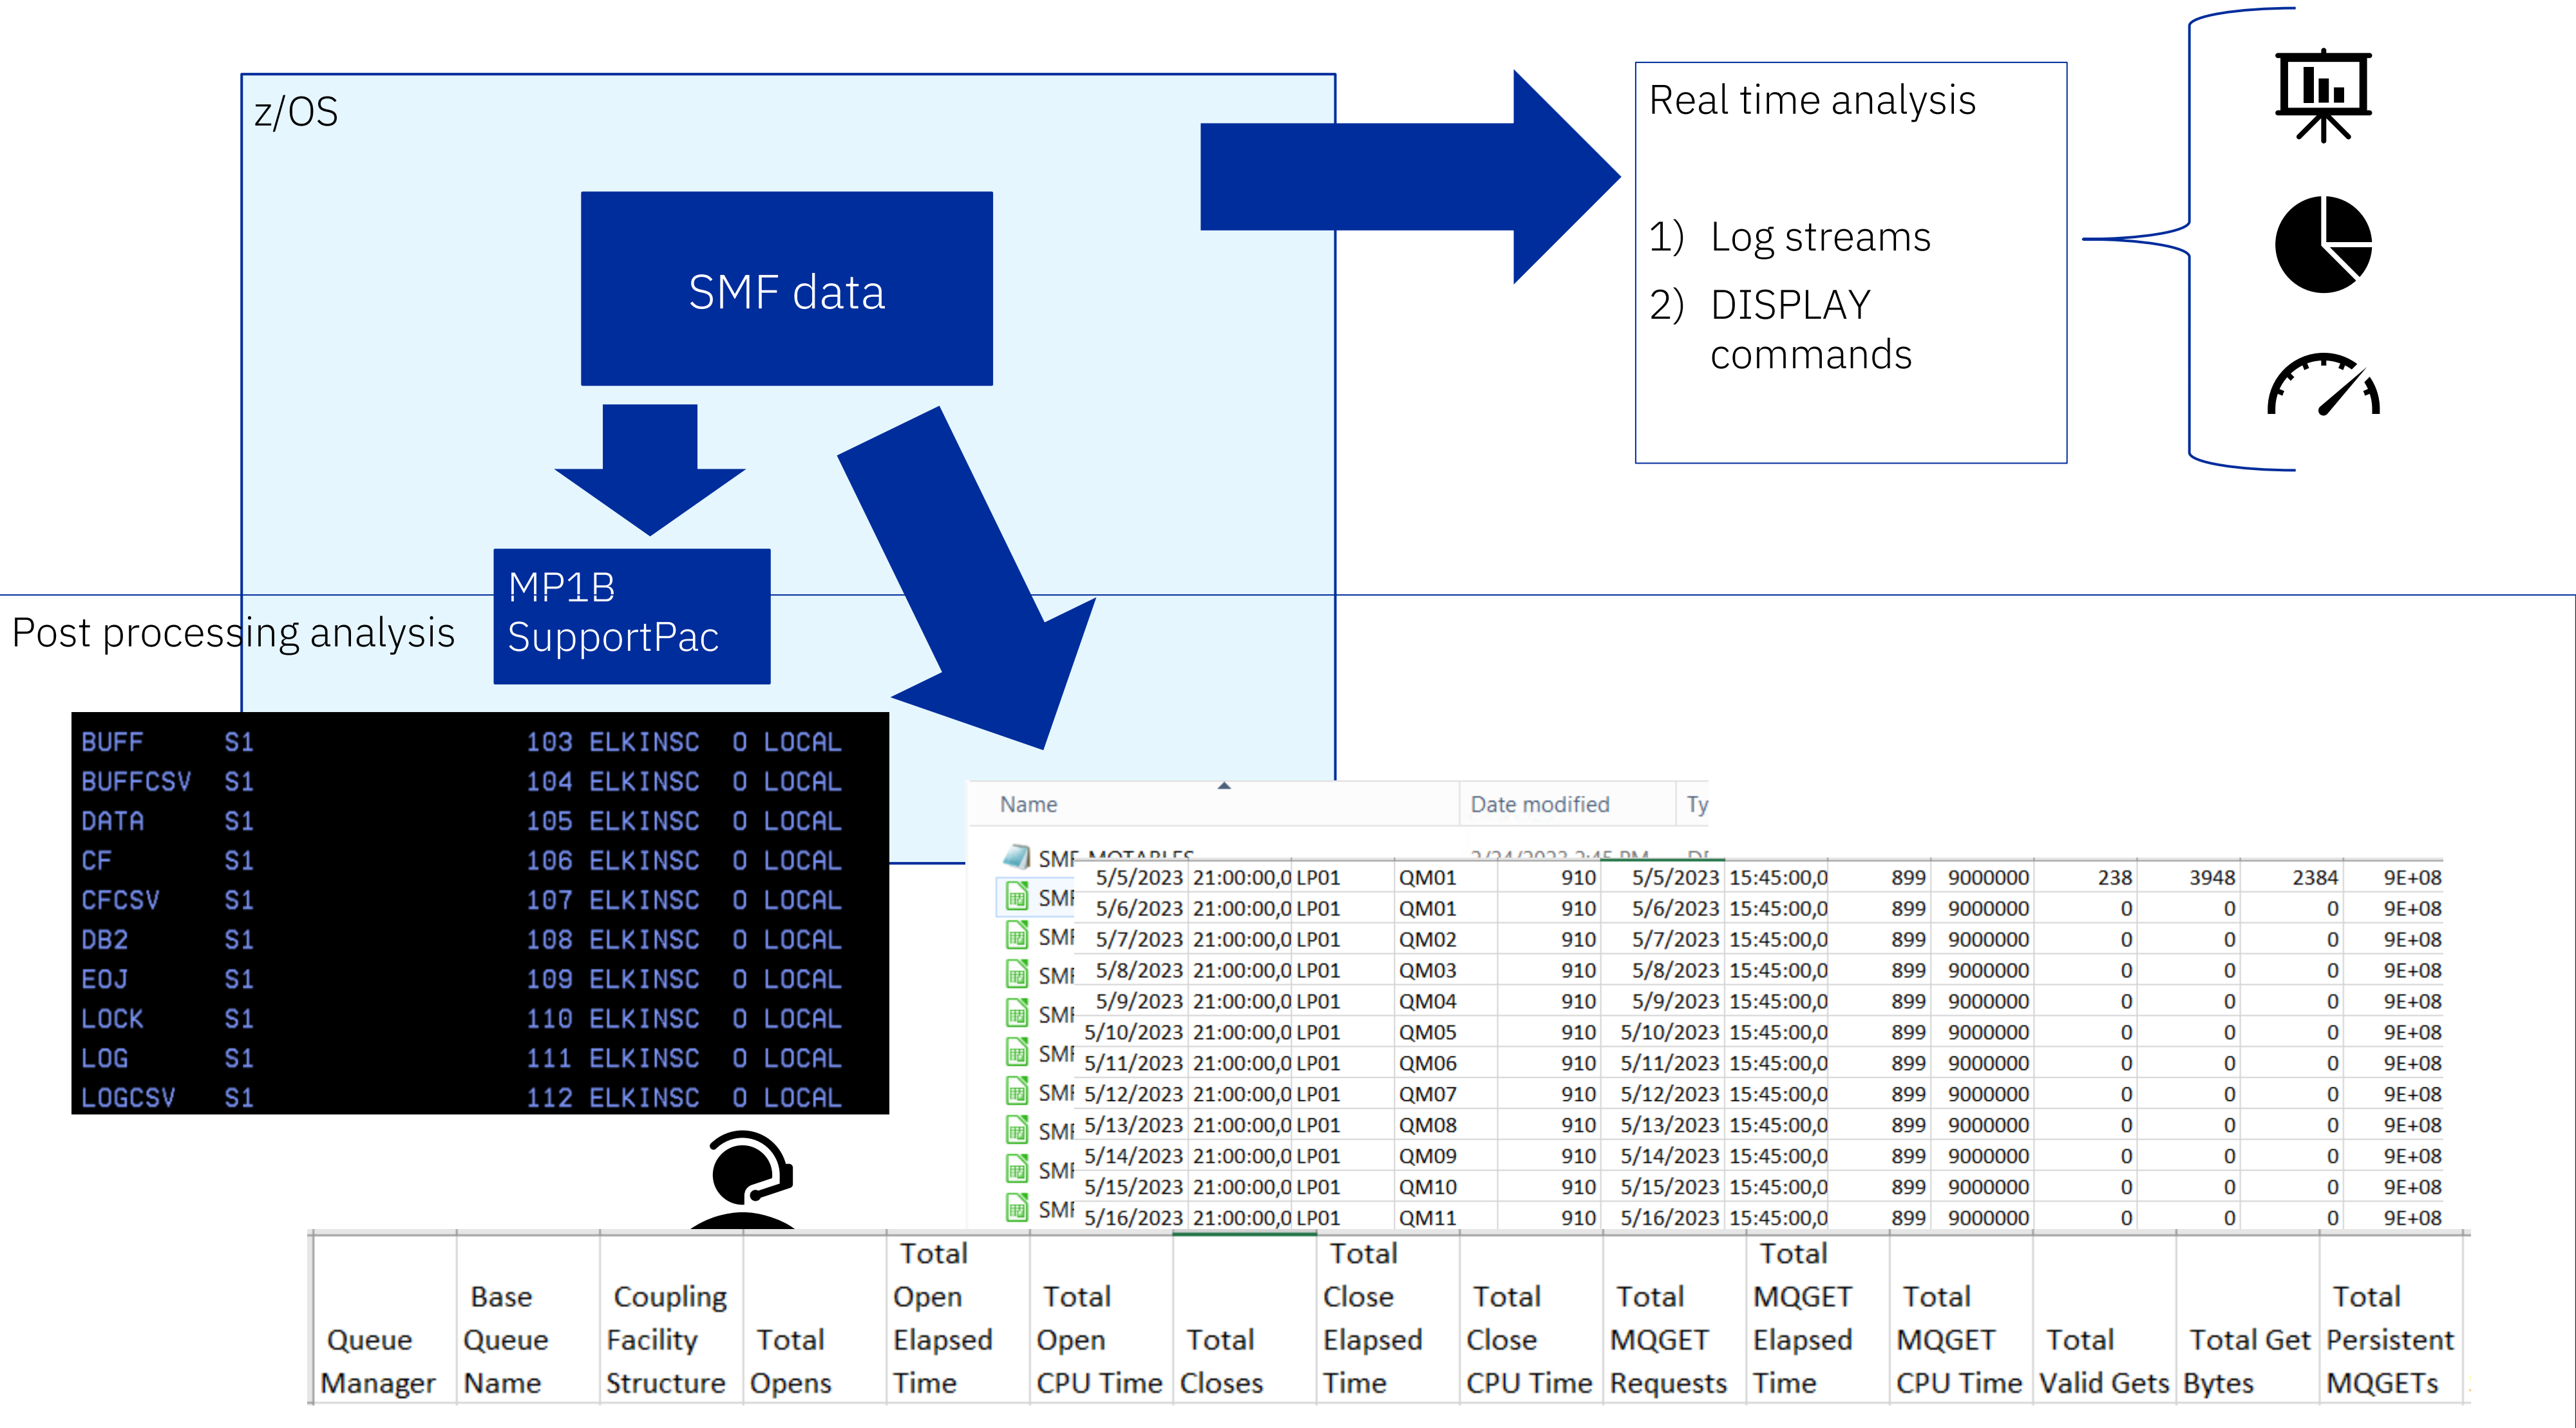
SMF-QCST.csv – Channel Accounting

SMF-WQ.csv – Task Queue Accounting

SMF-WTAS.csv – Task Accounting

SMF-WTID.csv – Task ID Accounting


How do you look at SMF data?



CICS Transaction Rate Achieved

All transactions monitored with MQ Accounting Class(3) enabled

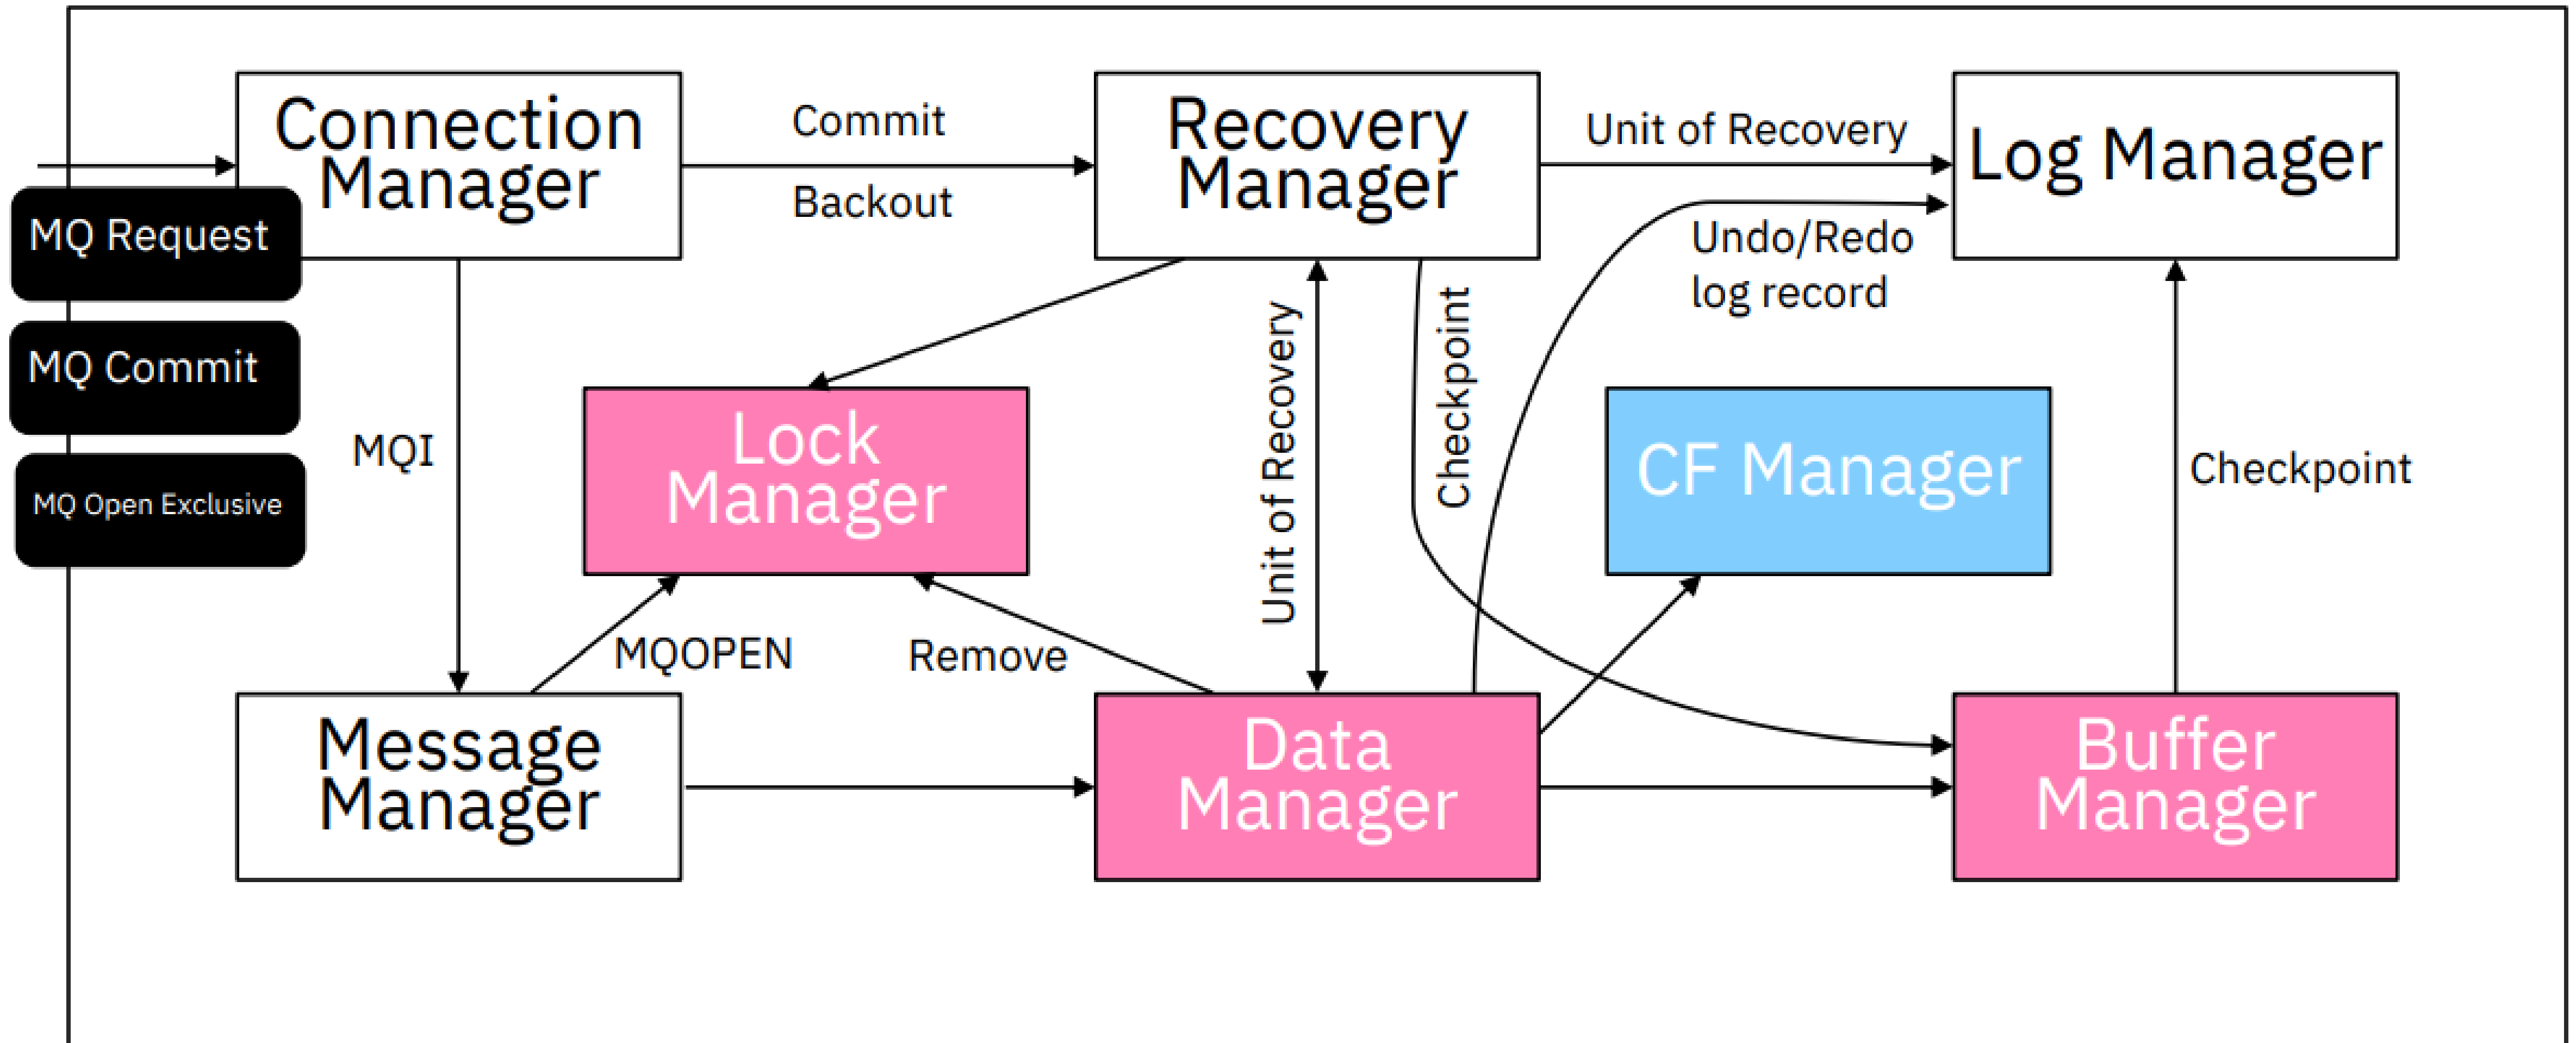
- SMF written to Dataset
- SMF written to LogStream
- SMF written to LogStream with hardware compression



Interpreting SMF 115 data for private queues

Building Blocks

Resource Managers



First Line Managers

The threads within the QMMSTR address space who do the real work within each queue manager by interacting with applications and the underlying z/OS resource managers

They include:

Connection Manager – not the channel initiator, but local connections

Recovery Manager

Log Manager

Message Manager

Topic Manager

Data Manager

Buffer Manager

Lock Manager

Storage Manager

CF Manager

Security Manager

Buffer Manager

The buffer manager remains the largest area of opportunity for performance tuning for private queues.

- 1) Keeping private queue messages in-memory (buffer pools), especially for responsiveness driven workload, can be critical for meeting those goals.
- 2) Buffer pool contention, that is too much activity on a pool, can also cause delays due to the volume of work being done.



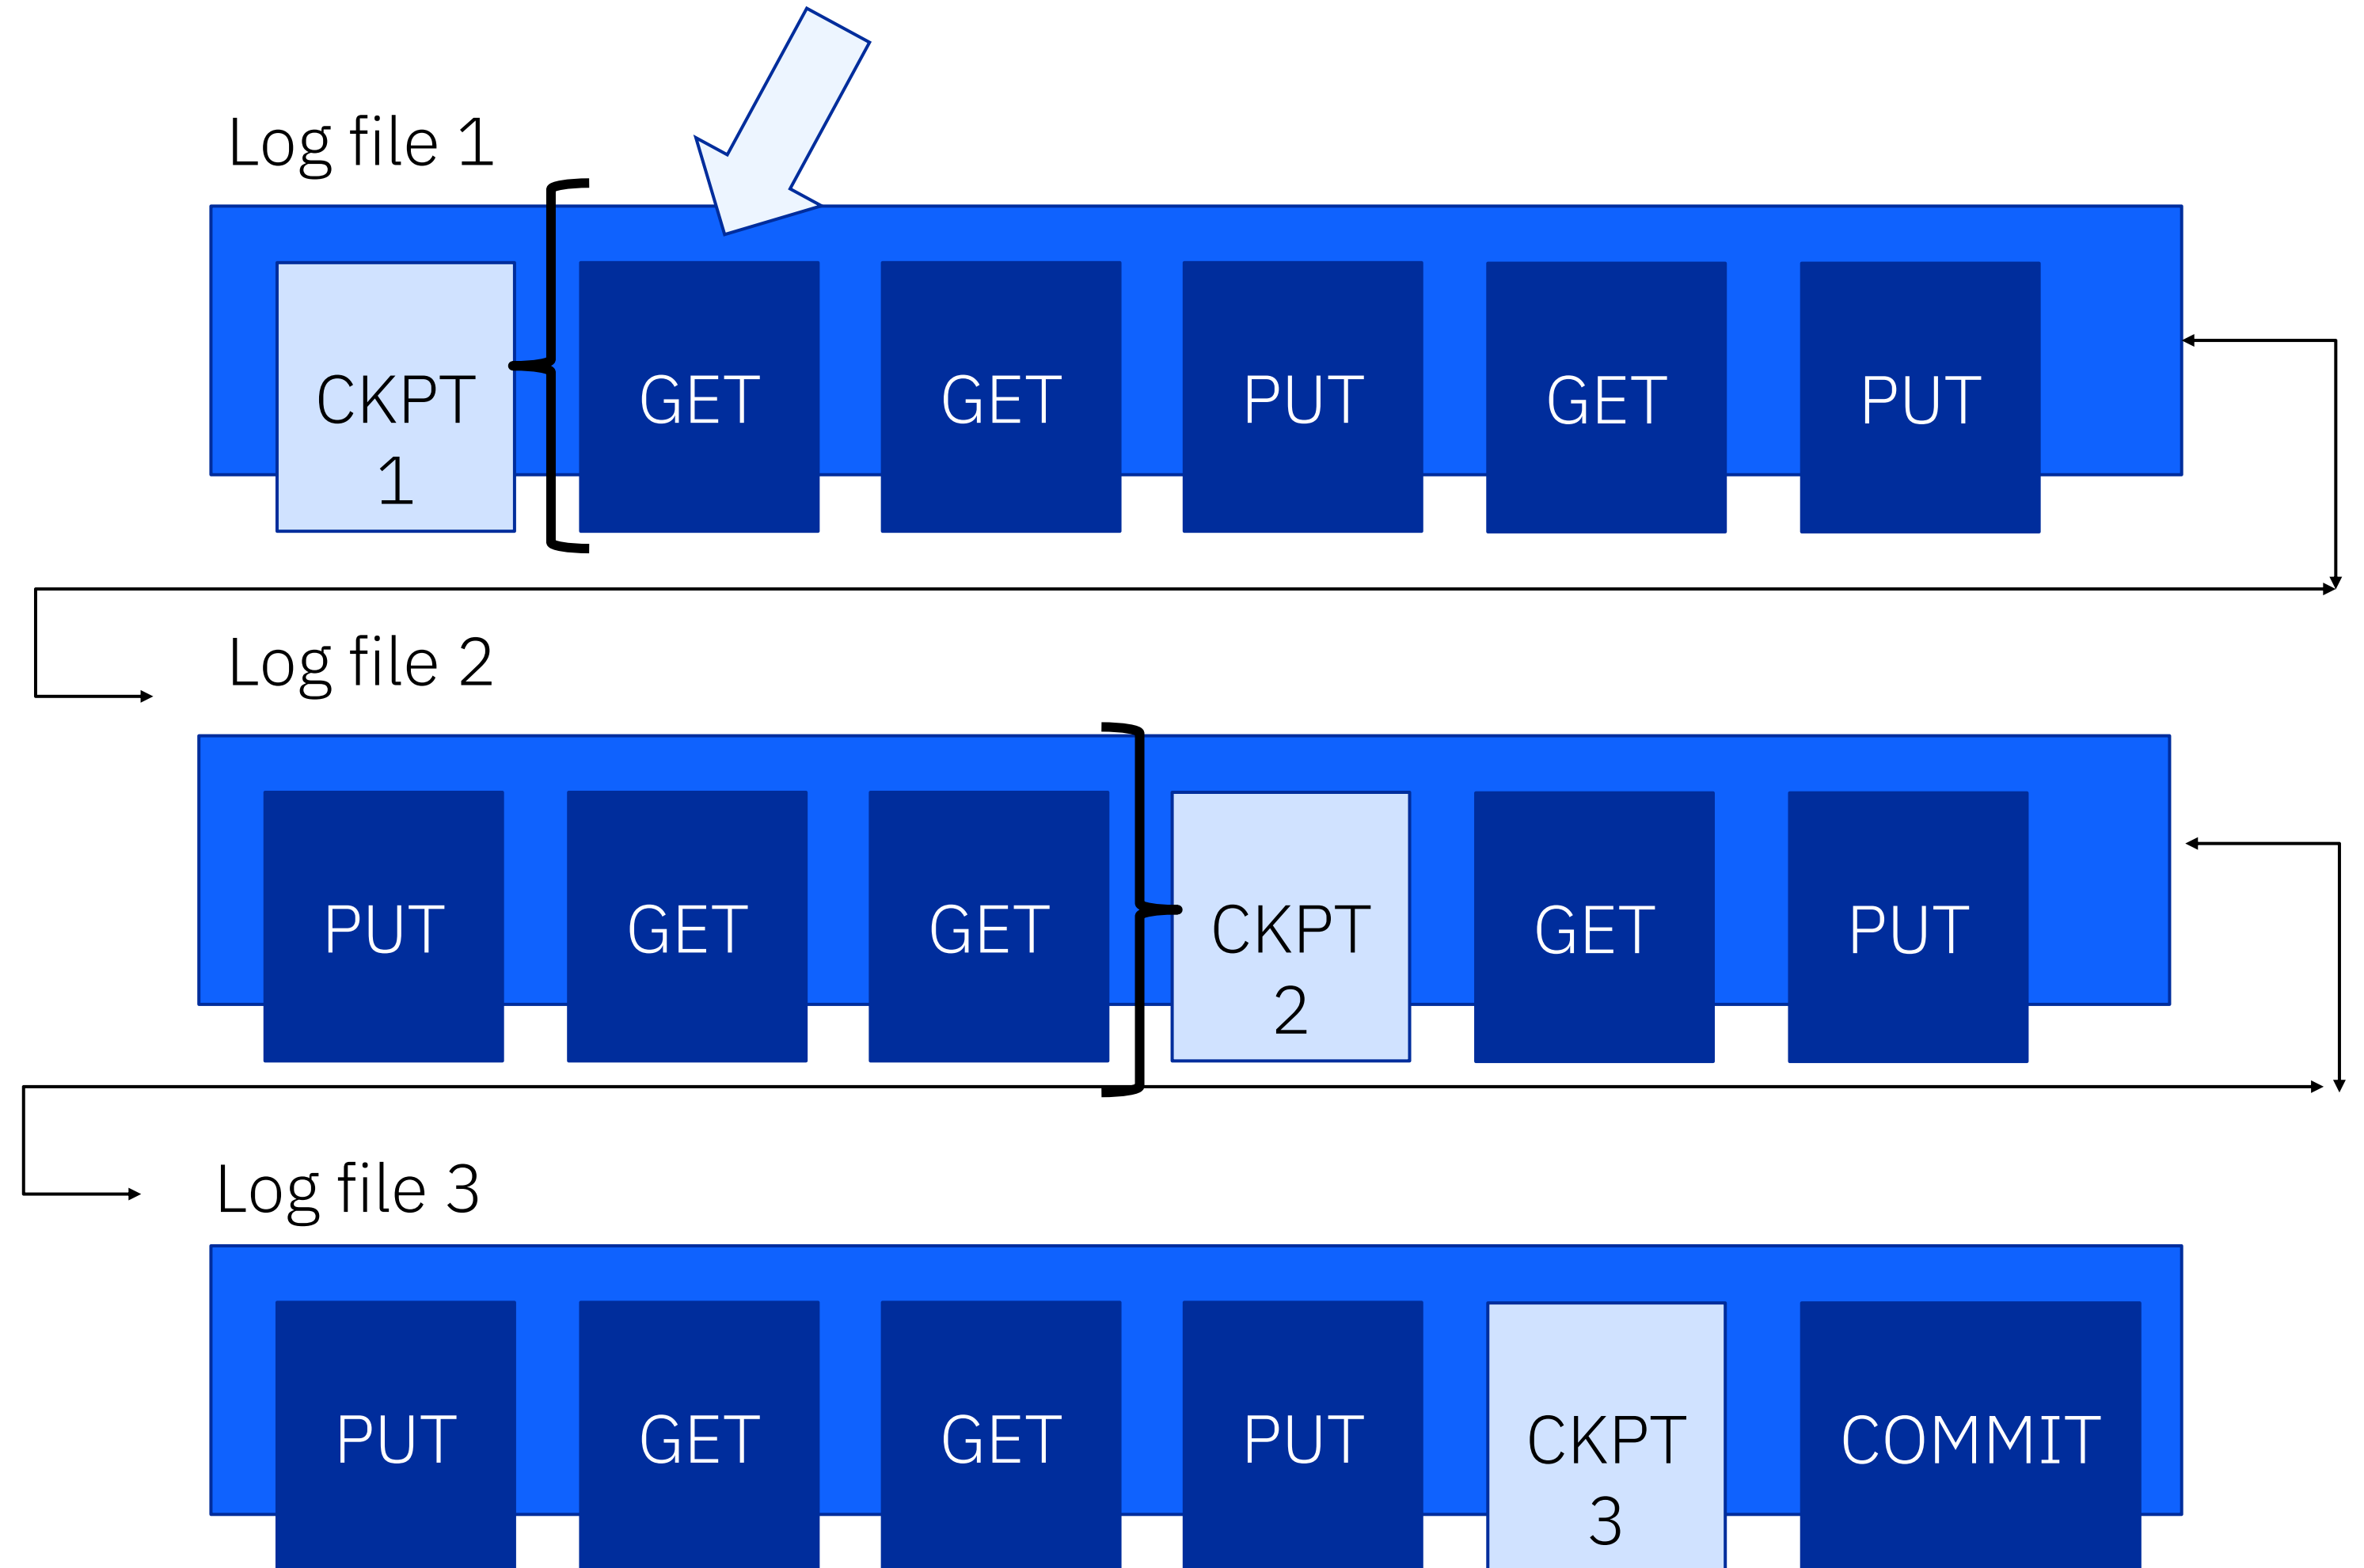
SMF-QPST.csv – Buffer Manager Statistics

- Buffer Pool
- Buffer Count
- Highest Used Percent
- Sync Writes
- Defer Write Thold Reached
- Sync Write Thold Reached
- Suspend No Buffers
- Pages Written
- Date
- Time
- LPAR
- QMgr
- MQ Version
- Interval Duration
- Location
- Pagefixed
- Current Stealable
- Getp Old Requests
- Getp New Requests
- DASD Read
- Set Write Pages
- DASD Write
- Lowest Stealable
- Highest Used
- Buffer Steals
- Buffer Steals Hash Changes

Log Manager

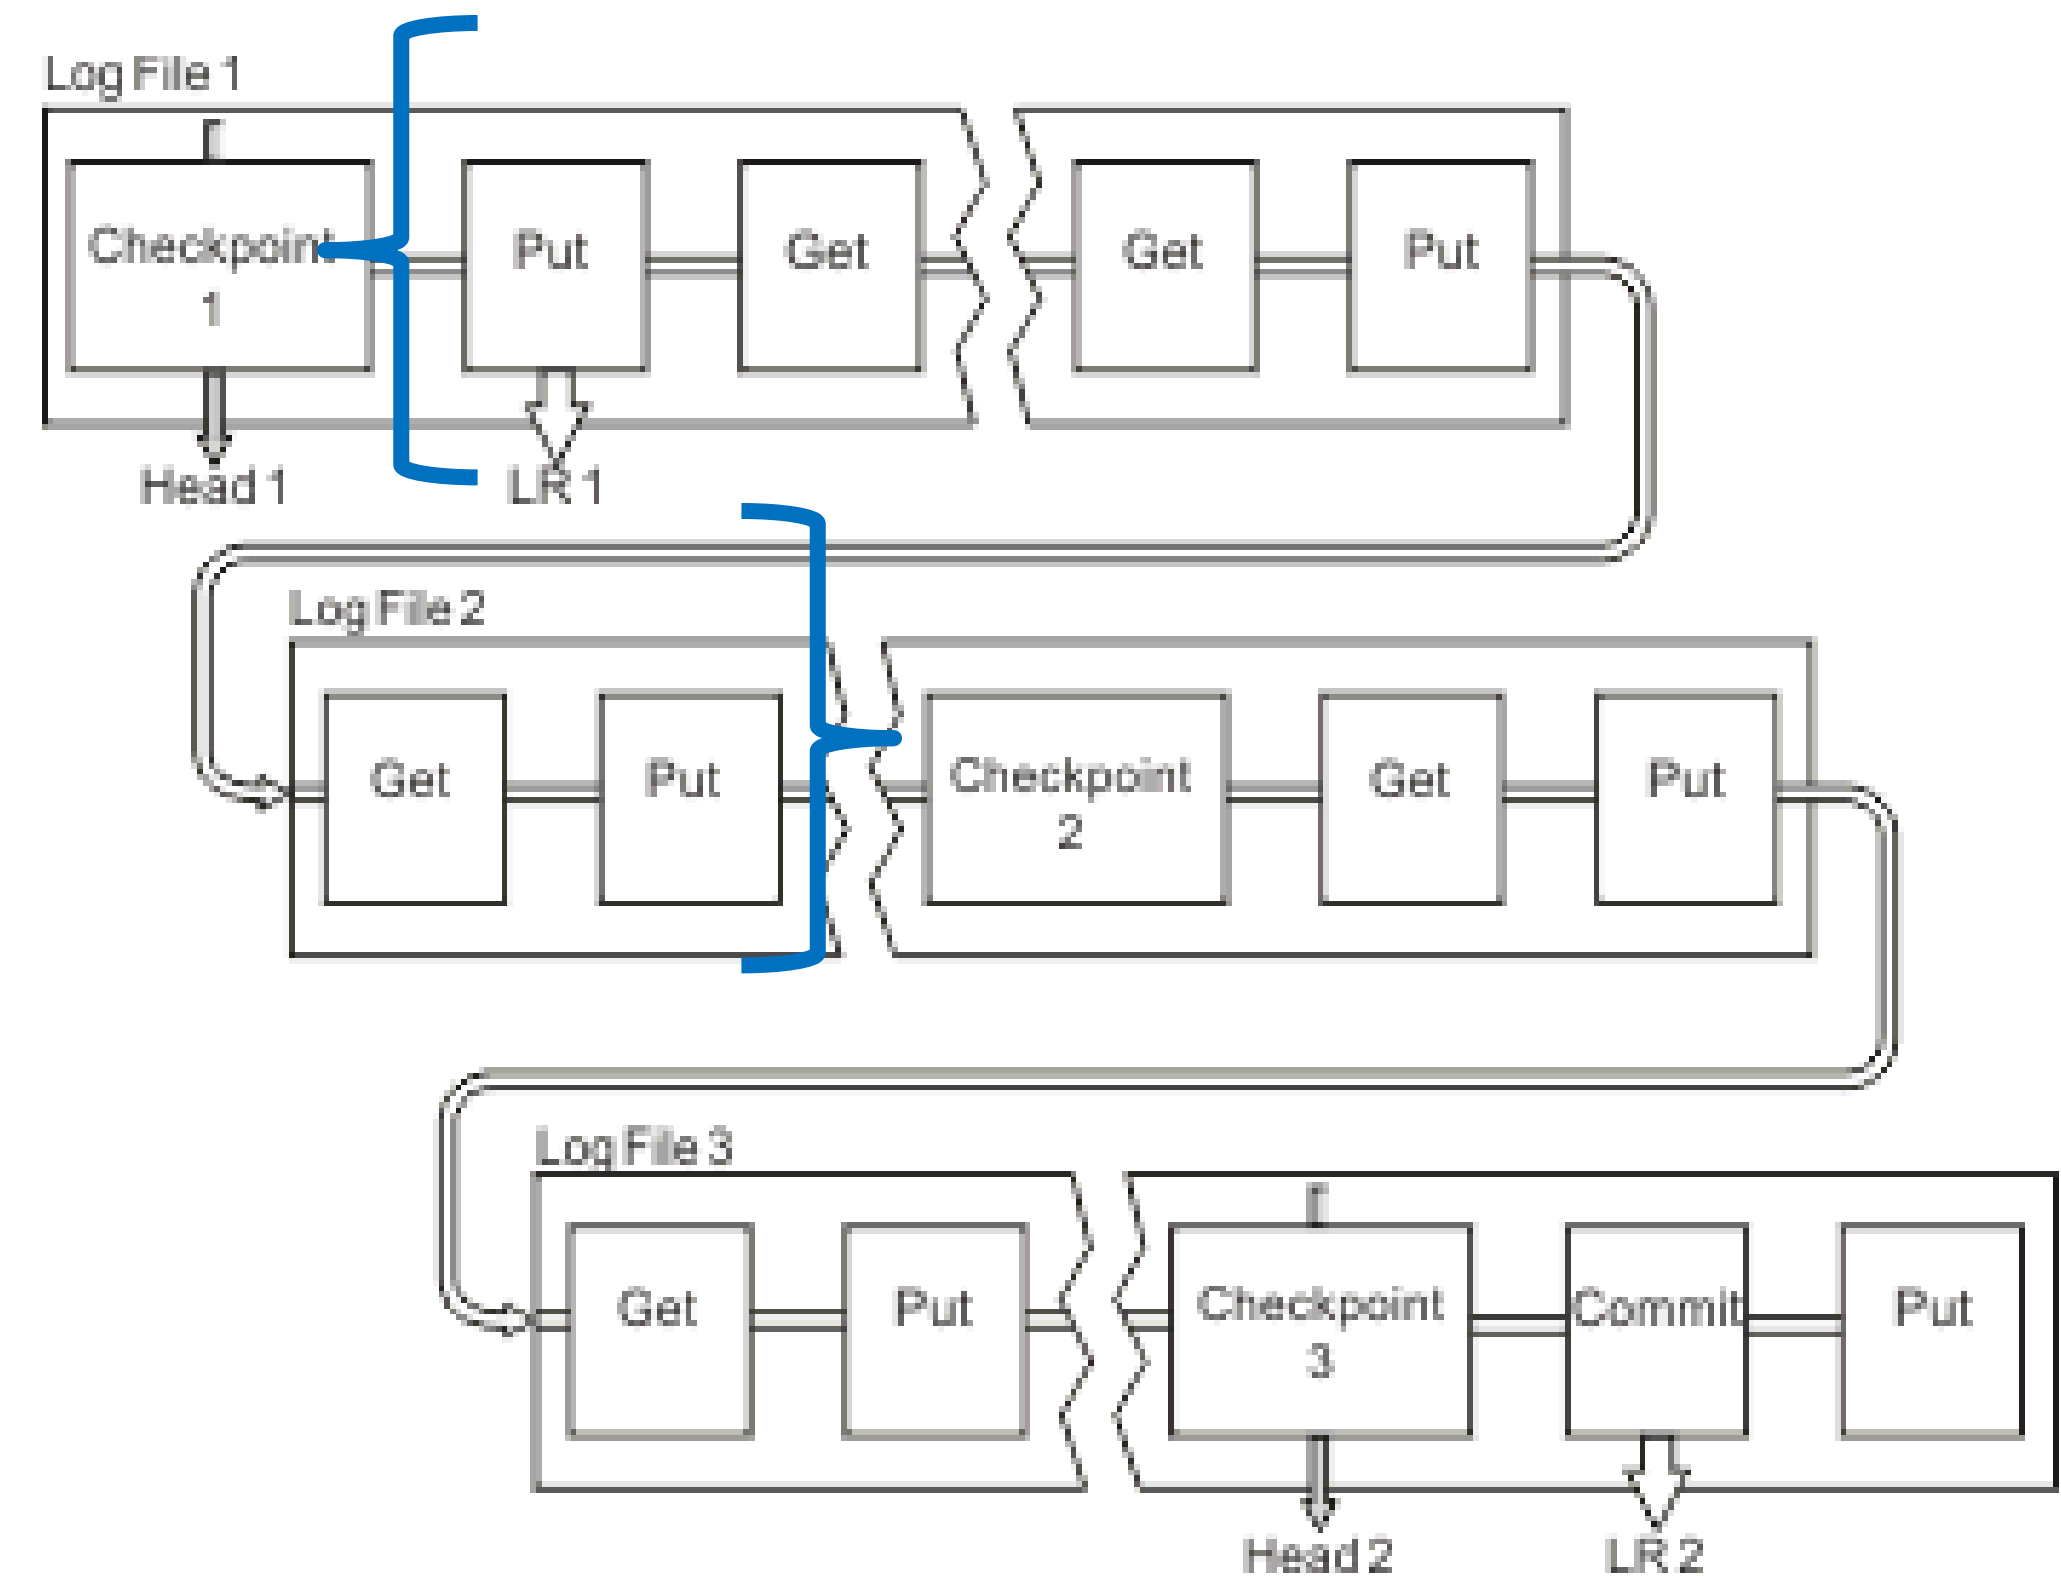
Traditionally, the act of logging persistent messages is the primary limiting factor of a queue manager. This is due to several constraints:

1. Single threaded work
2. Limited log buffers
3. Backing out work being done
4. MQ I/O limitations
5. Checkpointing



SMF-QJST.csv – Log Manager Statistics

- Log Reads
- Log Task Busy
- Checkpoints
- Unavailable buffer count
- I/O Max Duration Log Copy – how long it takes to do physical I/O



SMF-QMST.csv – Message Manager Statistics

- Records API requests
- Less useful for problem determination
- More useful for observing workload volume and peak periods
- Critical metrics here:
 - **MQPUT** – When an application made a put request onto the QM
 - **MQPUT1** – When an application has made a request to do an open and a put together

SMF-QTST.csv – Publications Statistics

- Publications with no Subscriber to Topic
- Total Publication Request Count
- Date Time LPAR QMgr
- Interval Duration
- Total Publication API Count
- Administrative Publications - Total Proxy Publications - Total
- High point of Publications Publications - Low Point
- Longest ET for publication in microseconds
- Total ET for publications in microseconds

SMF-QIST.csv – Data Manager Statistics

- Message Mgr MQGETs
- Data Mgr MQGETs
- MQGET Difference
- Message Mgr MQPUTs; Message Mgr MQPUT1s; Total Message Mgr Puts
- Data Mgr MQPUTs
- PUT Difference
- MSG_COUNT
- OBJECT CREATE; OBJECT DELETE;
- OBJECT PUT; OBJECT GET; OBJECT LOCATE


Implications:

GET difference can indicate scrolling

PUT difference can indicate put to waiting getter
performance enhancement


Object creates and deletes shows temporary
dynamic queue utilization

Positive put difference
MM puts > DM puts




Put to waiting getter
advantage

Negative put difference
DM puts > MM puts




Publications or generated
messages from triggering

Positive get difference
MM gets > DM gets



Target queue is empty
Not a problem

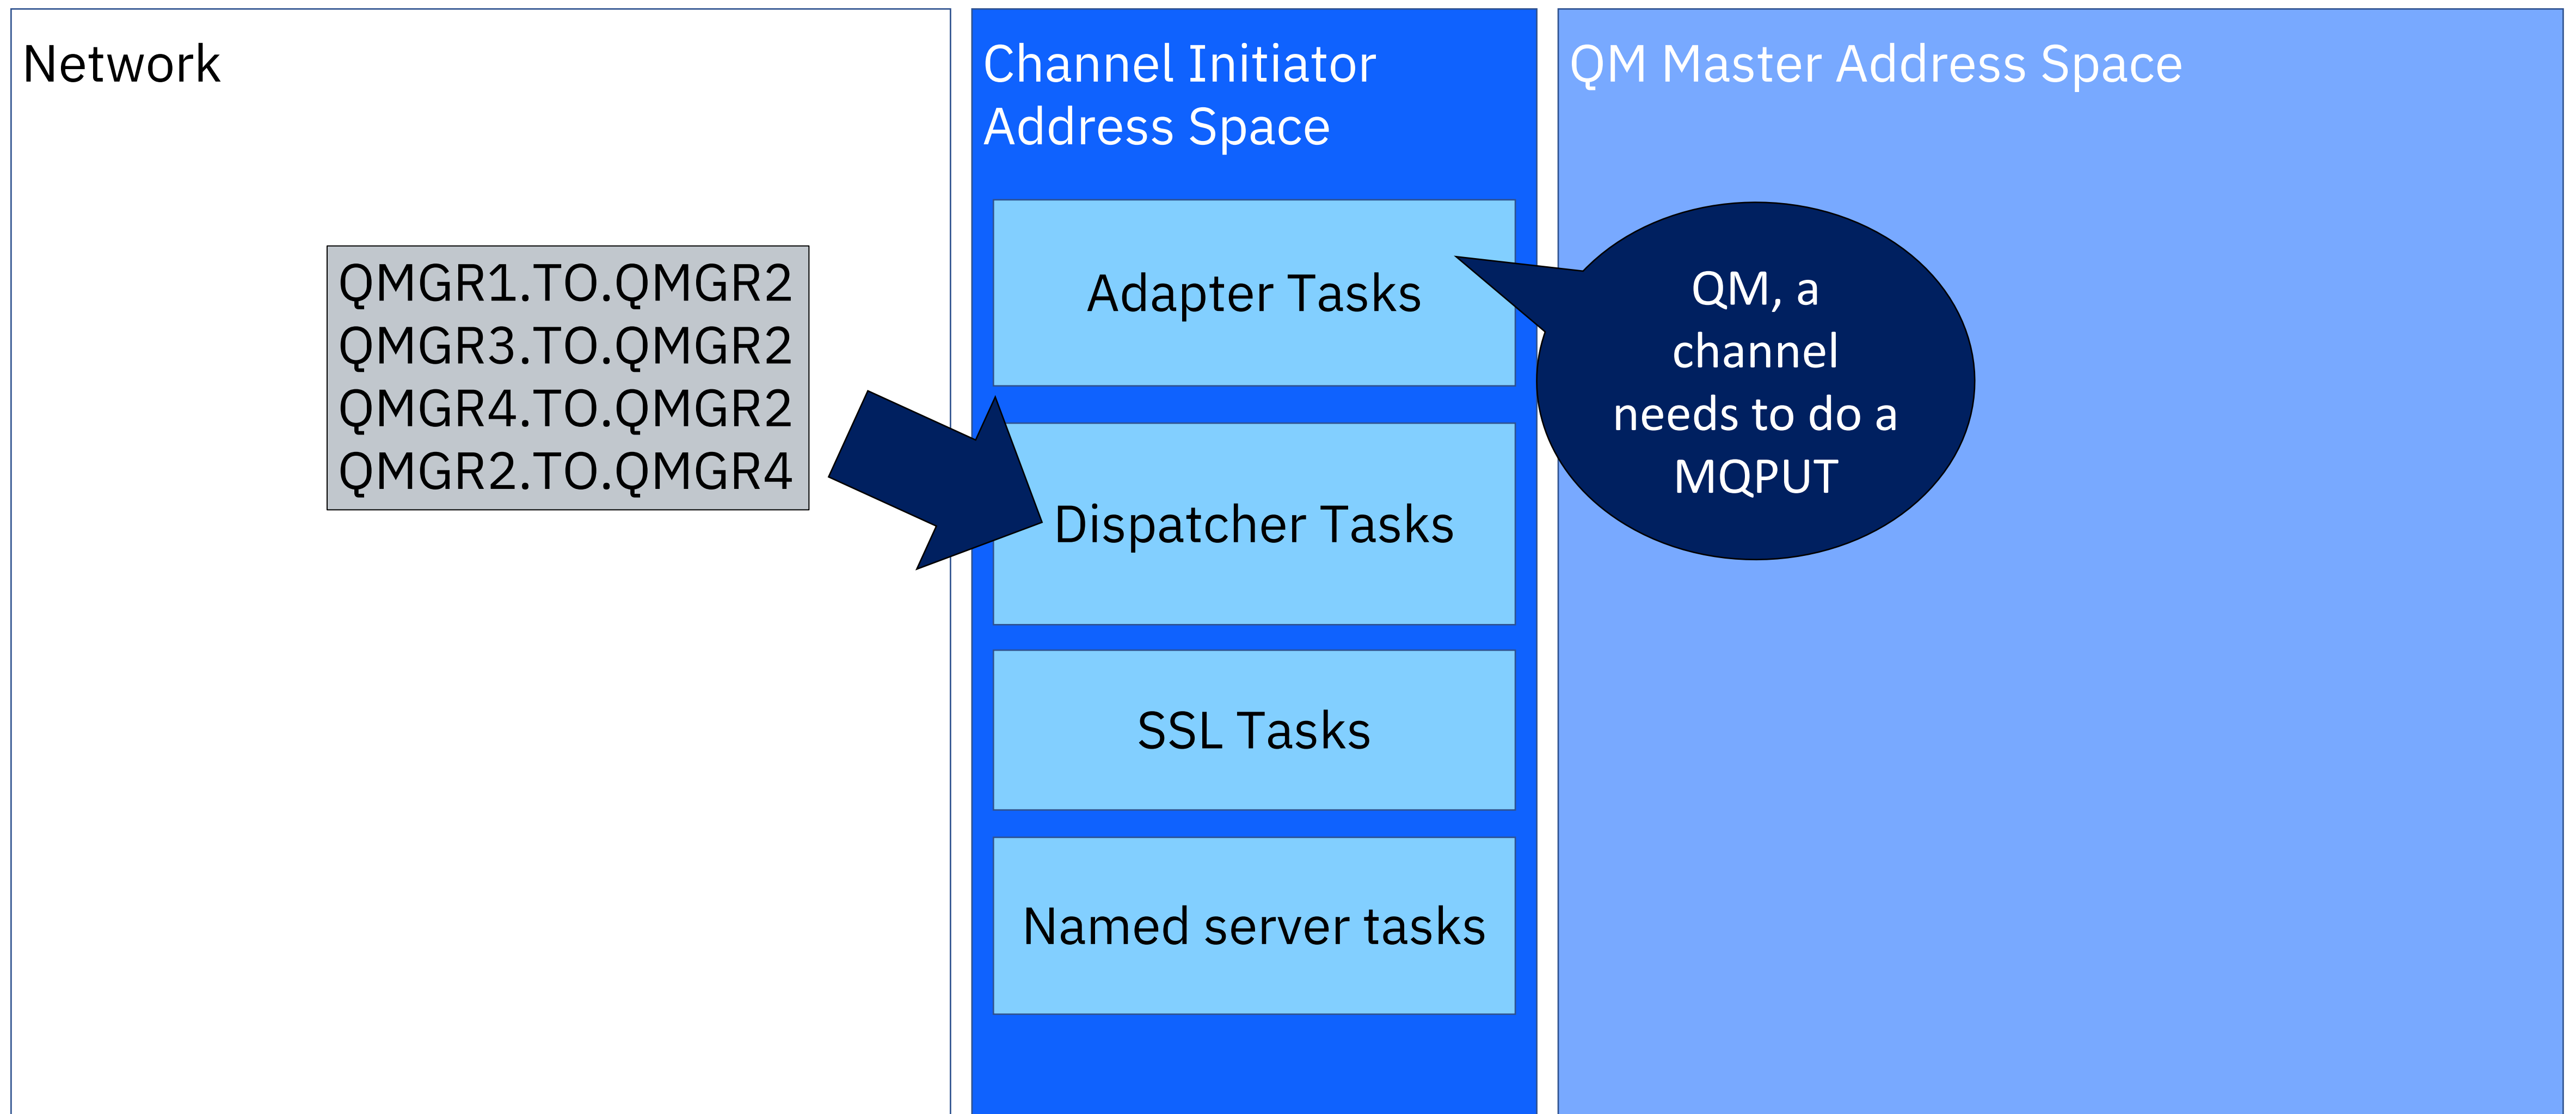
Negative get difference
DM gets > MM gets



May indicate scrolling
Look for skipped
messages

Interpreting SMF 115 channel information

CHINIT Address Space Structure



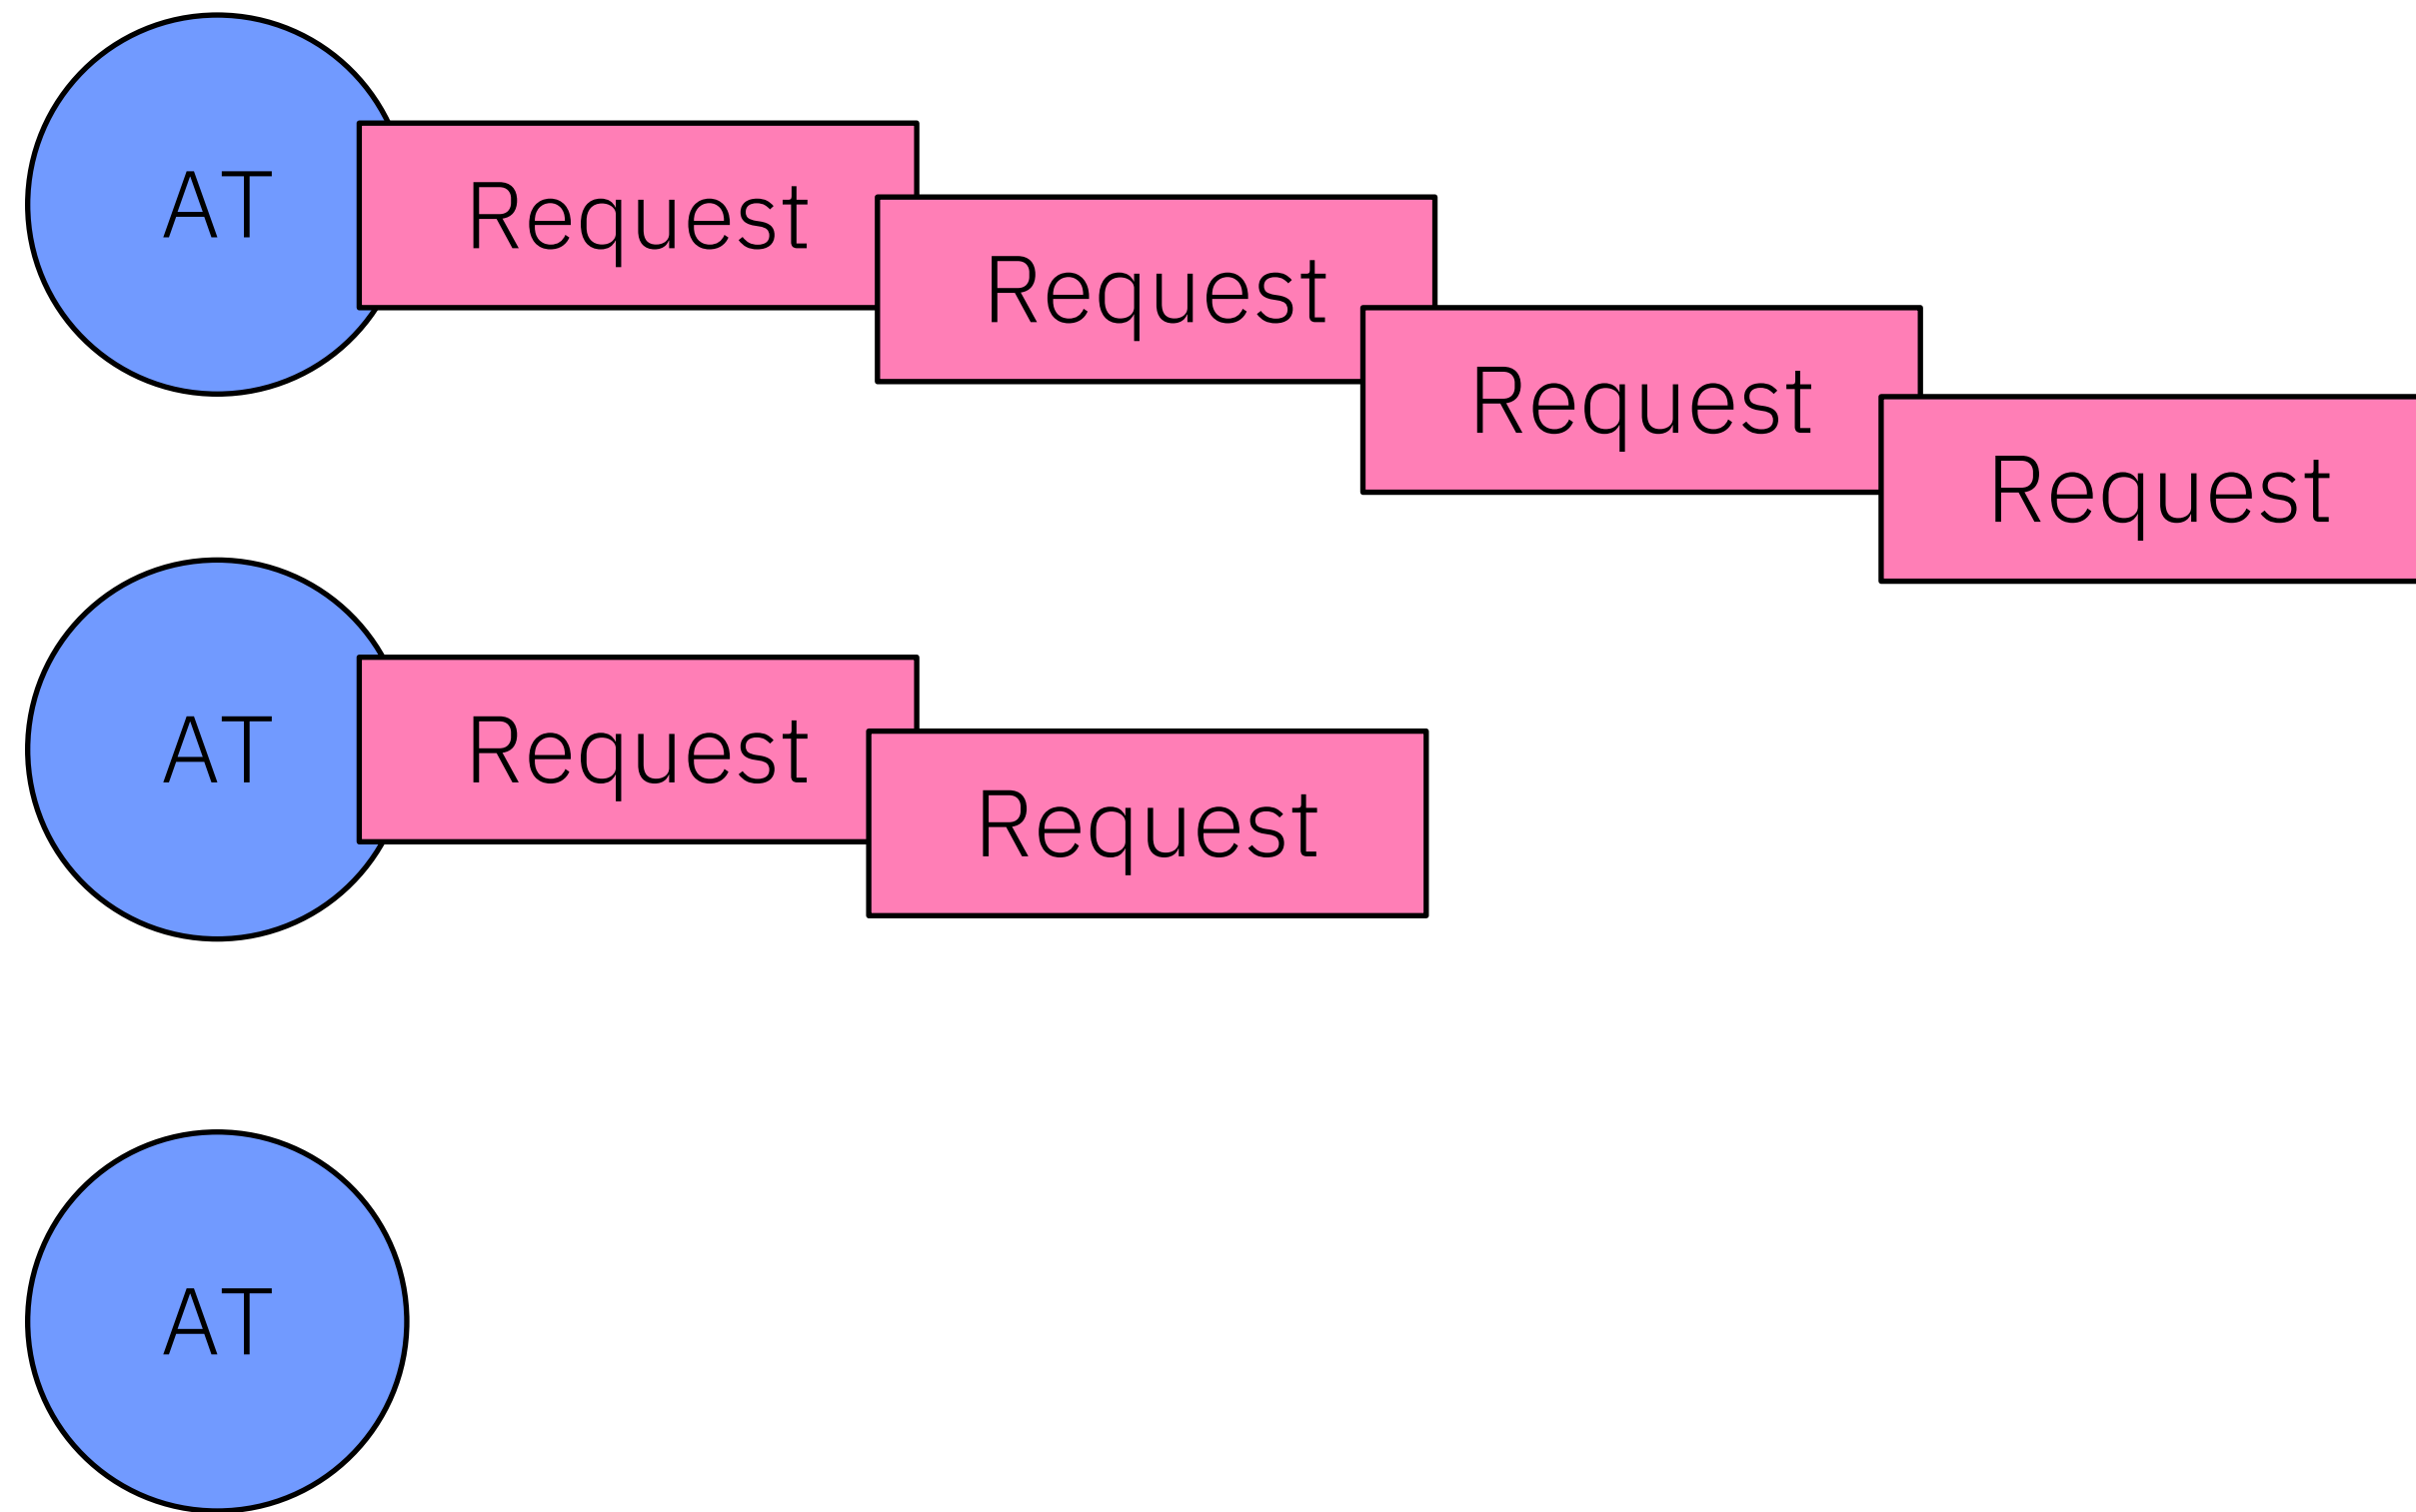
SMF-QCCT.csv – Channel Statistics

- CHINIT job name
- **QSG name if it is in a QSG**
- Peak number of current channels
- Peak number of active channels
- **MAXCHL - maximum permitted current channels**
- **ACTCHL - maximum permitted active channels**
- TCPCHL - maximum permitted
- TCP/IP channels
- LU62CHL - maximum permitted LU62 channels
- **Storage used by CHINIT**

SMF-QCTADP.csv – Adapter Task Statistics

- Date LPAR QMgr
- Adapter Task Number – always unique
- Total Requests for this Adapter Task
- Total CPU for this Adapter Task
- Total Elapsed Time for this Adapter Task – could be in a wait state
- Total Wait Time for this Adapter Task

How adapter tasks are assigned



SMF-QCTDSP.csv – Dispatcher Task Statistics


- Date LPAR QMgr
- Dispatcher Task Number
- Total Requests for this Dispatcher Task
- Total CPU for this Dispatcher Task
- Total Elapsed Time for this Dispatcher Task
- Total Wait Time for this Dispatcher Task
- Maximum channels for this dispatcher task for Day

How dispatcher tasks are assigned to channels


Scenario 1:

$200 / 5 = 40$ | $40 > 10$ (from the rule of 10) | SO, 10 channels will be assigned to each dispatcher task

Active Channel Max of 200



Dispatcher Tasks allocated



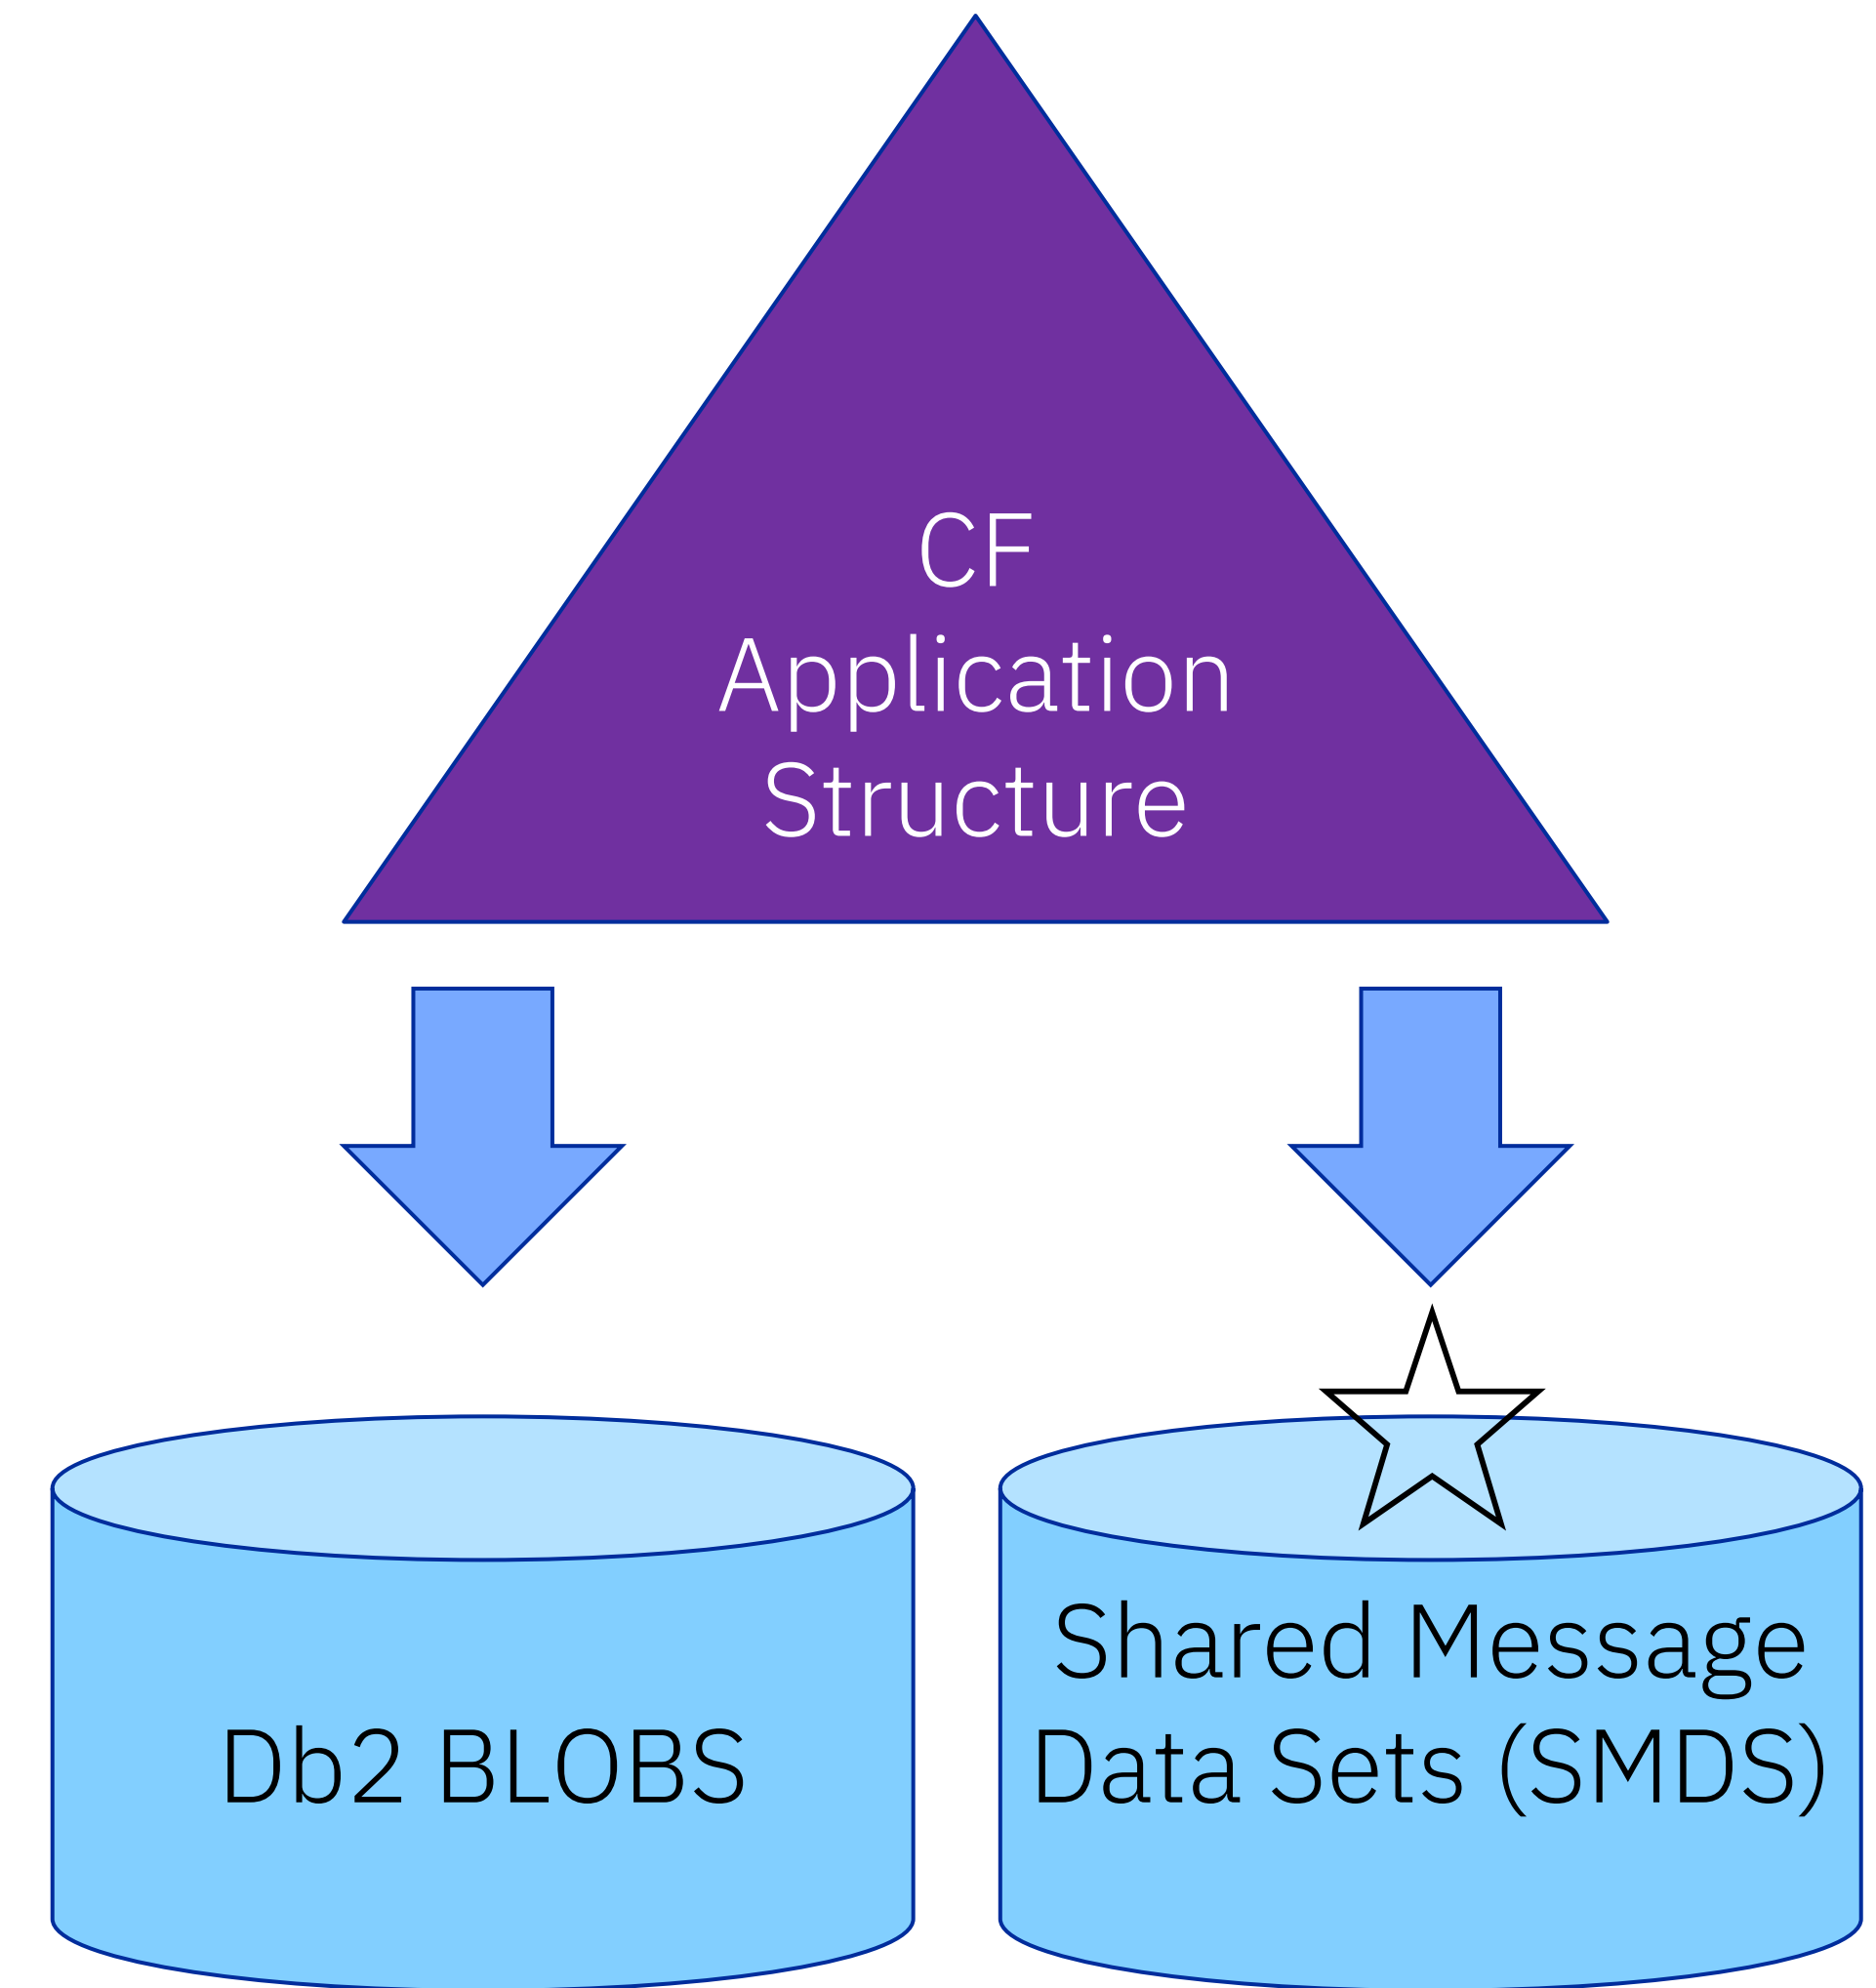
SMF statistics for shared queues

SMF-QEST.csv – Coupling Facility Statistics

- DATE TIME LPAR
- QMGR MQ_VERSION
- INTERVAL_START_DATE
- INTERVAL_START_TIME
- INTERVAL_DURATION
- **STRUCTURE_NAME**
- STRUCTURE_NUMBER
- IXLLSTE_CALLS
- IXLLSTM_CALLS
- IXLLSTE_REDRIVES
- IXLLSTM_REDRIVES
- STRUCTURE_FULL

SMF-Q5ST.csv – BLOB Statistics

- Date Time LPAR QMgr
- BLOB SELECTS
- BLOB INSERTS
- BLOB UPDATES
- BLOB DELETES
- BLOB LISTS




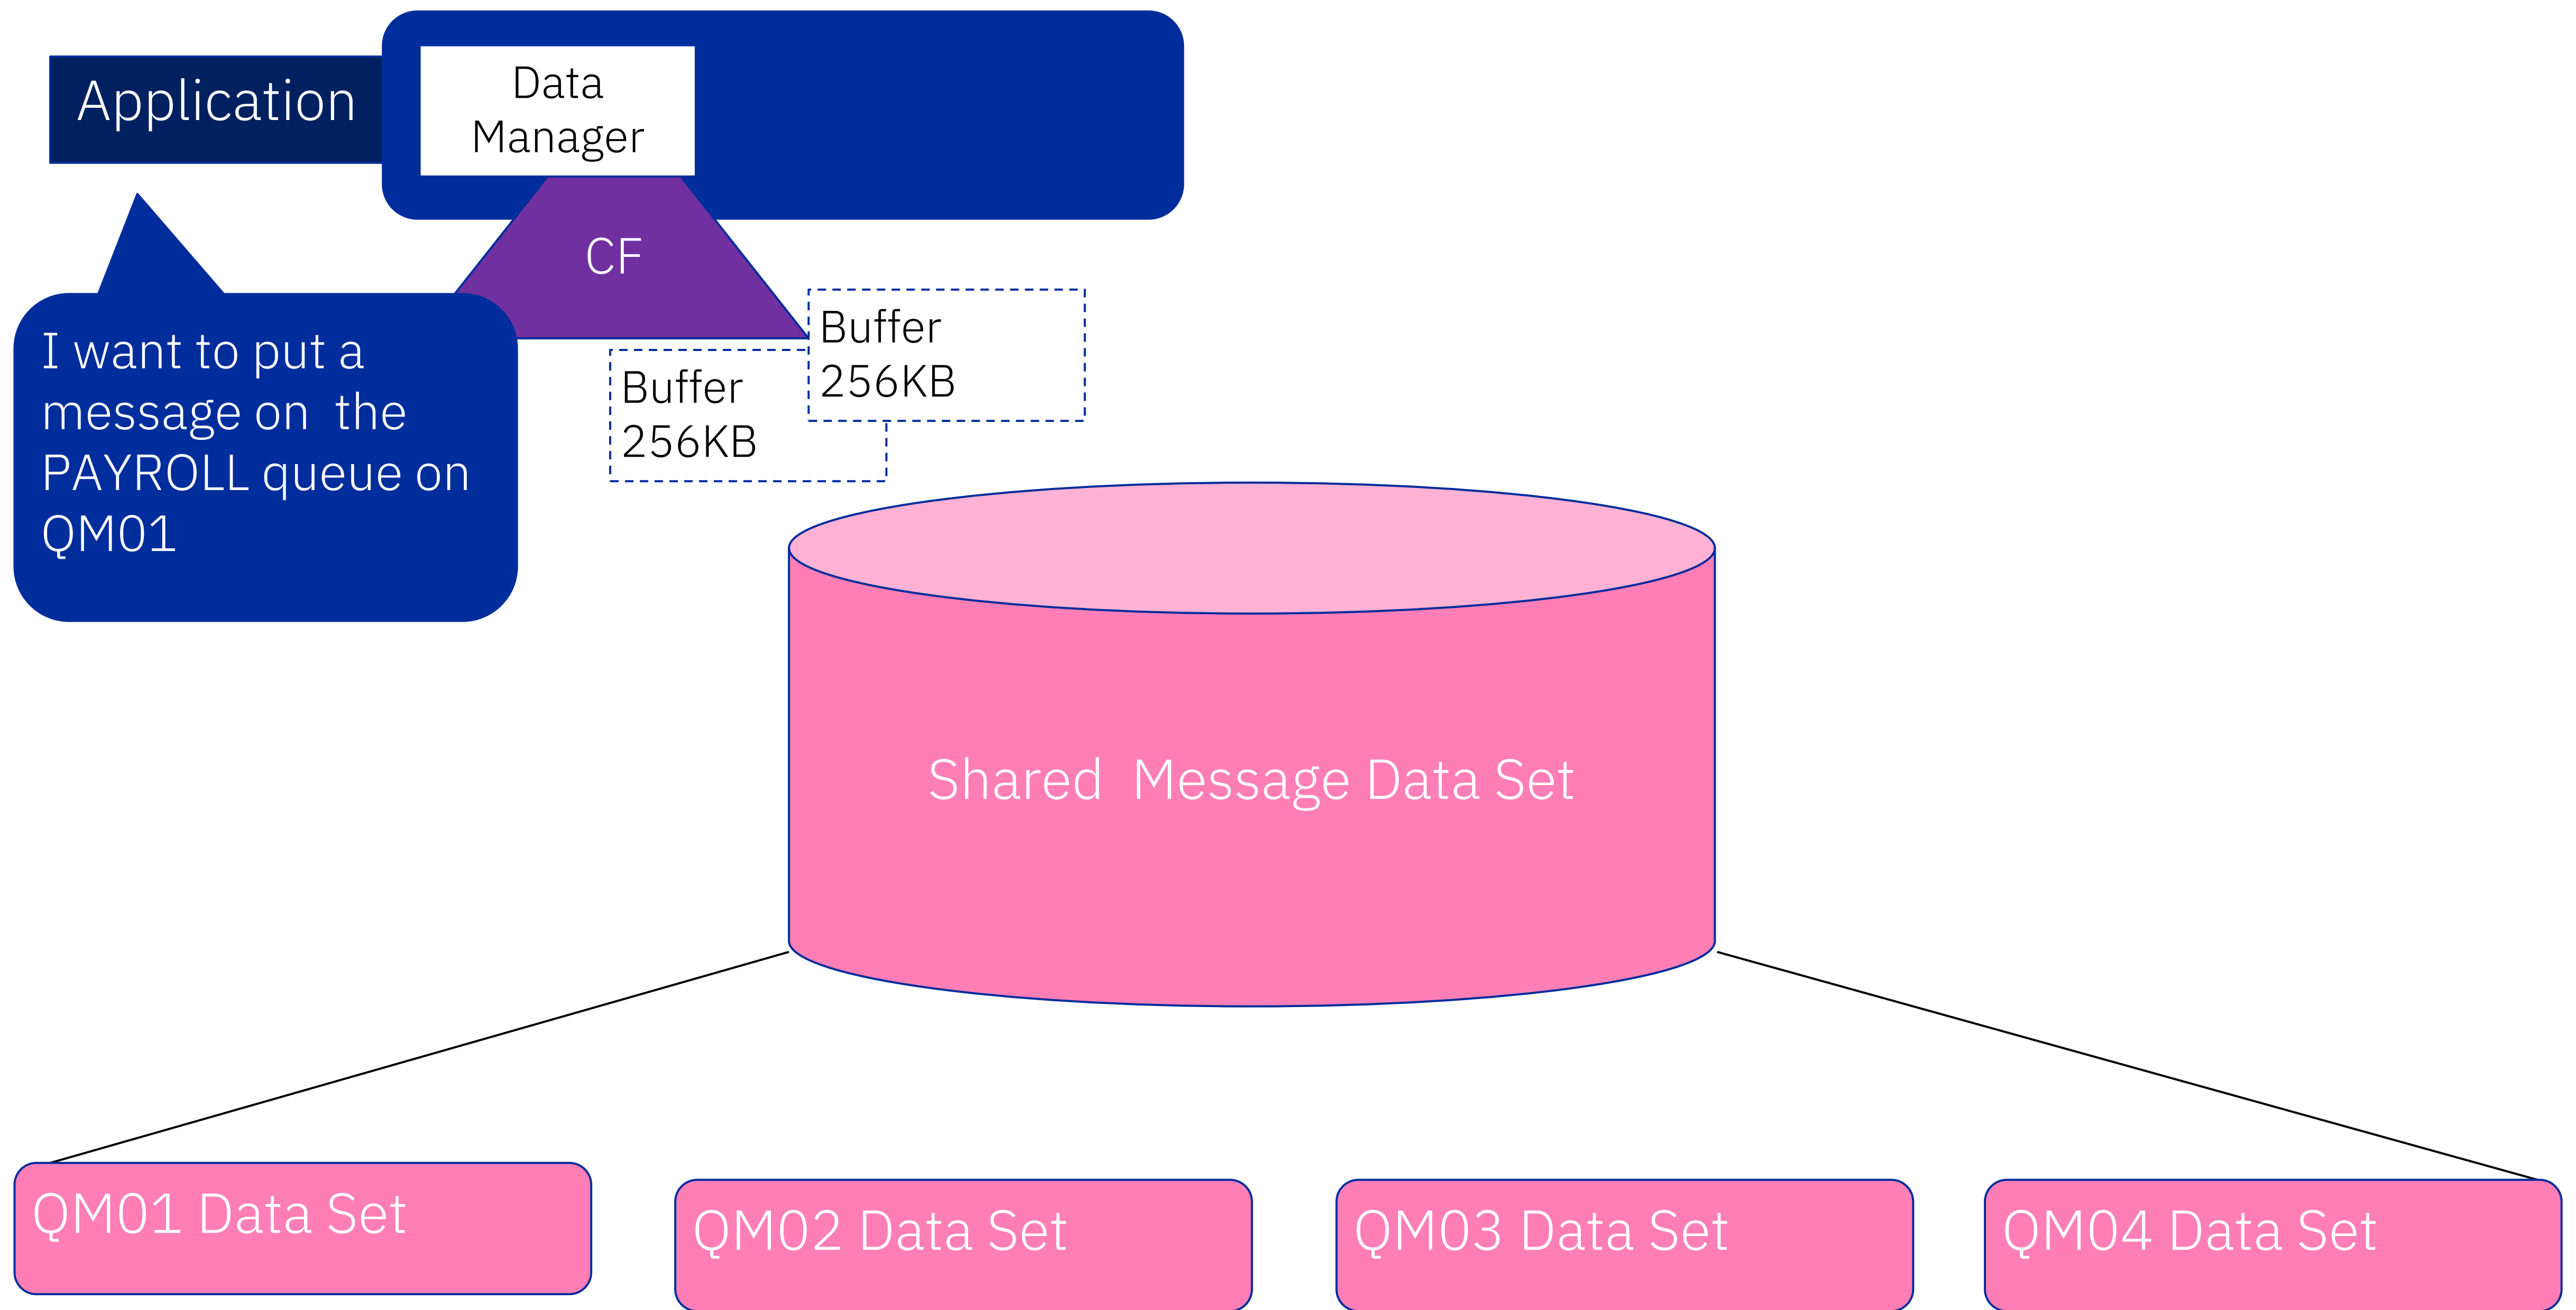
Single and multi-entry Redrives

Redrives, both single and multiple entry, are when the coupling facility detects that we need a larger buffer for returning a message.

Once queue manager gets the notification that it needs a larger buffer, it will use the buffer size necessary.

Using redrives allows for more flexibility in message sizes





SMF 116 Accounting Data

SMF-QCST.csv – Channel Accounting

- QMGR
- CHL_NAME
- TYPE
- CONNECTION_NAME
- START_DATE
- DISPOSITION
- COLLECTION_TIME_DATE DATE
- COLLECTION_TIME_TIME CHAR(19)
- STATUS
- STAT_SETTING
- And more....

SMF-WTID.csv – Task ID Accounting

- DATE TIME LPAR QMGR
- MQ_VERSION
- WTAS_CORRELATOR
- APPL_TYPE
- CONNECTION_NAME
- OPERATOR_ID
- NID
- CORREL_HEX
- CORREL_CHAR
- UOW_ID
- ACCOUNTING_TOKEN
- CHANNEL_NAME
- CHANNEL_CONNECTION_NAME
- CONTEXT_TOKEN
- MVS_USERID

SMF-WTAS.csv – Task Accounting

- QMgr Correl
 - Longest_Latch
 - Max Latch Wait Microseconds
 - Max Latch Wait Type
 - Start Date
 - Start Time
- QMgr
 - Correl
 - Type 11 Latch Wait Time (Over 5000 mics.)
 - Type 11 Wait Count
 - Task Start Date
 - Task Start Time

Latching types

Type Number	What it means
11	Typically - ensuring serialization across API requests on a queue
12	Buffer pool latching
15	First-open Last closed effect on shared queues
16	Serializing access to a PSID
19	Serializing access to a buffer page
21	Log buffer latching
23	Serializing access to the BSDS
24	Serializing access across waiting getters
25	To serialize access to a particular lock AND for creating Accounting
30	Task switching and security checks
31	Often, queue scanning when a queue is not indexed, occasionally when checking security and freeing storage
32	Serializing requests to get or free storage

SMF-WQ.csv – Task Queue Accounting

- Base Queue Name
- Open Name
- Bufferpool ID
- Pageset ID
- Coupling Facility Structure
- Total Opens
- Total Open Elapsed Time
- Total Open CPU Time
- Total Closes
- Total Close Elapsed Time
- Total Close CPU Time
- Total MQGET Requests
- Total MQGET Elapsed Time
- Total MQGET CPU Time
- Total Valid Gets
- Total Get Bytes
- Total Persistent MQGETs
- Total Messages Skipped
- Total Messages Expired
- Total MQPUT Requests
- Total MQPUT Elapsed Time
- Total MQPUT CPU Time
- Total MQPUT1 Requests
- Total MQPUT1 Elapsed Time
- Total MQPUT1 CPU Time
- Total Valid MQPUTs
- Total Bytes Put
- Total Puts to Waiting Getter
- Total Put1s to Waiting Getter
- Total Generated Messages
- Total Persistent MQPUTs
- Total Persistent MQPUT1s
- Max Depth on Queue
- Max Time on Queue
- Min Time on Queue
- Total Inquiries
- Total Sets
- Get percent unfulfilled
- Index Type

To recap...

SMF 115 data

SMF 115 for shared
queues specifically

SMF 116 data

Concept check

What is a redrive?

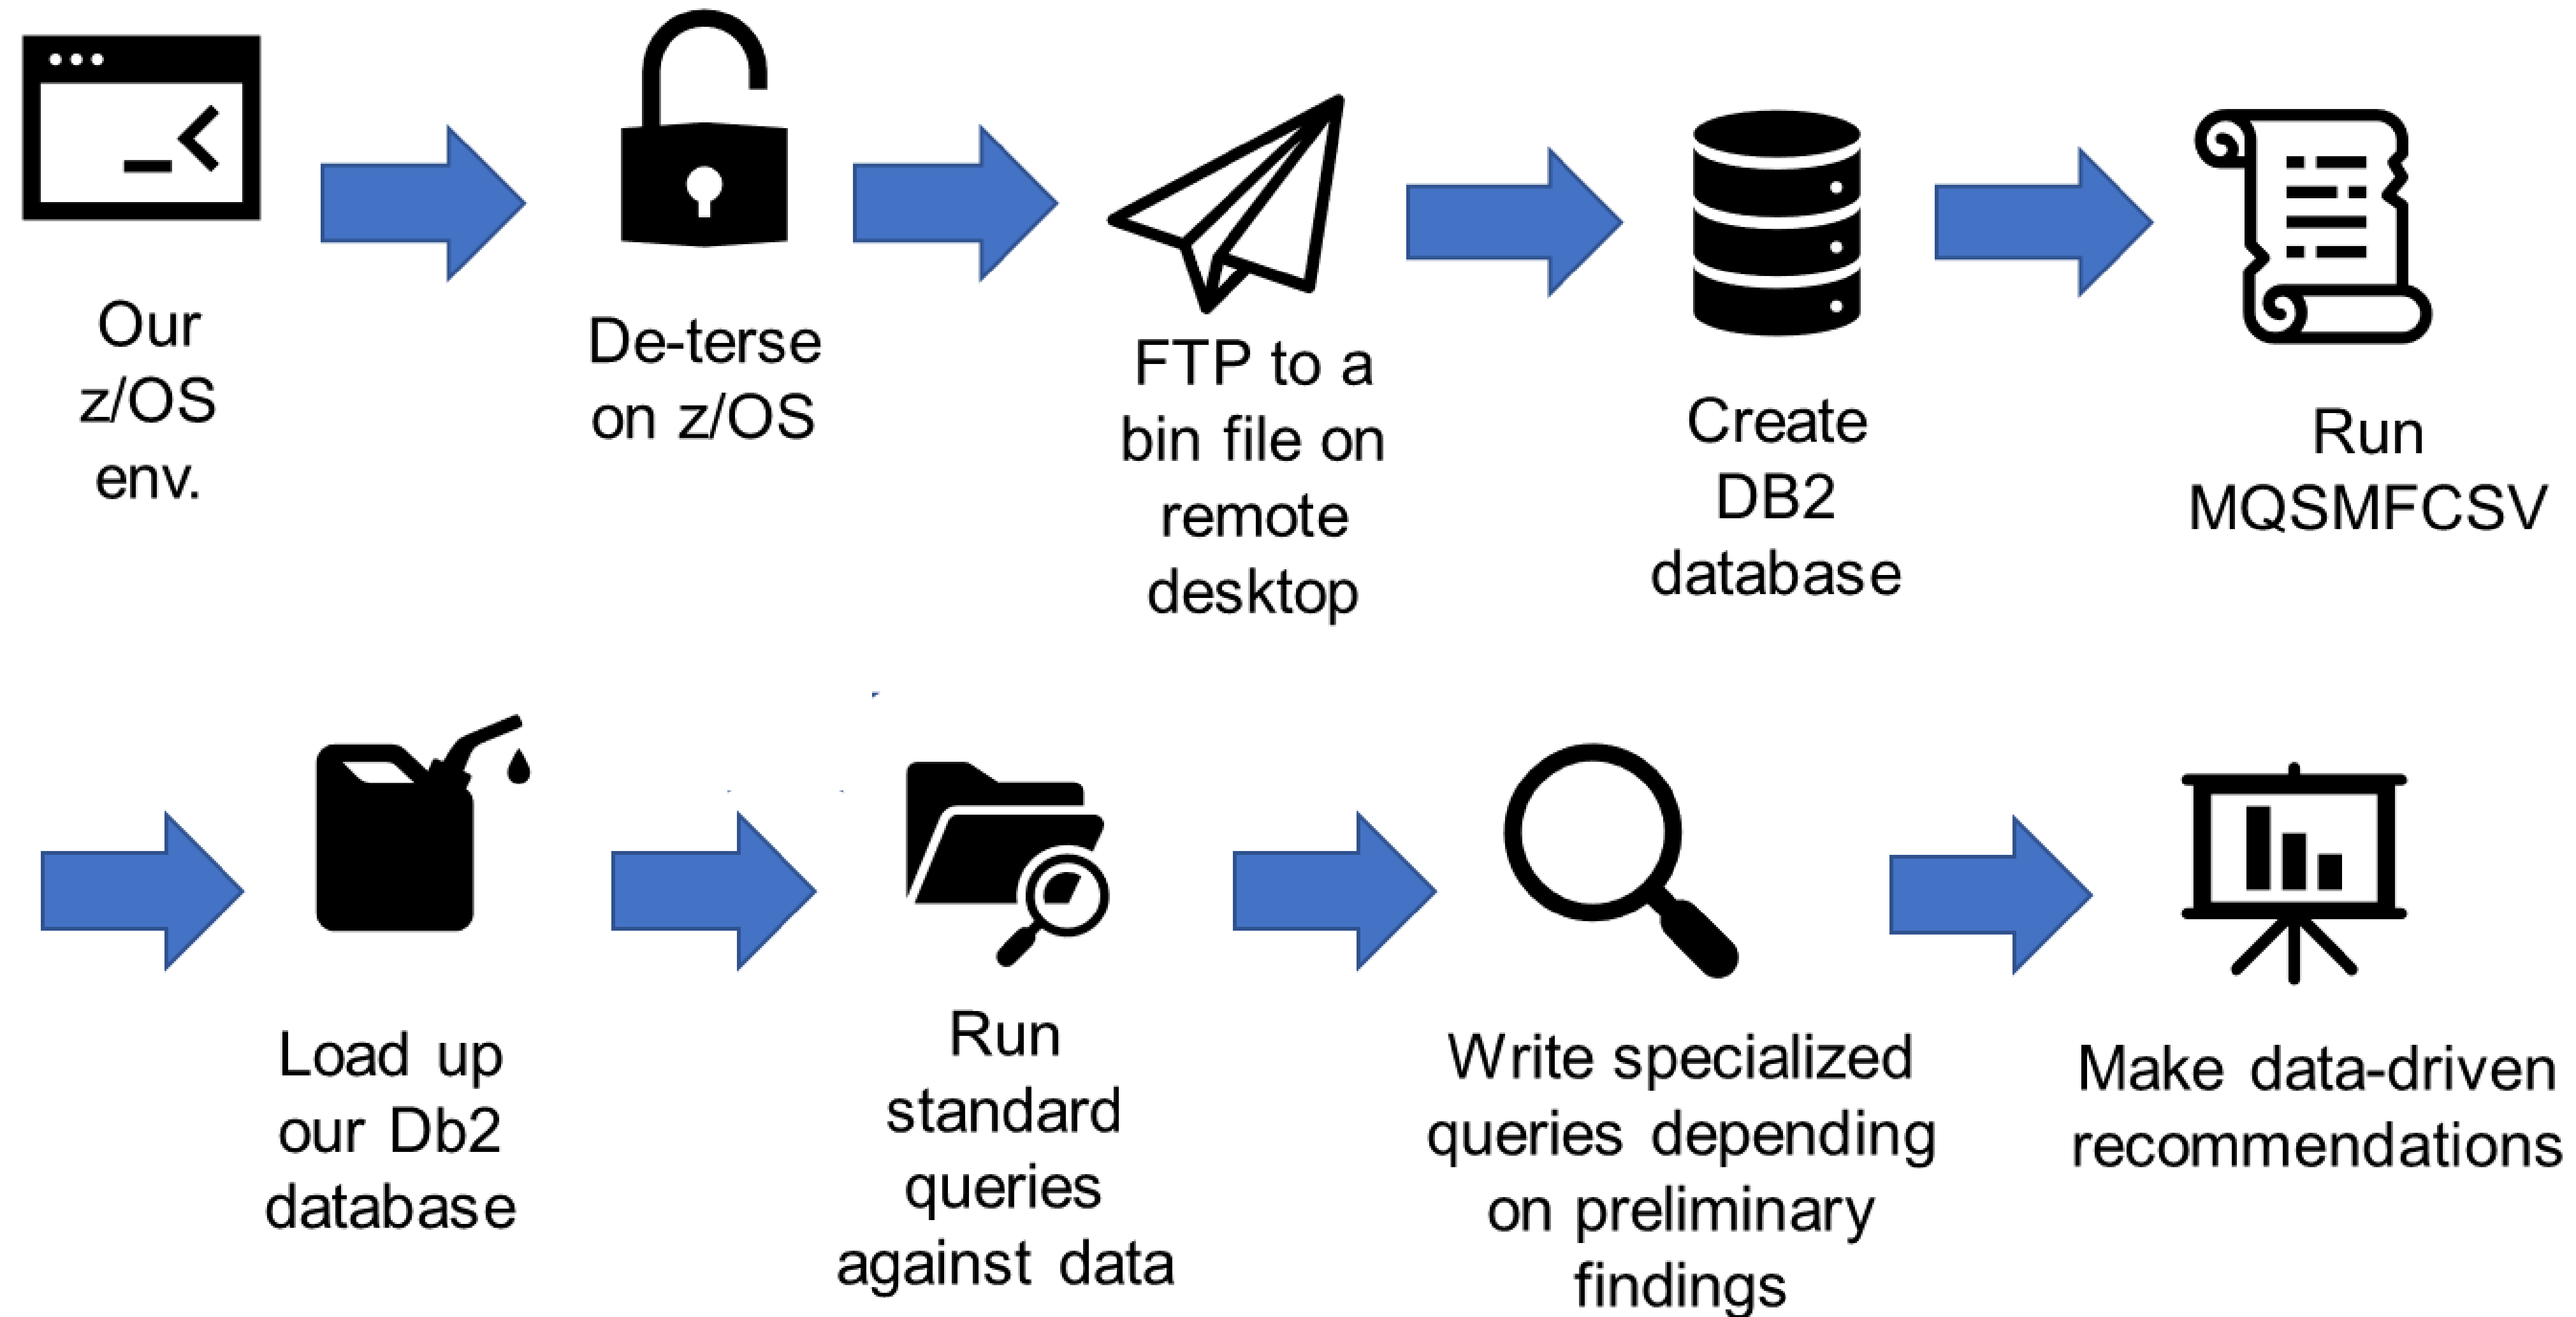
- 1) When you have to drive the same route twice
- 2) When the queue manager makes a call to the CF and the CF responds saying you need a larger buffer
- 3) When an adapter and dispatcher task interfere with one another

What is the difference between assignment for adapter and dispatcher tasks?

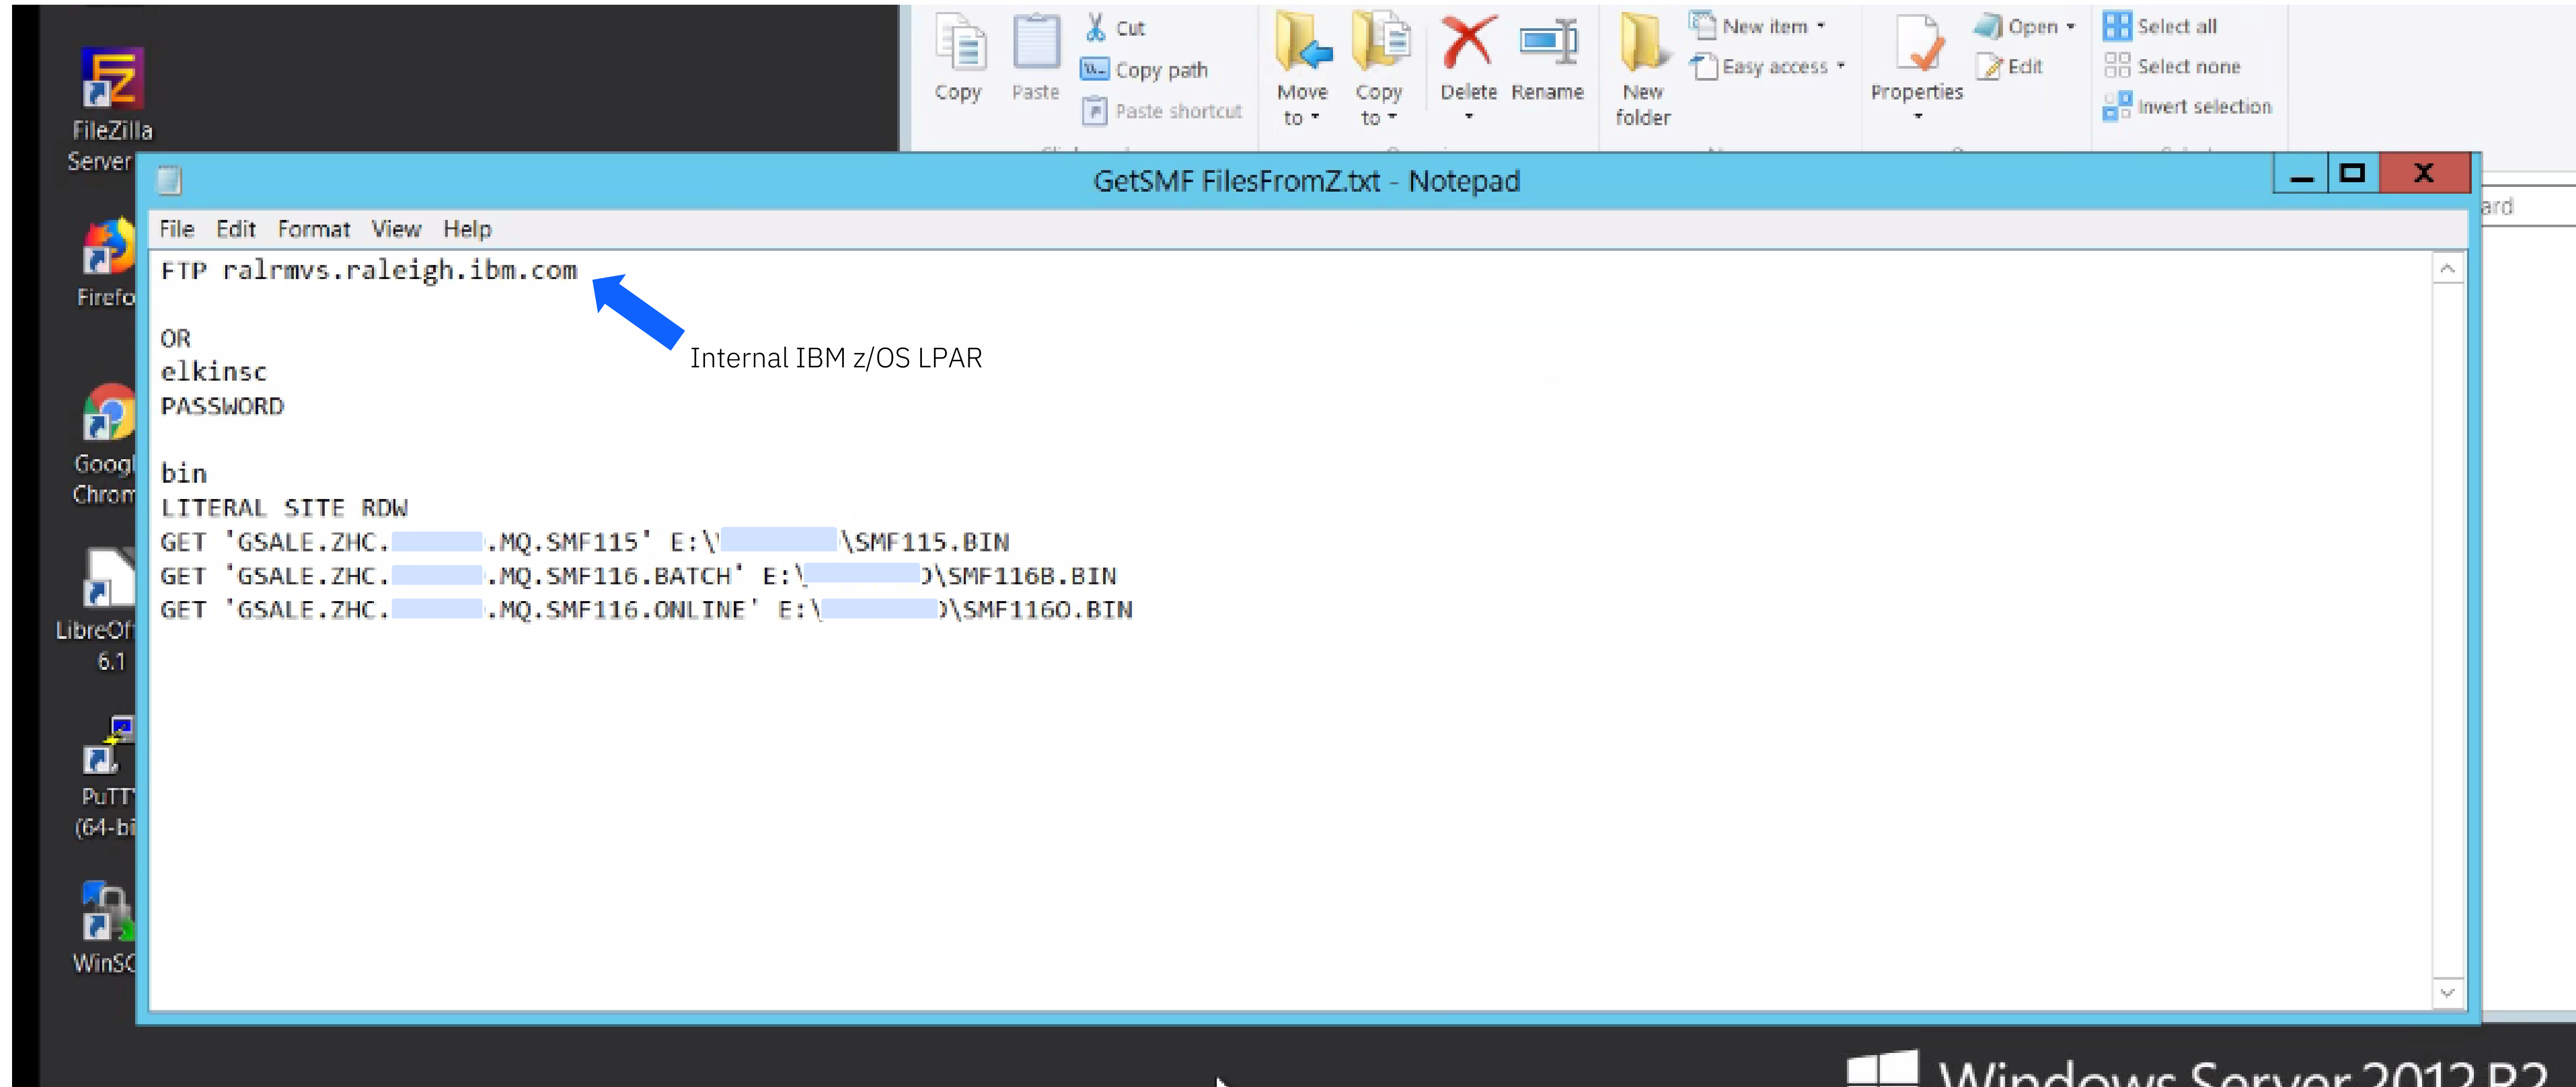
- 1) Dispatcher tasks are worker tasks for channel code to run on. Adapter task issue MQ API calls on behalf of channels
- 2) Dispatcher tasks connect the queue manager to the CHINIT address space. Adapter tasks connect the network to the CHINIT address space.
- 3) Adapter tasks can become more evenly distributed by reducing the MAXACTCHL.

Demonstration

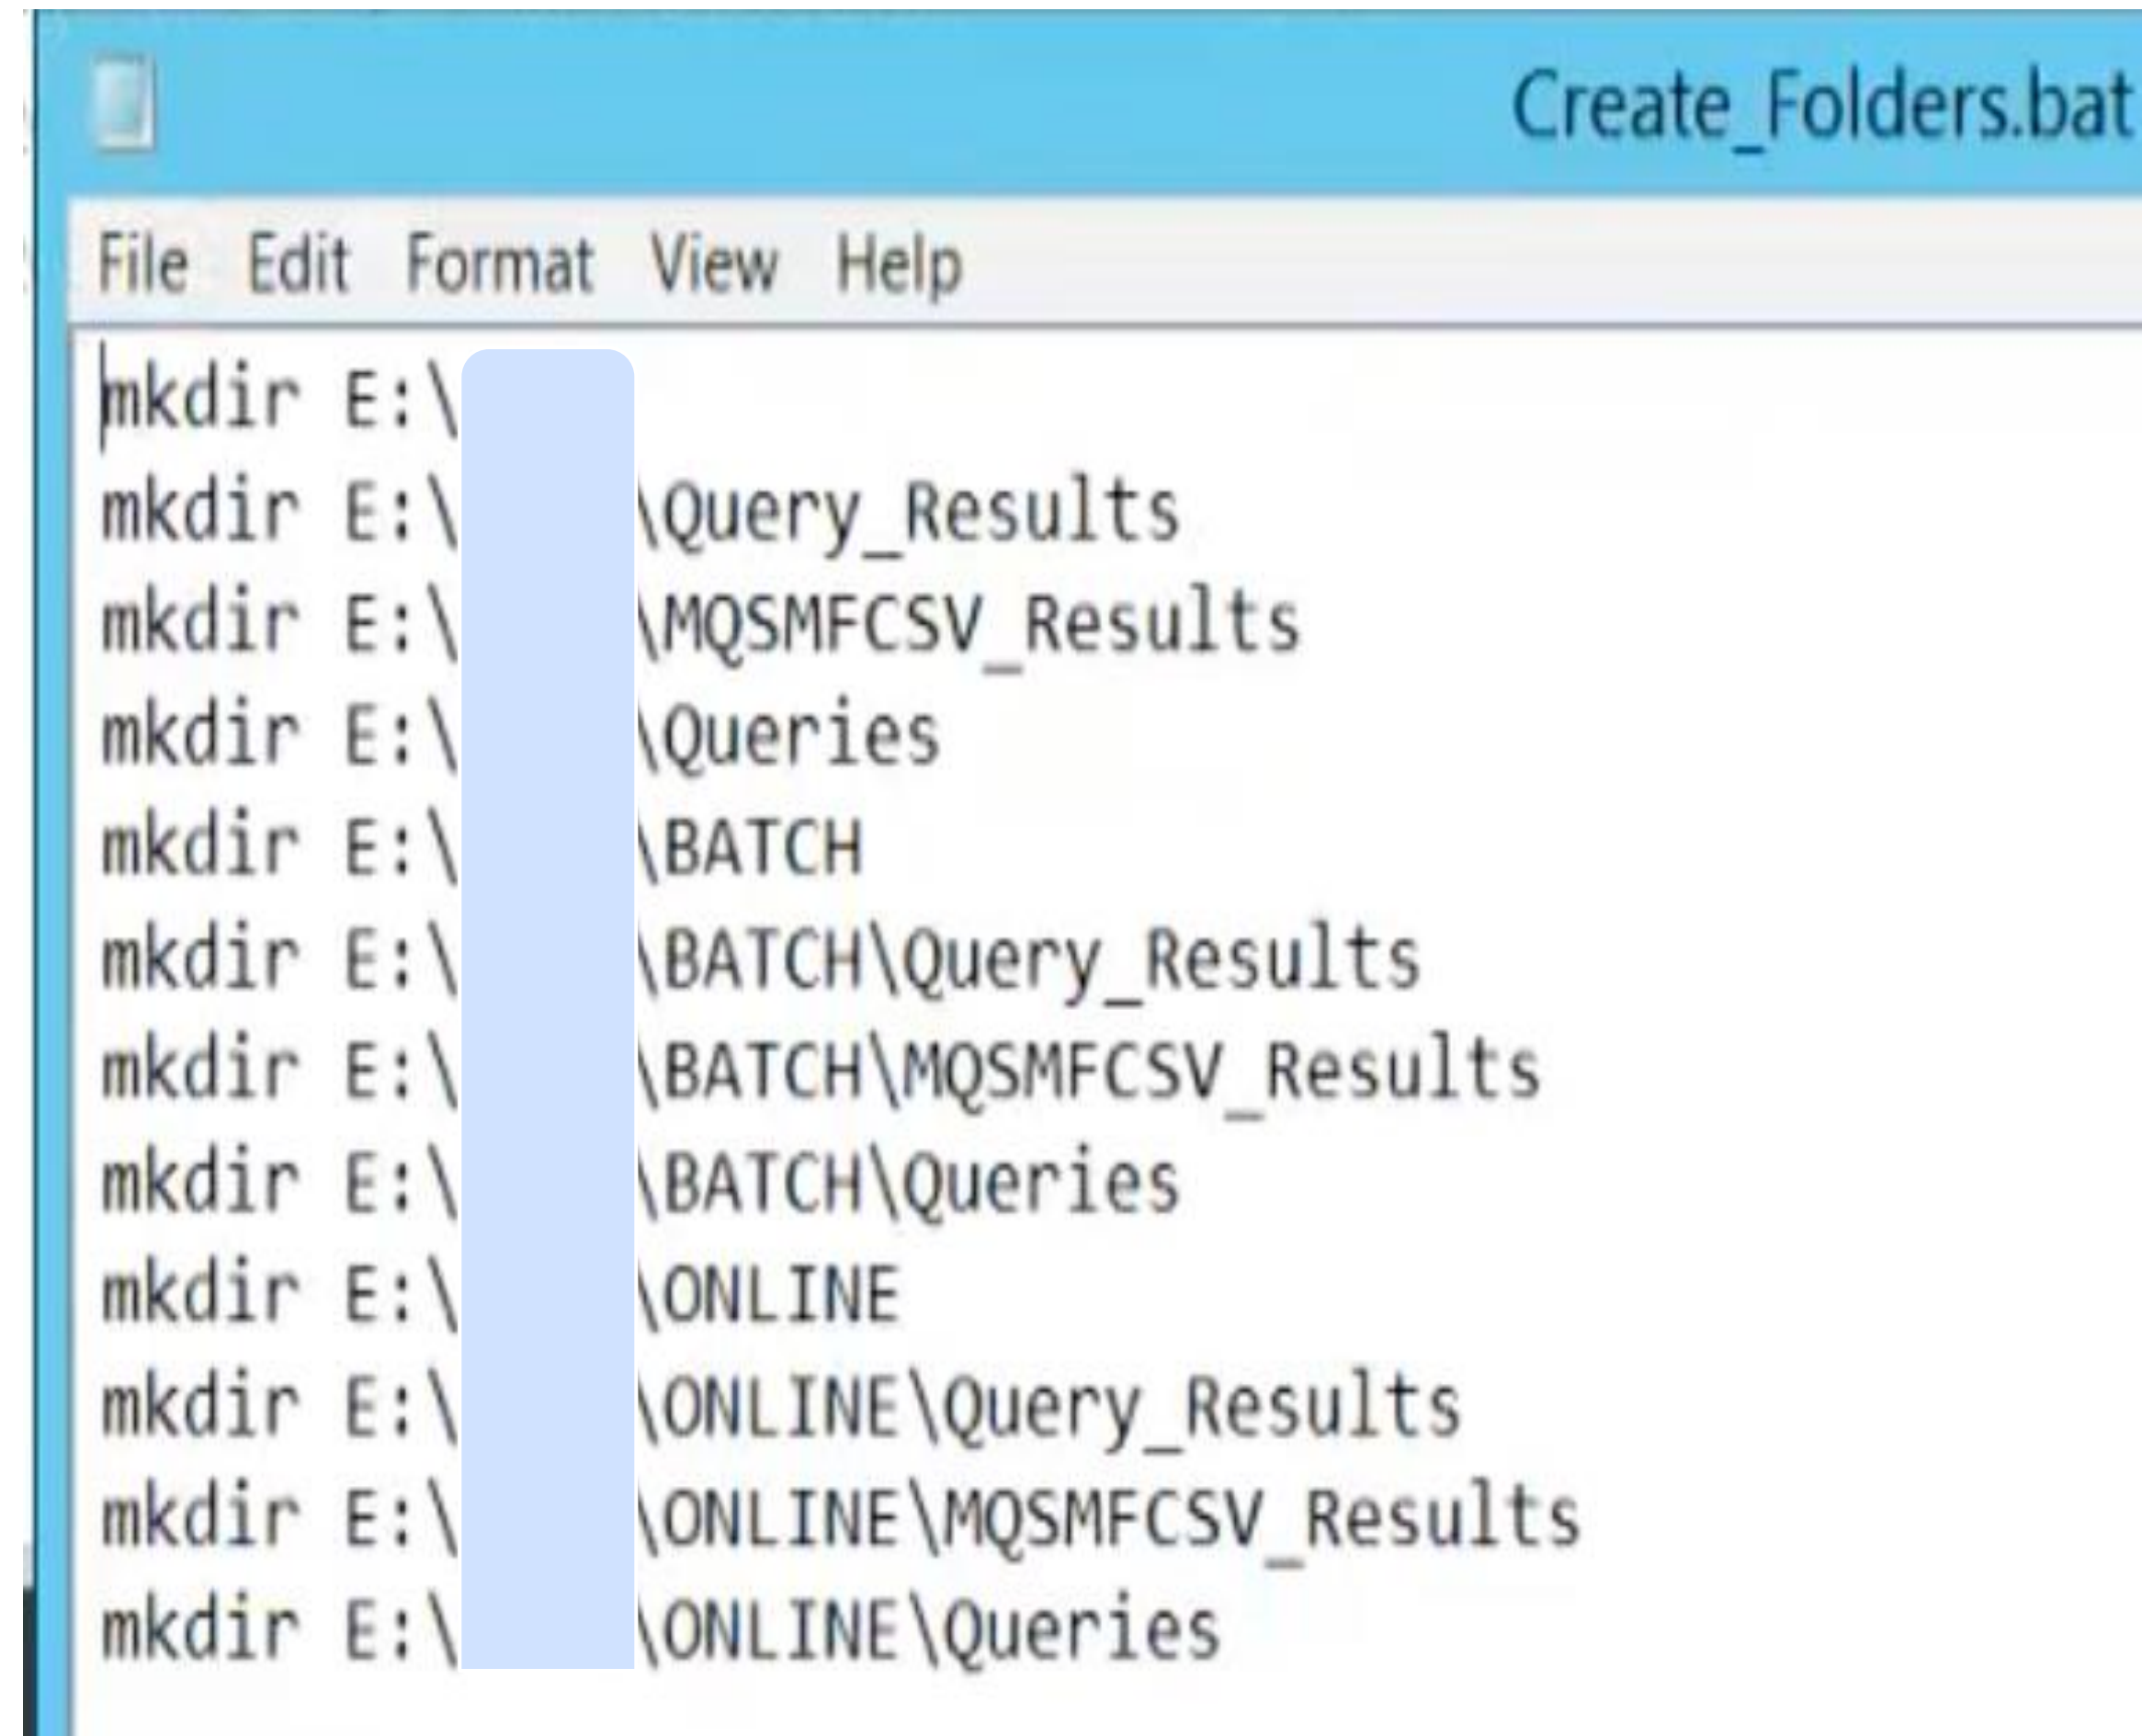
At a bird's eye



Pull data down from z/OS



Create file system



```
Create_Folders.bat
File Edit Format View Help
mkdir E:\
mkdir E:\Query_Results
mkdir E:\MQSMFCSV_Results
mkdir E:\Queries
mkdir E:\BATCH
mkdir E:\BATCH\Query_Results
mkdir E:\BATCH\MQSMFCSV_Results
mkdir E:\BATCH\Queries
mkdir E:\ONLINE
mkdir E:\ONLINE\Query_Results
mkdir E:\ONLINE\MQSMFCSV_Results
mkdir E:\ONLINE\Queries
```

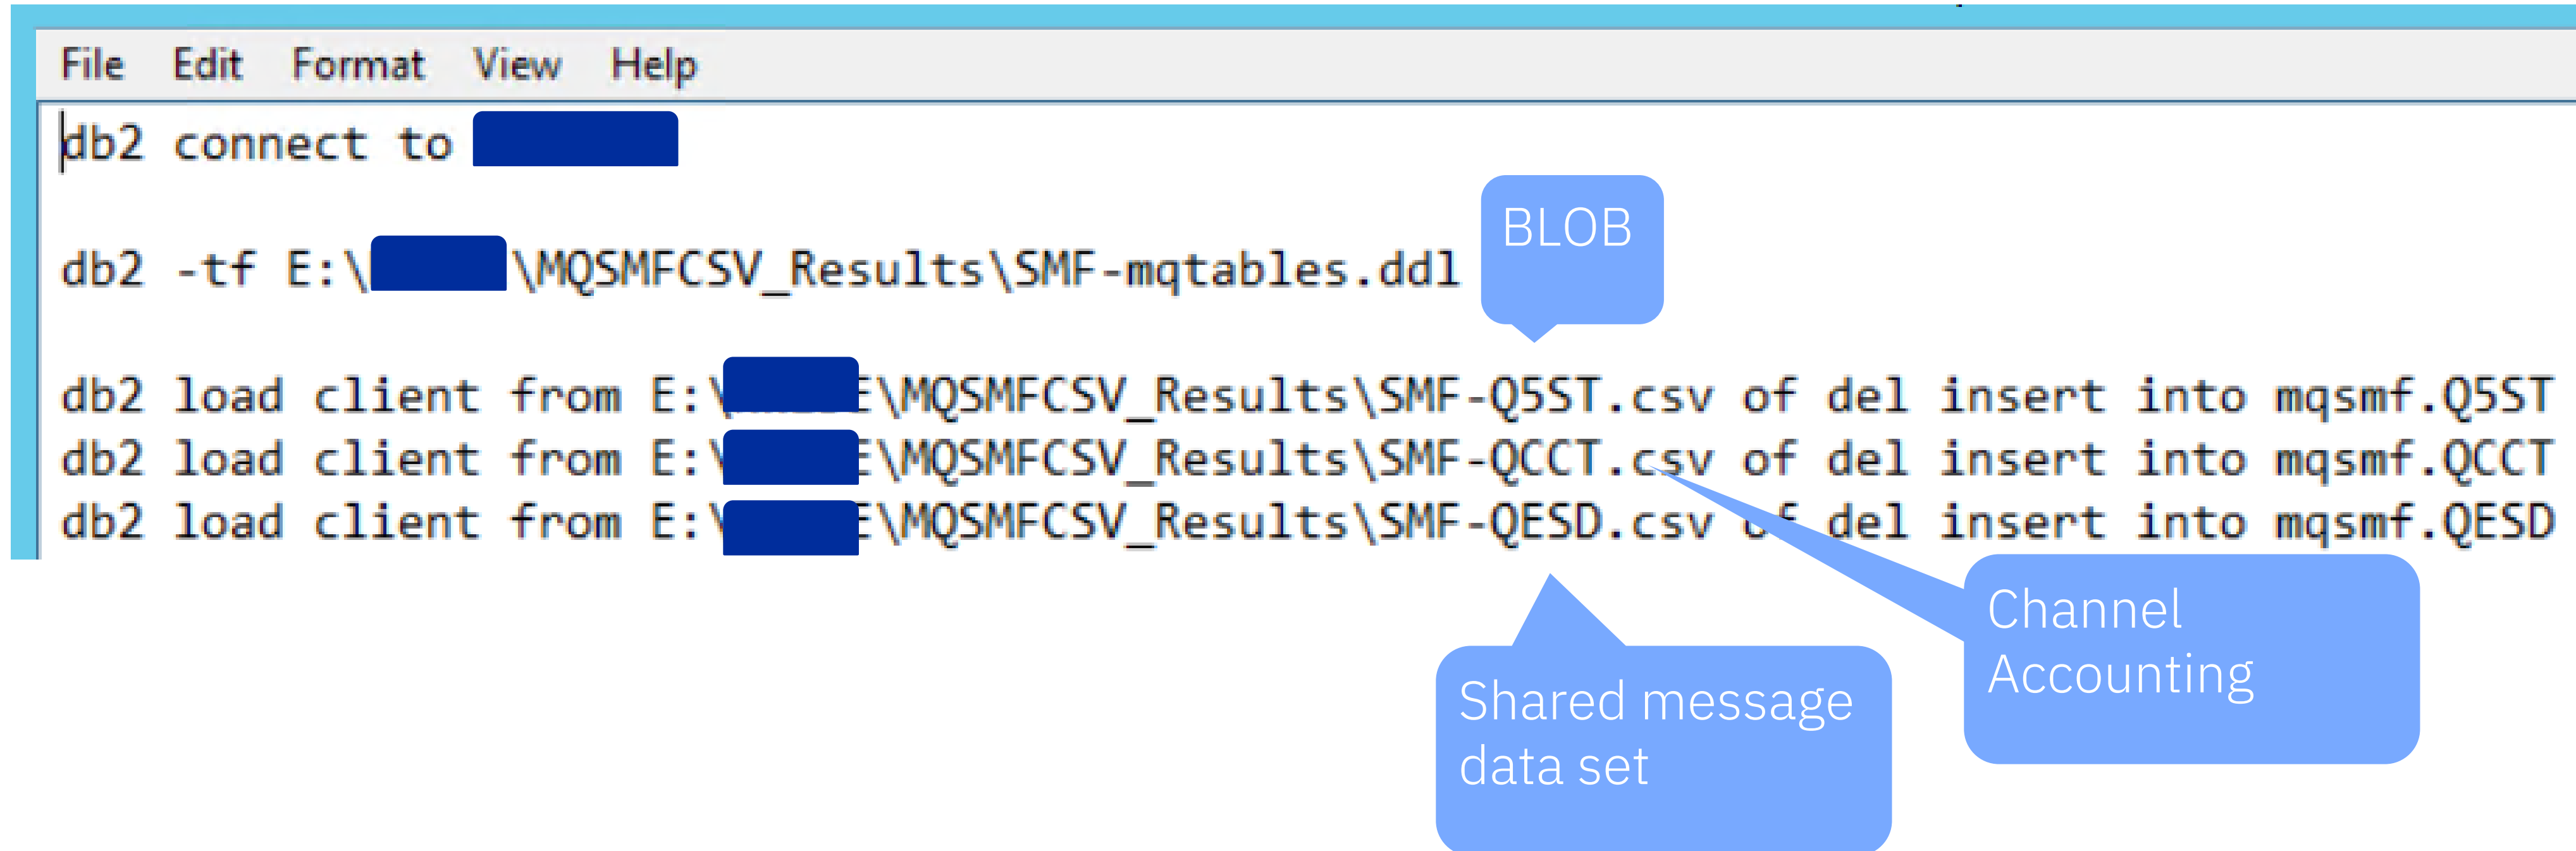
Create DB2 database(s)

```
Db2 CREATE DATABASE XXXXXX AUTOMATIC STORAGE YES ON 'E:\XXXXXX'  
DBPATH ON 'C:' ;
```

Run MQSMFCSV

```
mqsmfcsv -f sql -i E:\customerdir\MQ115.bin -o E:\customerdir\MQSMFCSV_Results
```


Load the database(s)

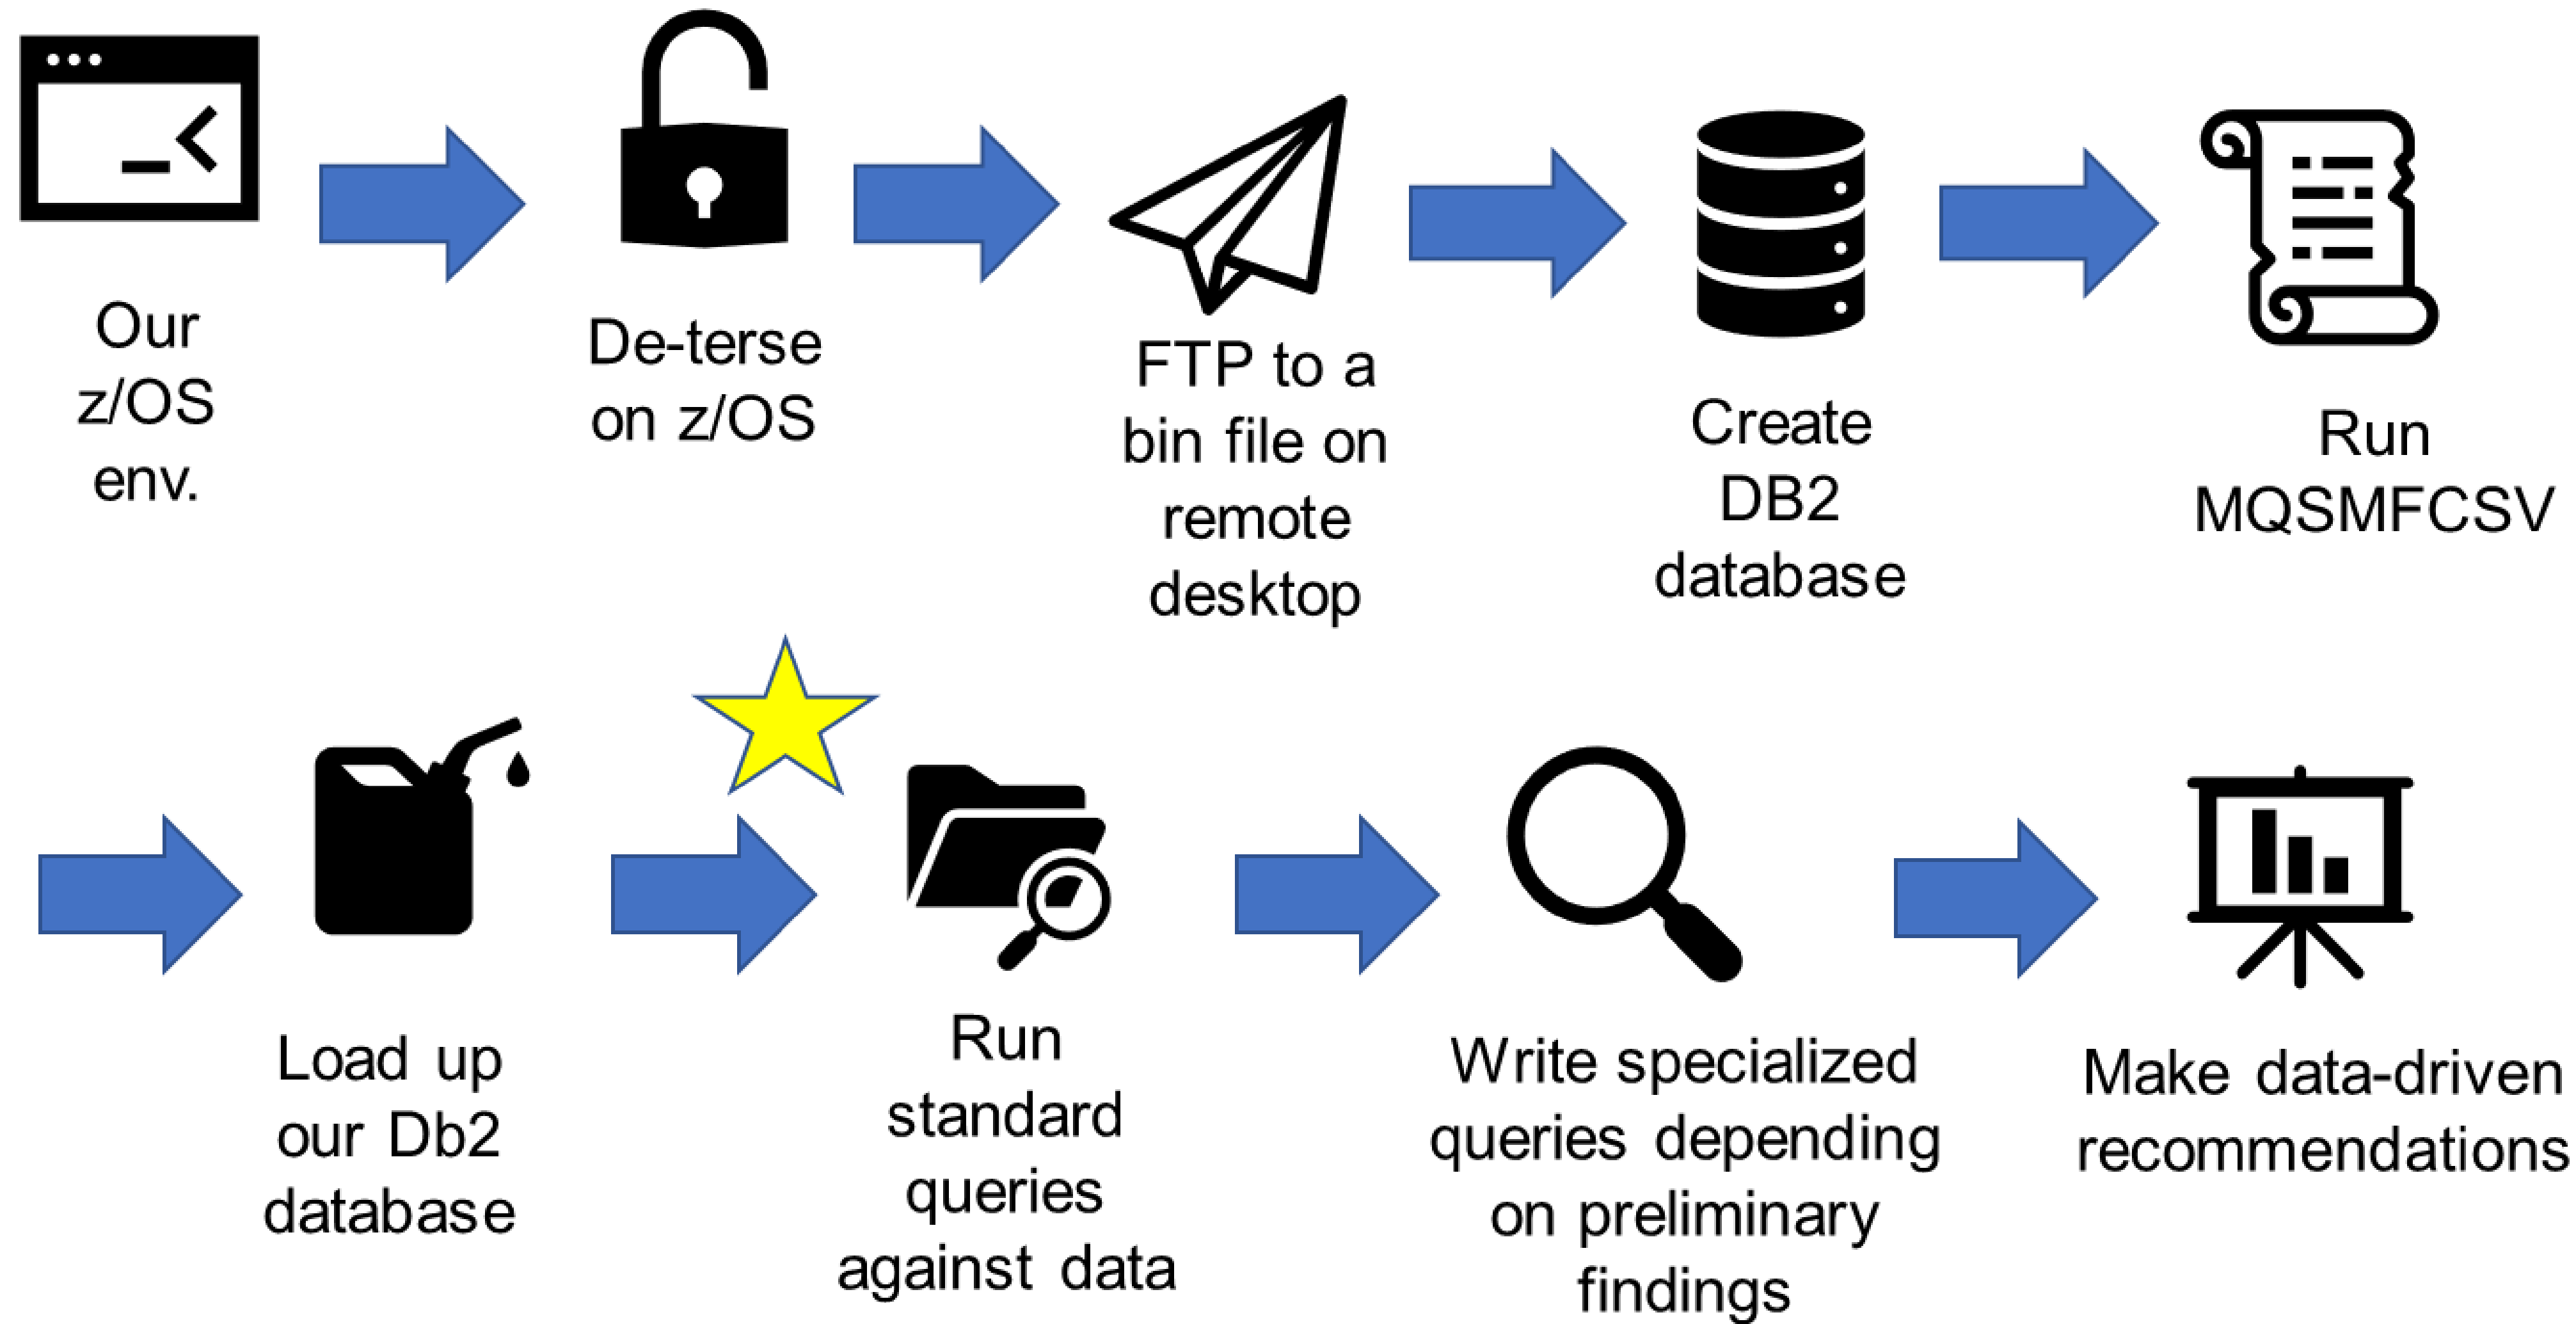


The screenshot shows a DB2 command window with a menu bar (File, Edit, Format, View, Help) and a command prompt. The commands entered are:

```
db2 connect to [REDACTED]  
  
db2 -tf E:\[REDACTED]\MQSMFCSV_Results\SMF-mqtables.ddl  
  
db2 load client from E:\[REDACTED]\MQSMFCSV_Results\SMF-Q5ST.csv of del insert into mqsmf.Q5ST  
db2 load client from E:\[REDACTED]\MQSMFCSV_Results\SMF-QCCT.csv of del insert into mqsmf.QCCT  
db2 load client from E:\[REDACTED]\MQSMFCSV_Results\SMF-QESD.csv of del insert into mqsmf.QESD
```

Annotations in blue callout boxes:

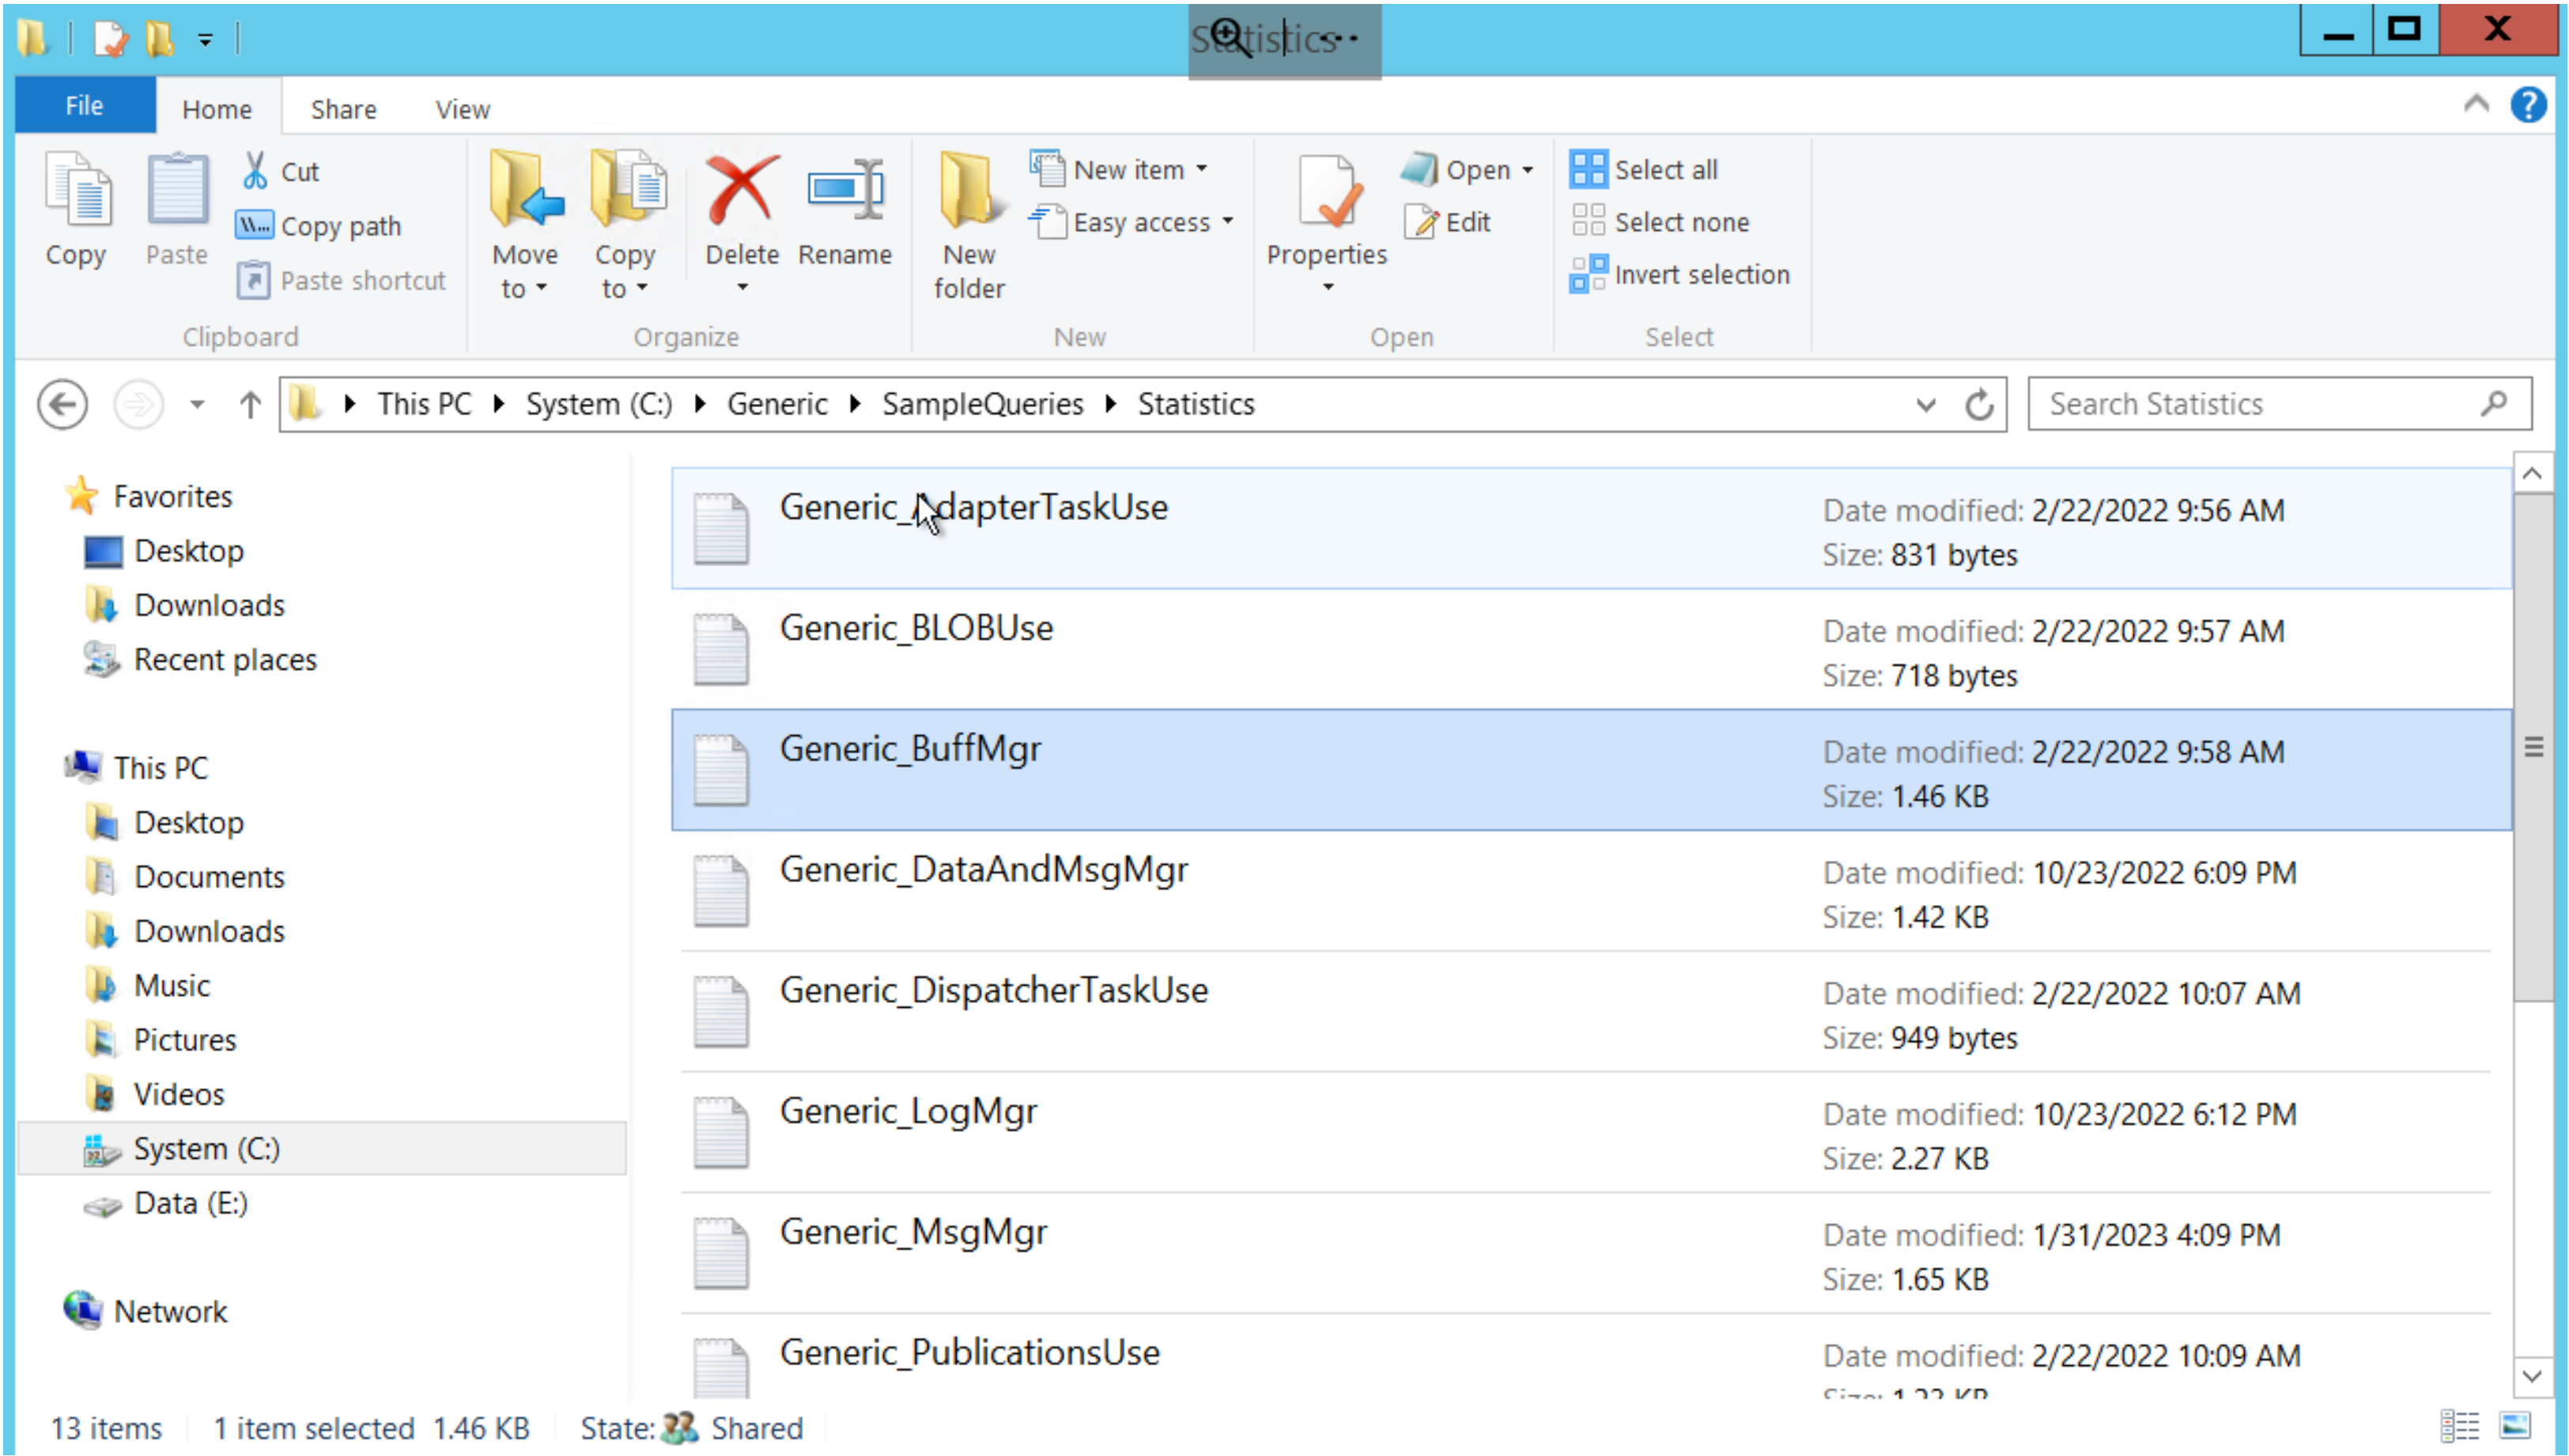
- A box labeled "BLOB" points to the `ddl` file in the second command.
- A box labeled "Shared message data set" points to the `SMF-Q5ST.csv` file in the third command.
- A box labeled "Channel Accounting" points to the `SMF-QCCT.csv` file in the fourth command.



Check who is here

```
SELECT DISTINCT LPAR,  
QMGR  
FROM MQSMF.QMST (message manager)  
ORDER BY LPAR.
```

Customize some queries



Run the queries against the data

```
DB2 -TVF E:\[redacted]\Queries\I[redacted]_AdapterTaskUse.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_BuffMgr.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_DataAndMsgMgr.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_DispatcherTaskUse.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_LogMgr.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_LongLatches.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_MsgMgr.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_OpenCloseCF.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_PublicationsUse.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_QueueSumm.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_SharedQueueSumm.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_Type11Latches.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_AdapterTaskUse.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_BuffMgr.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_DataAndMsgMgr.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_DispatcherTaskUse.txt
DB2 -TVF E:\[redacted]\Queries\I[redacted]_LogMgr.txt
```

Good starting point scenarios

1. Check if buffer pools are highly utilized
 2. If you see a highly utilized buffer pool, investigate which queues are most active for given buffer pool
 3. Move queues to less busy buffer pool
1. Check for get difference on your data manager
 2. Verify get difference is due to skipped messages
 3. If get difference is negative, look at queue summary to identify which queues are busy and not indexed
 4. Index queues
1. Check adapter task utilization
 2. Check adapter tasks aren't filled up due to scrolling
 3. If adapter tasks are all full, allocate more adapter tasks within reason

Further reading and resources:

Understanding MQ SMF Data: MP1B MQSMF pdf

Capacity planning and tuning guide for IBM MQ on z/OS: MP16 pdf

<https://github.com/ibm-wsc/mq-wildfire-mqv9zos/>

[MQ and SMF - How might I process the data? \(ibm.com\)](#)

[GitHub - ibm-messaging/mq-smf-csv: Simple formatter for MQ's SMF records to assist with import to spreadsheets and databases](#)

[Dorothy-Quincy/generic_smf \(ibm.com\)](#)

[Dorothy-Quincy/smf_chinit_task_interpretation \(ibm.com\)](#)

Thank you!

Dorothy Quincy

Dorothy.quincy@ibm.com